

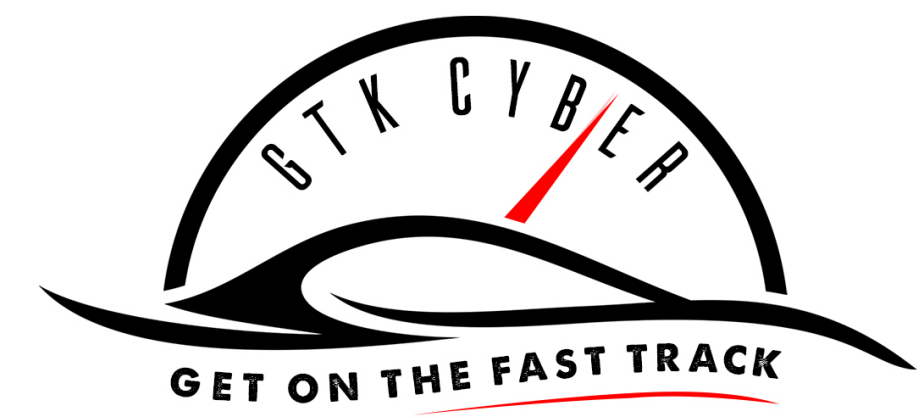






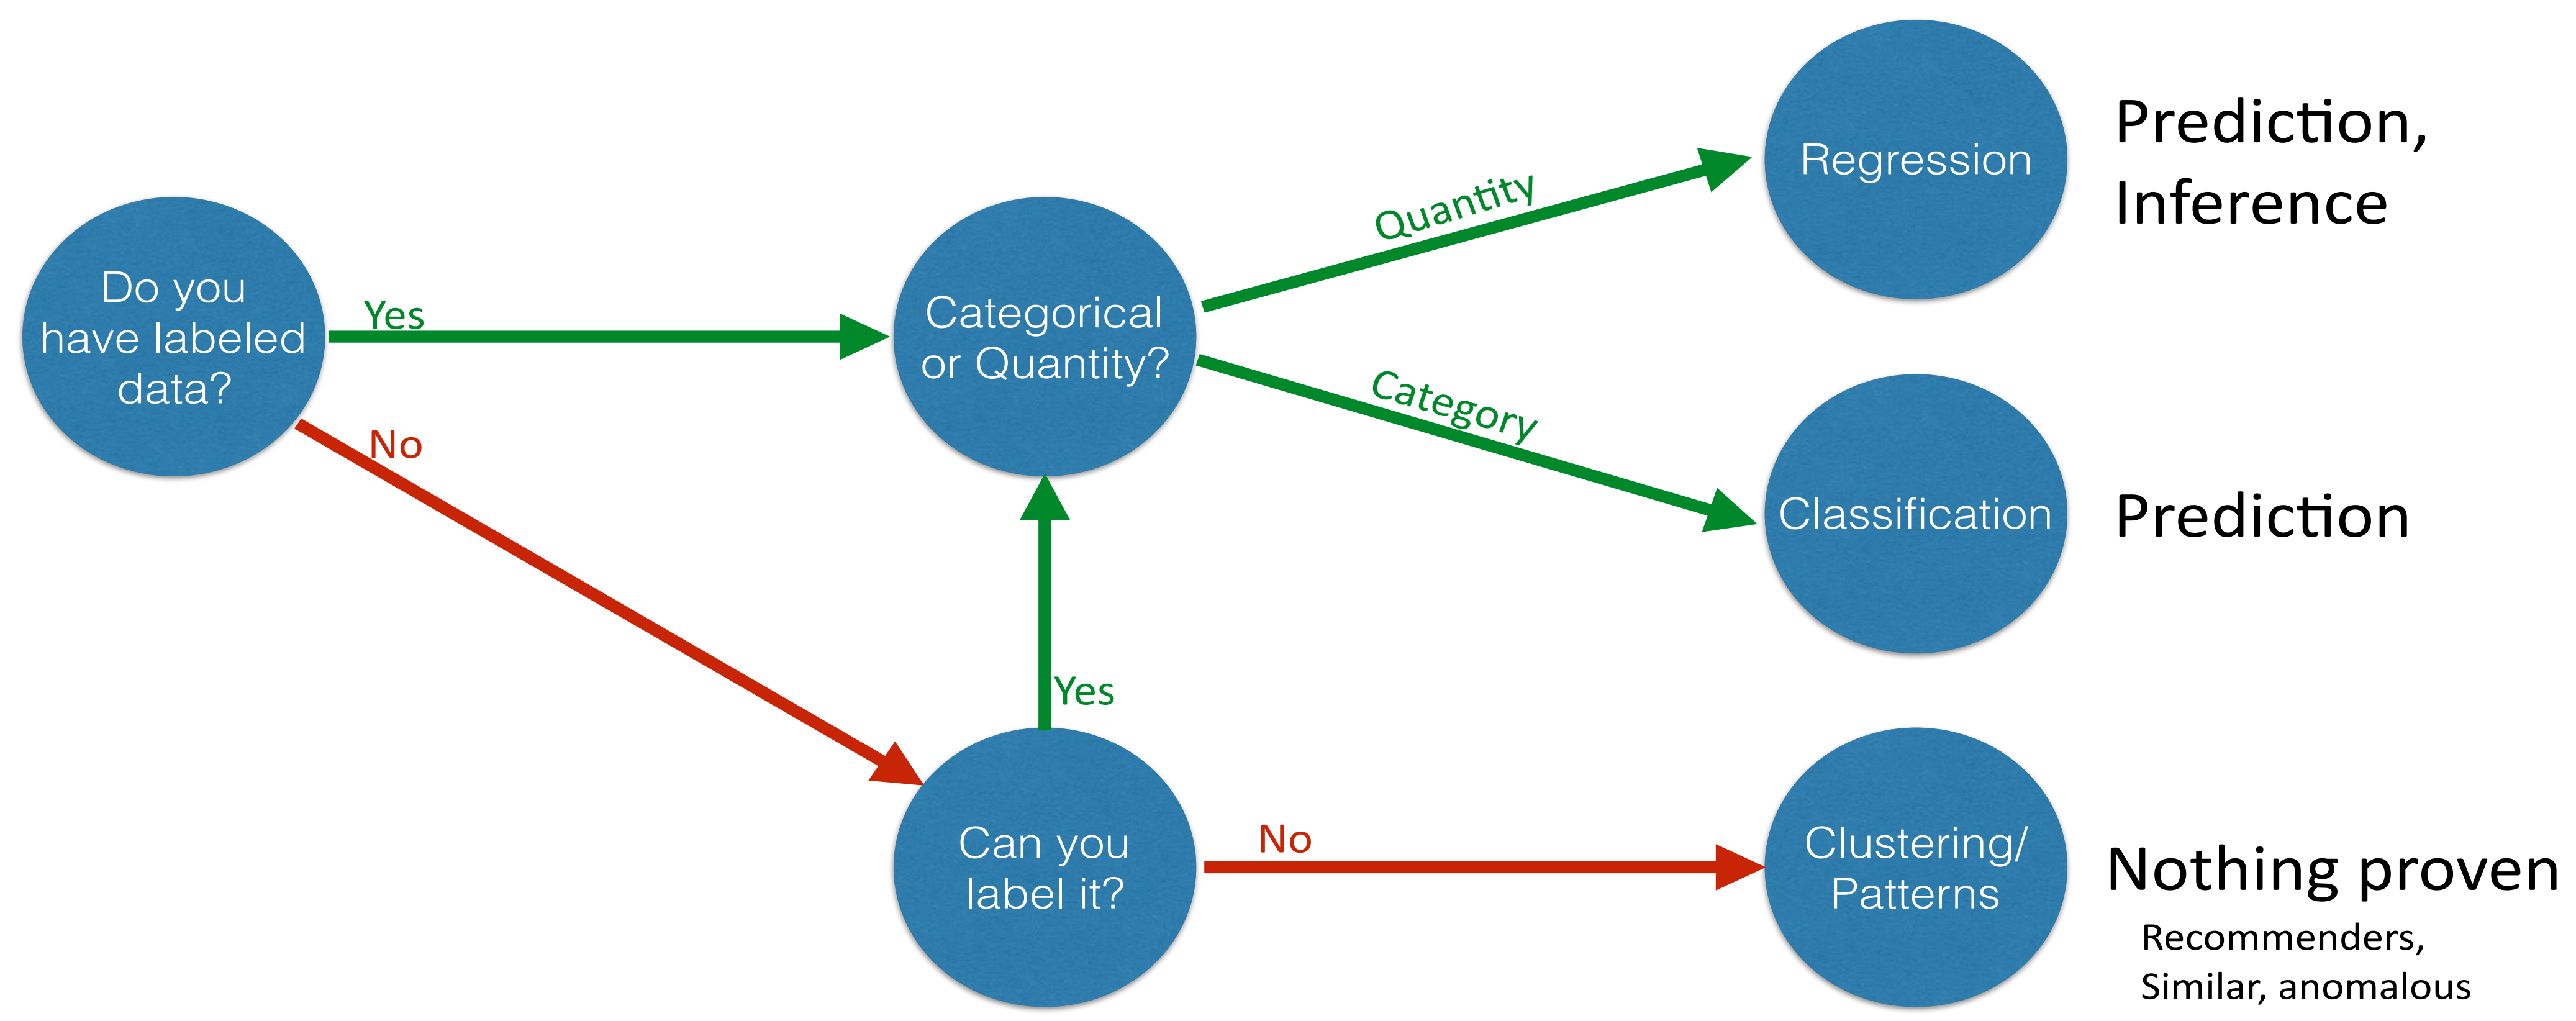
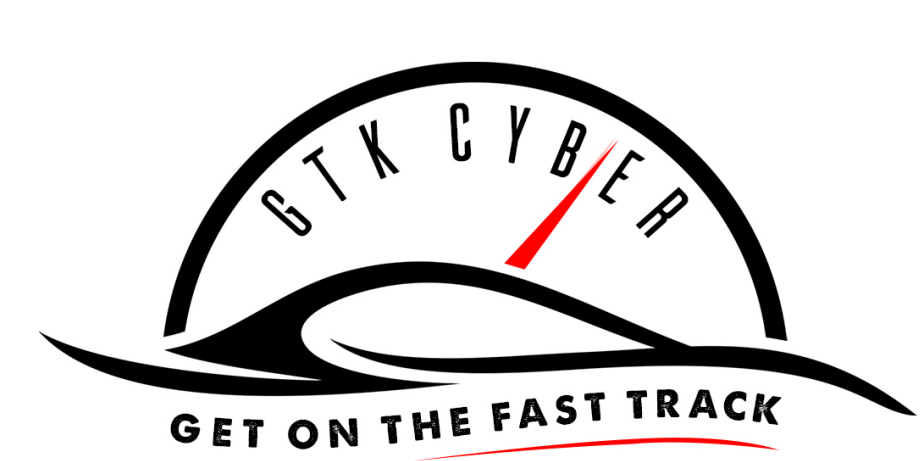
# Machine Learning for Security Professionals - Day 3

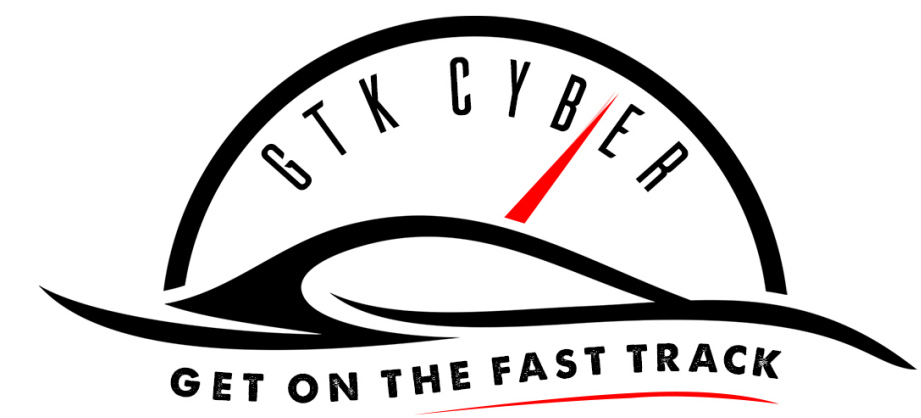
Unsupervised Learning: Clustering



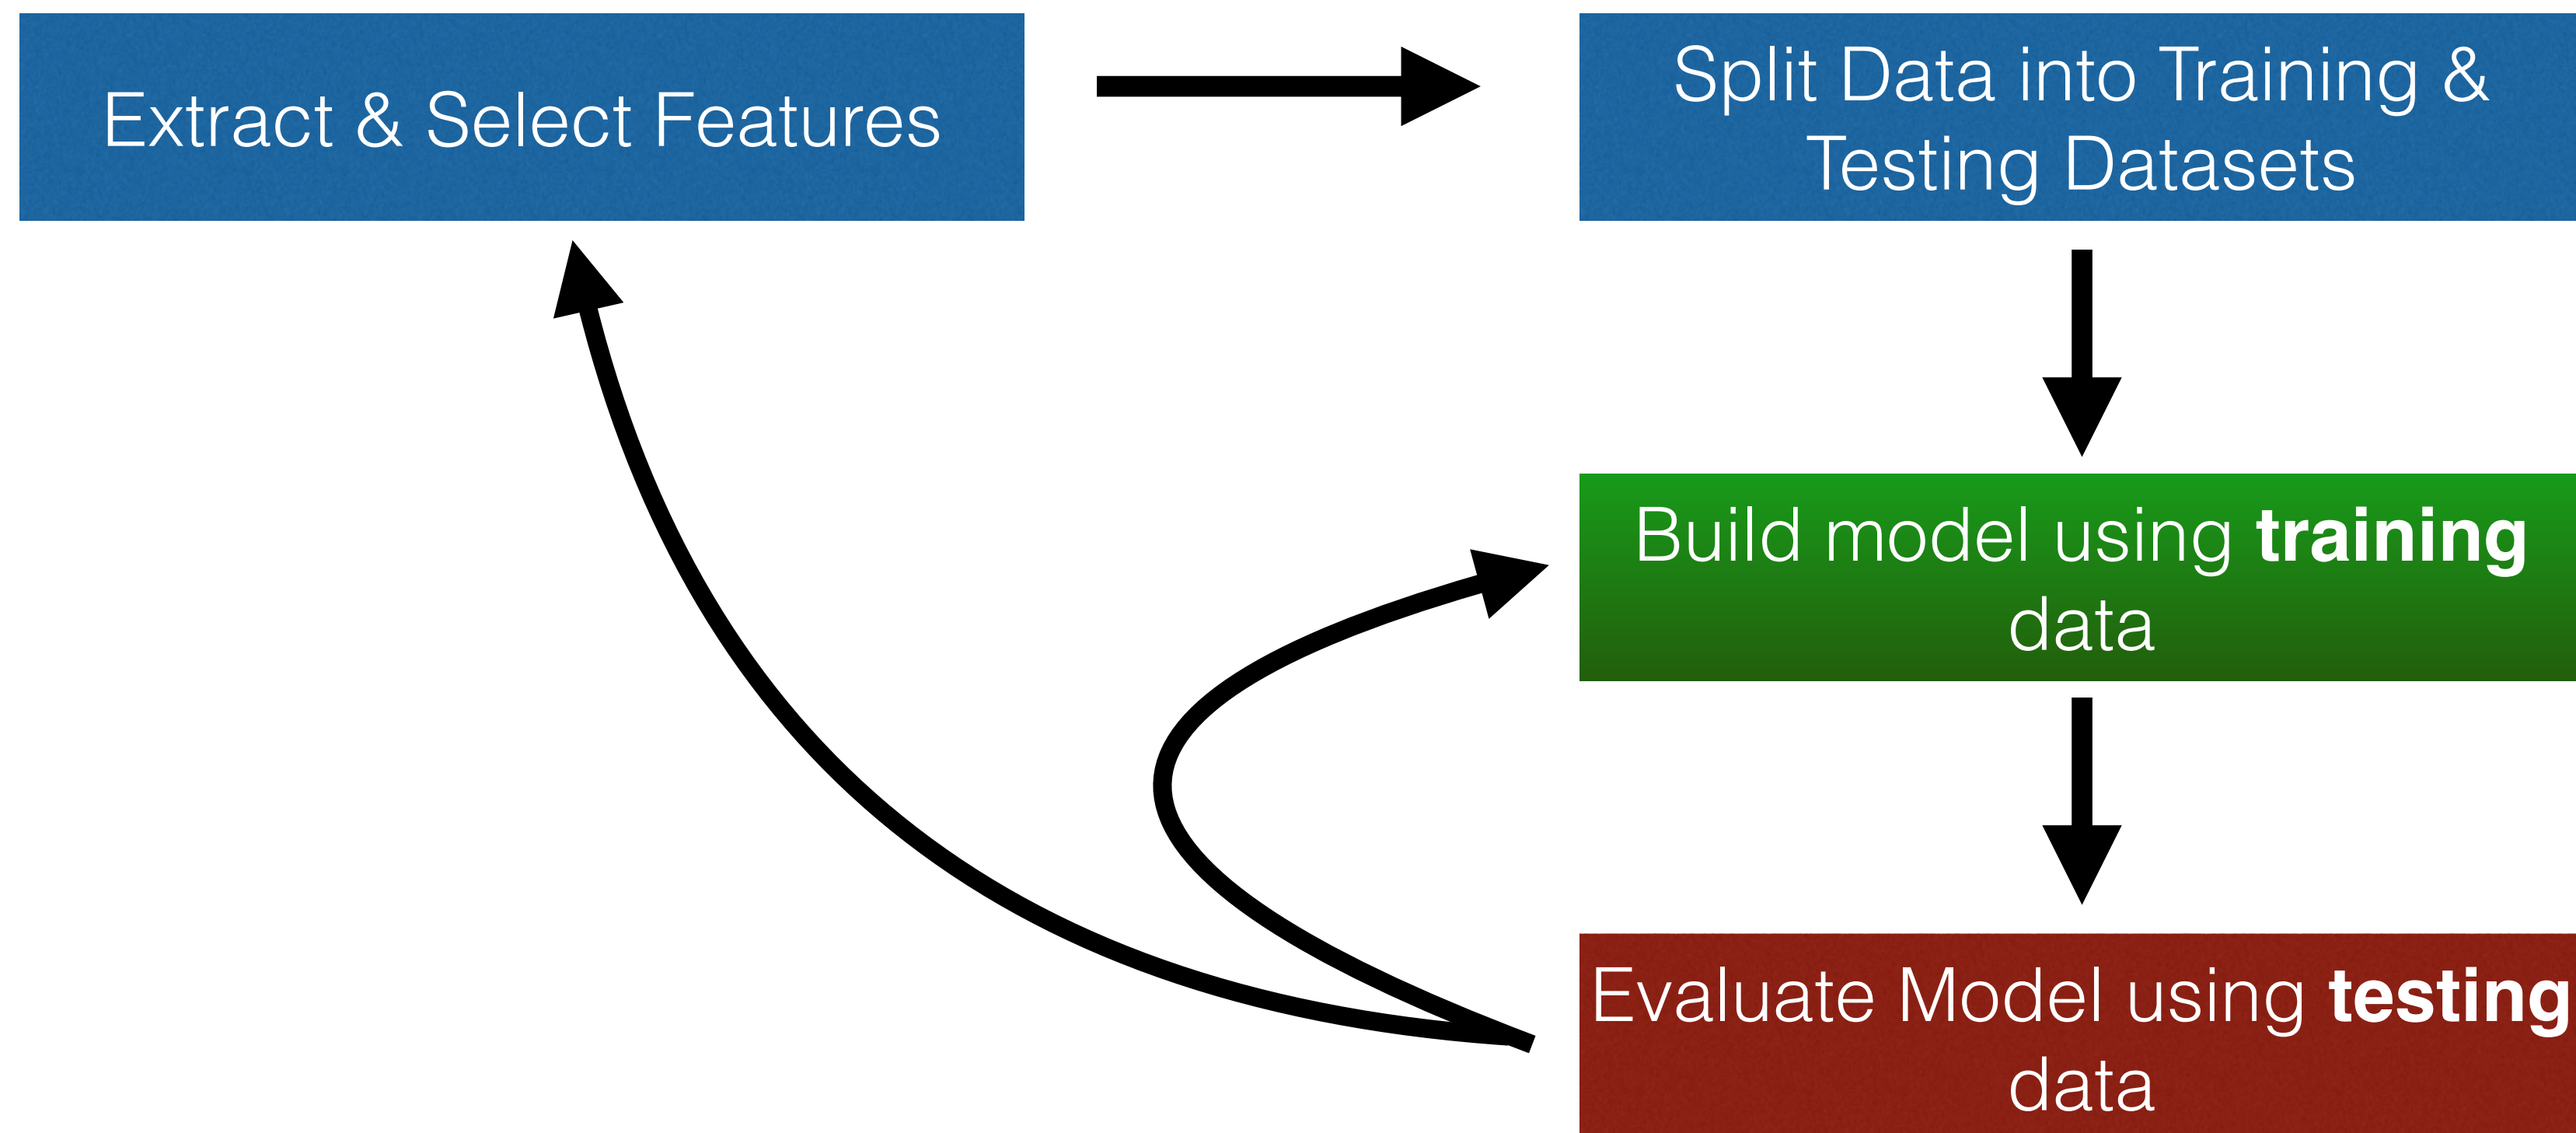
# Agenda for Today

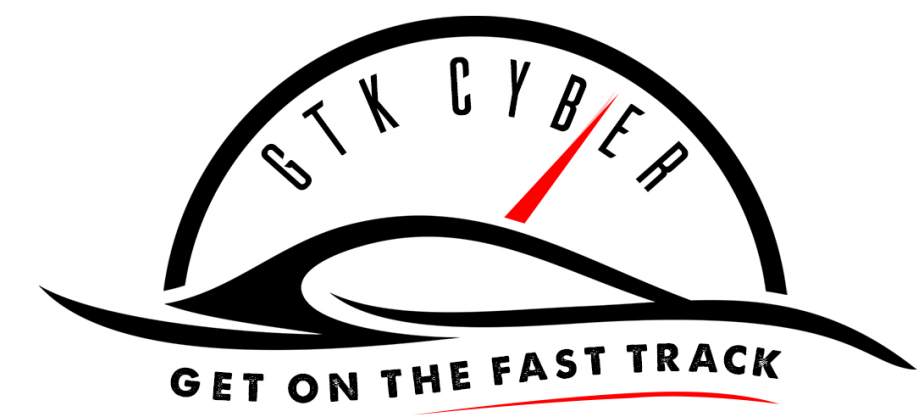
- Measuring Distances
- Math free overview of clustering techniques
- Pipelines and pickles






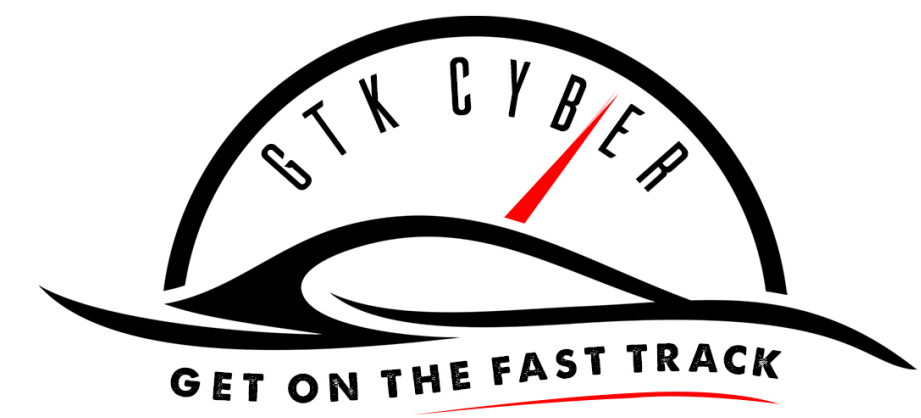
# Supervised ML Process





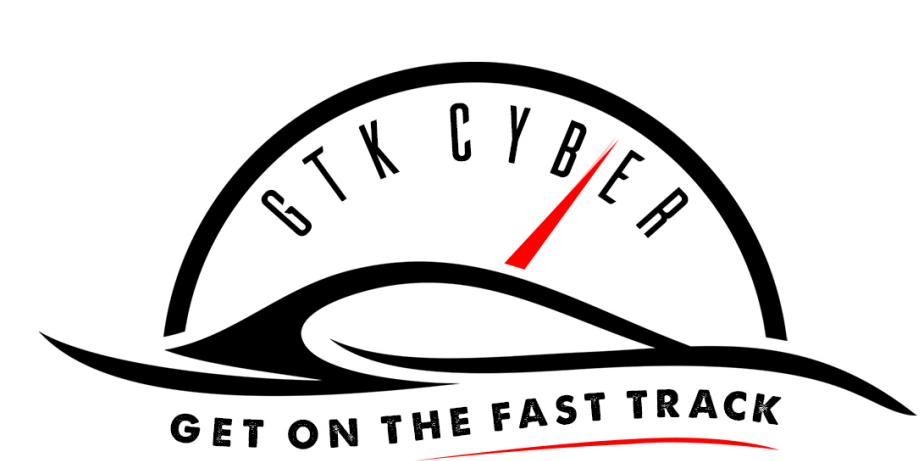
# Unsupervised ML Process





# Unsupervised Clustering Algorithm


1. Select Features
2. Calculate a distance measure
3. Apply a clustering algorithm
4. Validate?



# Which Departments are Similar?


	Malware events
Dept1	6
Dept2	1
Dept3	8





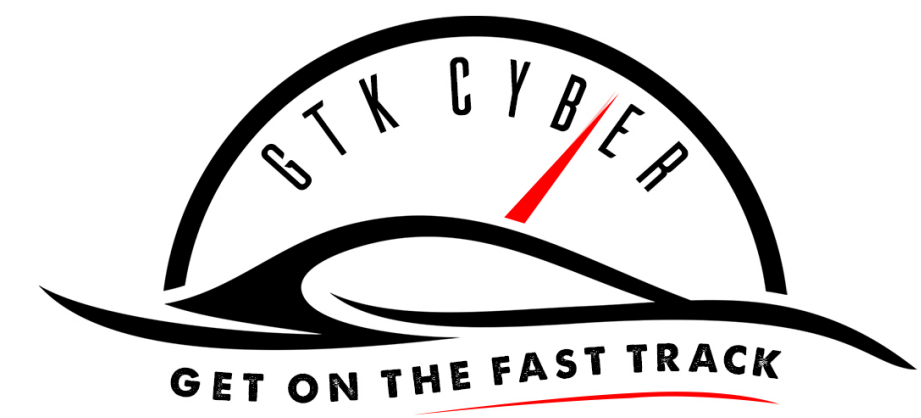
# Which Departments are Similar?

	Malware events	Phishing
Dept1	6	6
Dept2	1	2
Dept3	8	1



# Which Departments are Similar?

	Malware events	Phishing	Open Tickets
Dept1	6	6	3
Dept2	1	2	1
Dept3	8	1	9



# Computing Distance

	Malware events
Dept1	6
Dept2	1
Dept3	8

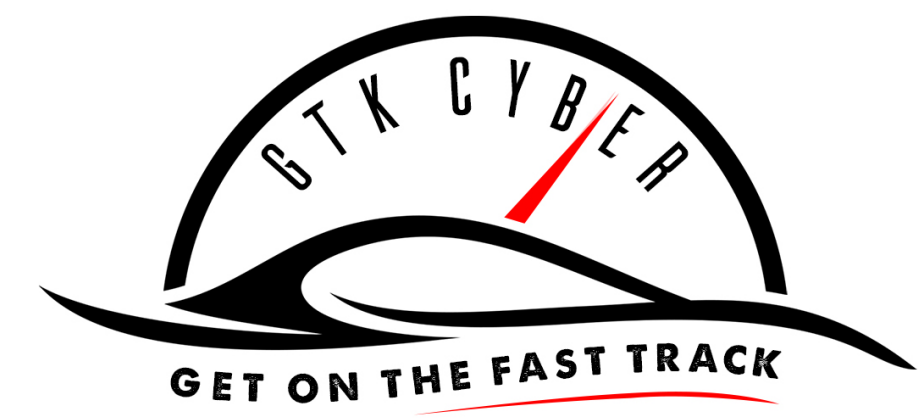
Compare:

Dept1 to Dept2:  $| 6 - 8 | = 2$

Dept2 to Dept3:  $| 1 - 8 | = 7$

Dept1 to Dept3:  $| 6 - 8 | = 2$






# Two-Dimensional Distance

	Malware events	Phishing
Dept1	6	6
Dept2	1	2
Dept3	8	1


## Multiple Distance methods

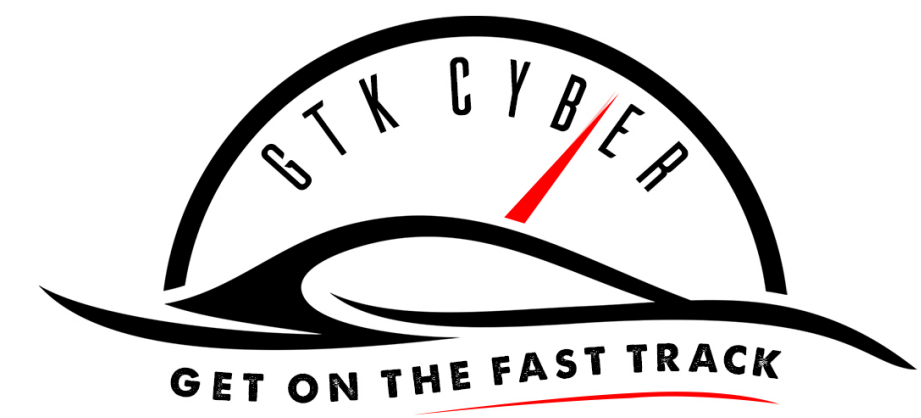
- Euclidean
  - Manhattan
  - Maximum
  - Canberra
  - Binary
  - Minkowski
- ... (to name a few)



# Two-Dimensional Distance

	Malware events	Phishing
Dept1	6	6
Dept2	1	2
Dept3	8	1






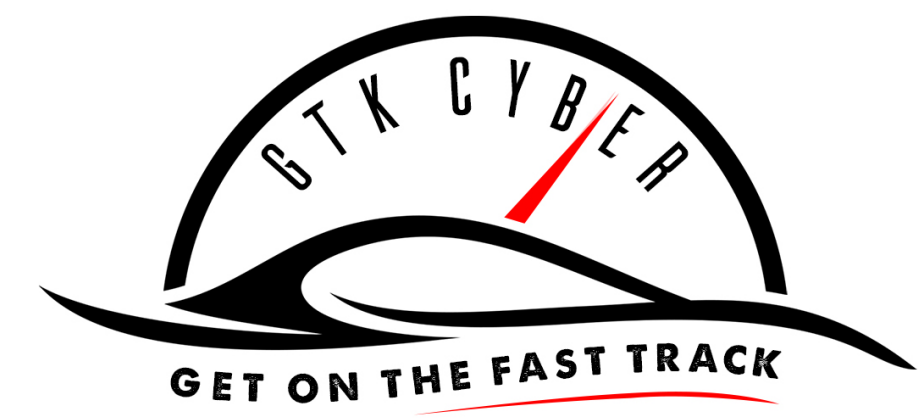
# Two-Dimensional Distance

Euclidean very common and easy to understand

	Malware events	Phishing
Dept1	6	6
Dept2	1	2
Dept3	8	1






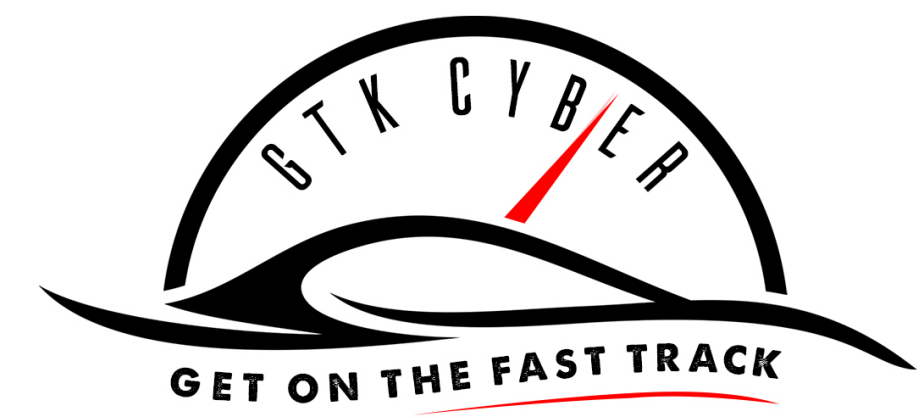


# Two-Dimensional Distance

Euclidean very common and easy to understand:  $a^2 + b^2 = c^2$

	Malware events	Phishing
Dept1	6	6
Dept2	1	2
Dept3	8	1




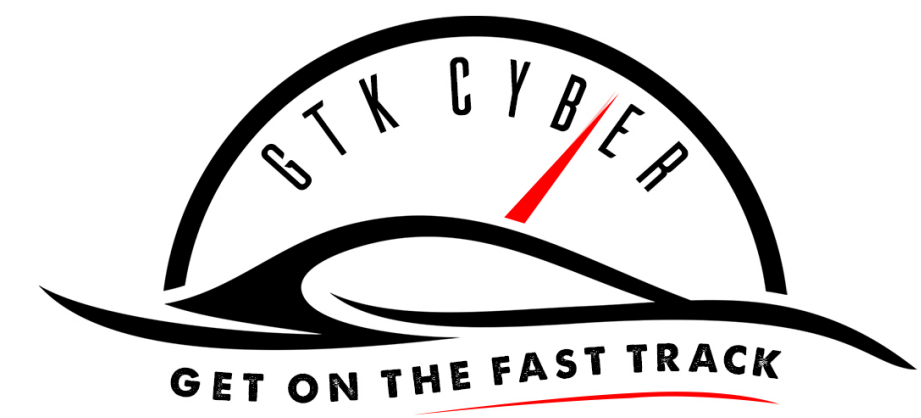


# Two-Dimensional Distance

Manhattan also easy to comprehend:  $a + b$

	Malware events	Phishing
Dept1	6	6
Dept2	1	2
Dept3	8	1





# Computing Distance

	Malware events	Phishing
Dept1	6	6
Dept2	1	2
Dept3	8	1

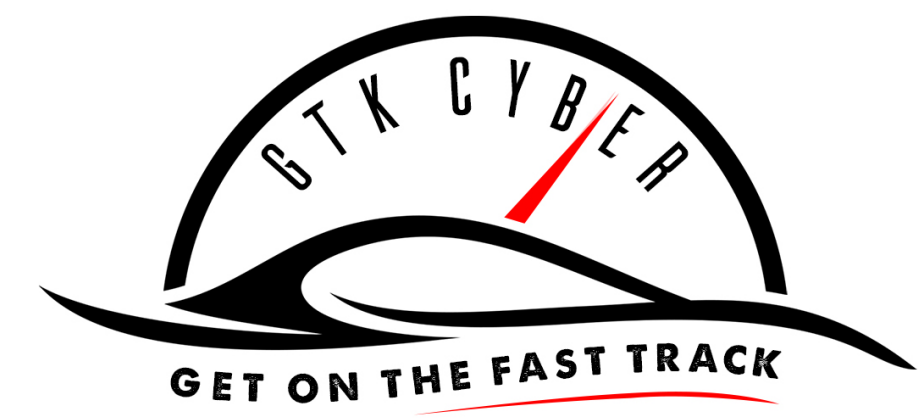
Compare:

Dept1 to Dept2:  $\sqrt{(6-1)^2 + (6-2)^2} = \mathbf{6.4}$

Dept2 to Dept3: ... = **7.1**

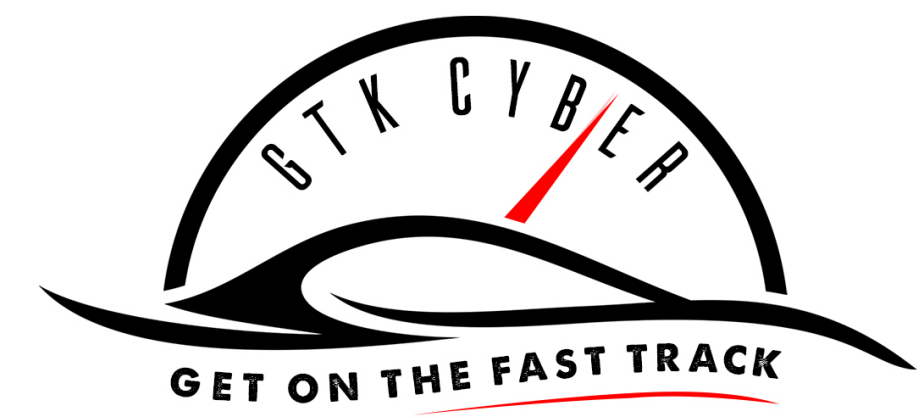
Dept1 to Dept3: ... = **5.4**






# Euclidean Distance calculations

```
def dist(x,y):  
    return np.sqrt(np.sum((x-y)**2))  
  
> mat = np.array([[ 6,6,3 ], [1,2,1], [8,1,9]])  
> dist(mat[0], mat[1])  
6.7082039324993694  
  
> dist(mat[1], mat[2])  
10.677078252031311  
  
> dist(mat[0], mat[2])  
8.0622577482985491
```



# Which Departments are Similar?


	Malware events	Phishing	Open Tickets	
Dept1	6	6	3	6.7 8.1 10.7
Dept2	1	2	1	
Dept3	8	1	9	




**Figure 19.**

Clustering on breach data  
across industries

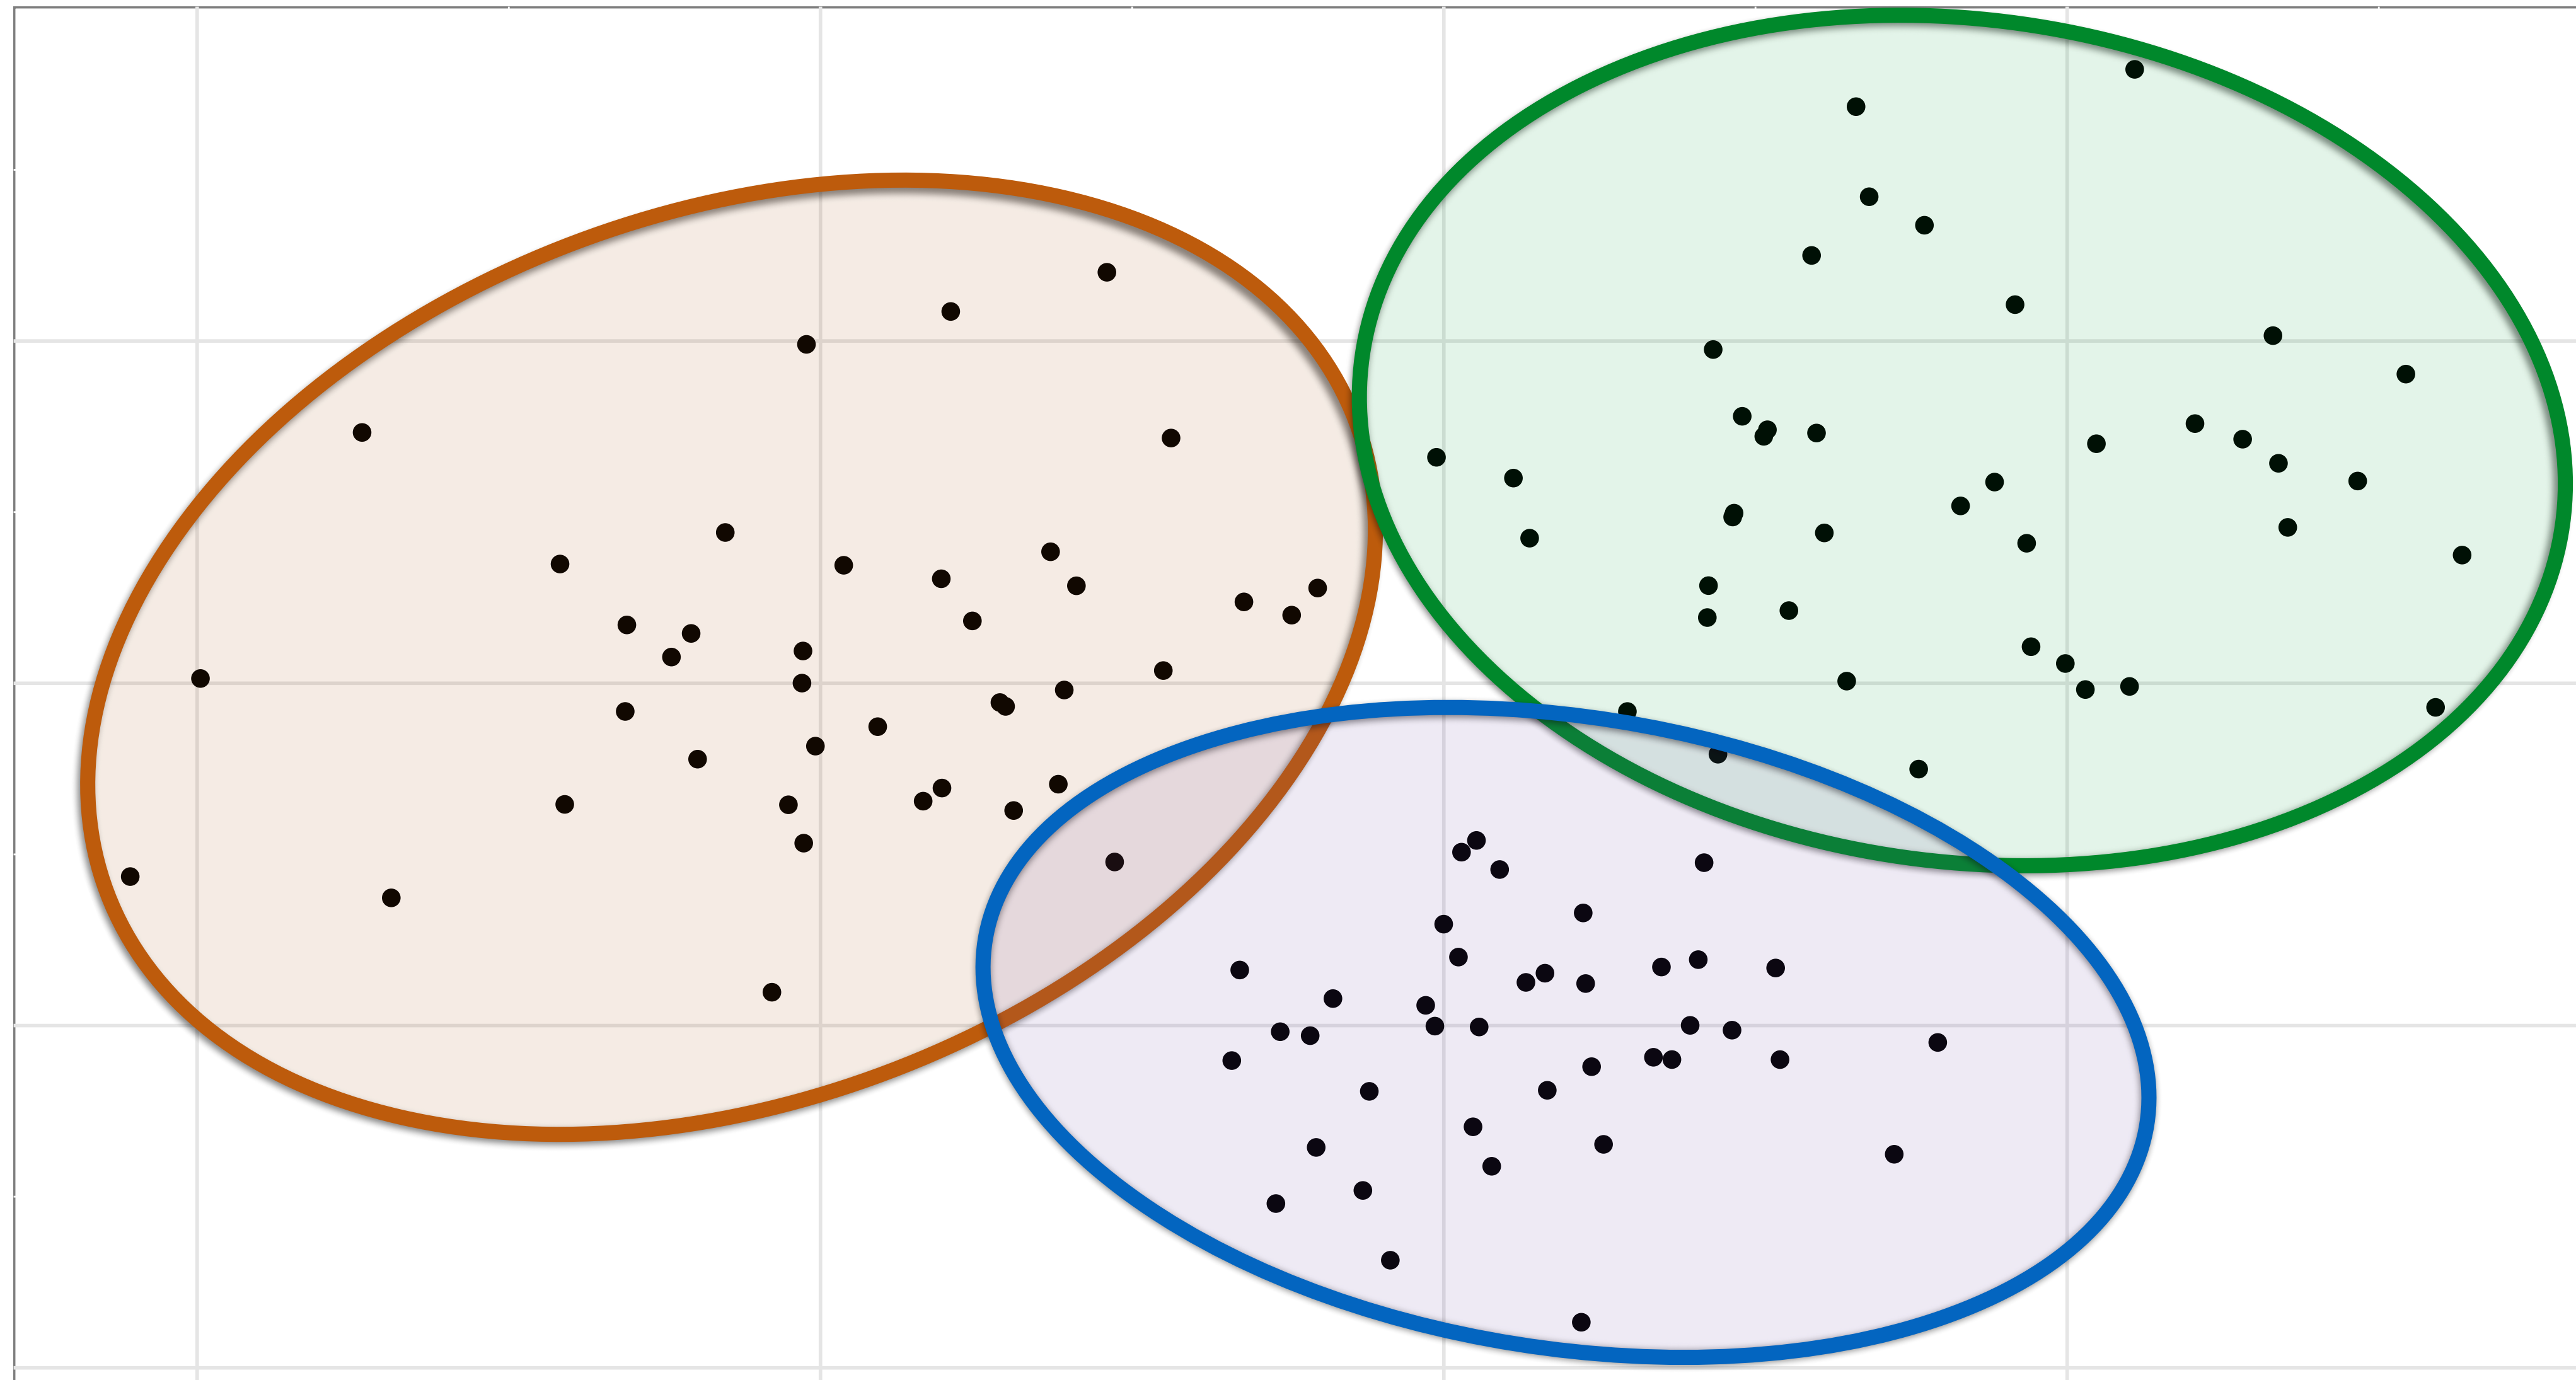




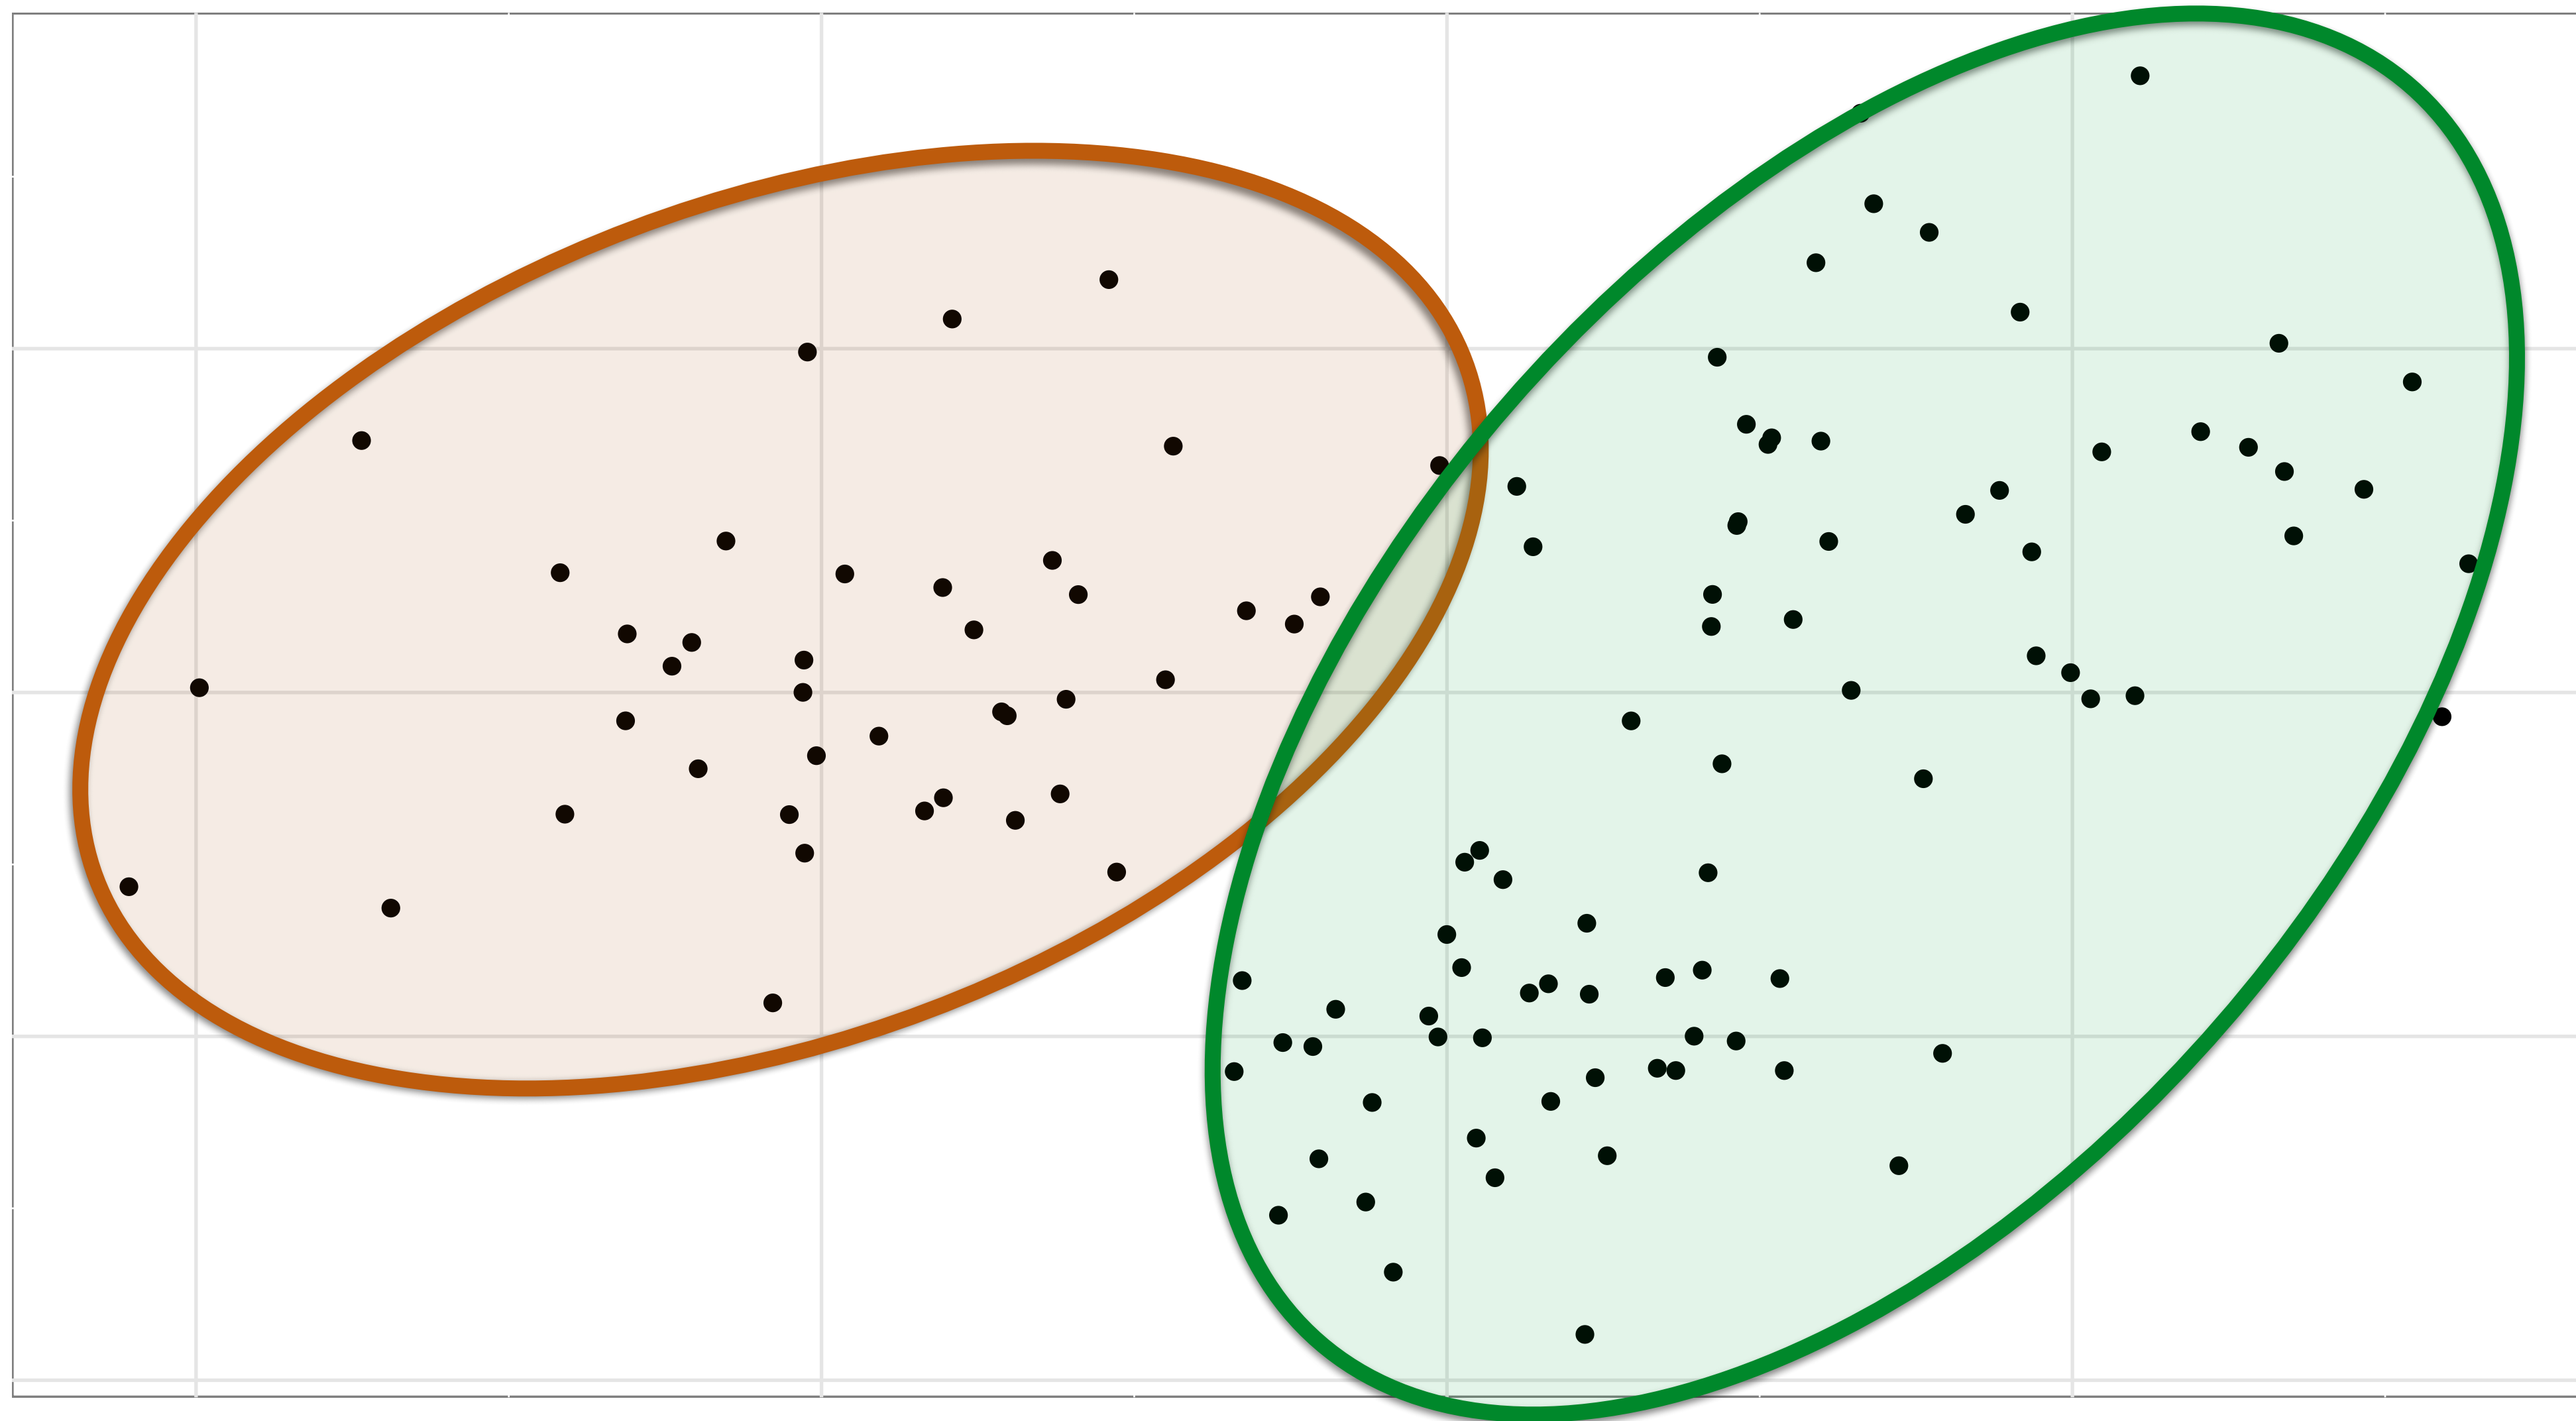
# Clustering...



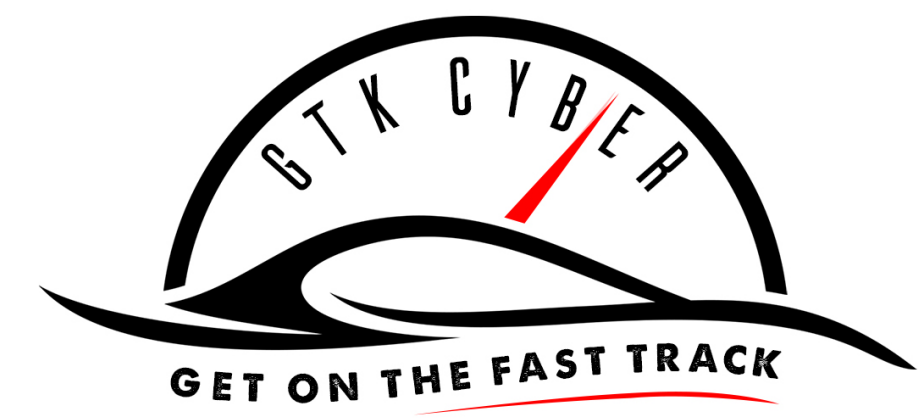
# Clustering...



# Clustering...









# K-Means

Before starting, pick the number of clusters,  $K$


1. Pick  $K$  random centroids within data range
2. Assign each data point to the nearest centroid
3. Move centroid to center of assigned points
4. Repeat steps 2 and 3 until centroid stops shifting




# Step 1: Pick 3 random centroids within data range




# Step 2: Assign each data point to the nearest centroid (1)



# Step 3: Move centroid to center of assigned points (1)




# Step 2: Assign each data point to the nearest centroid (2)







# Step 3: Move centroid to center of assigned points (2)




## Step 2: Assign each data point to the nearest centroid (3)




# Step 3: Move centroid to center of assigned points (3)



# Step 2: Assign each data point to the nearest centroid (4)




# Step 3: Move centroid to center of assigned points (4)






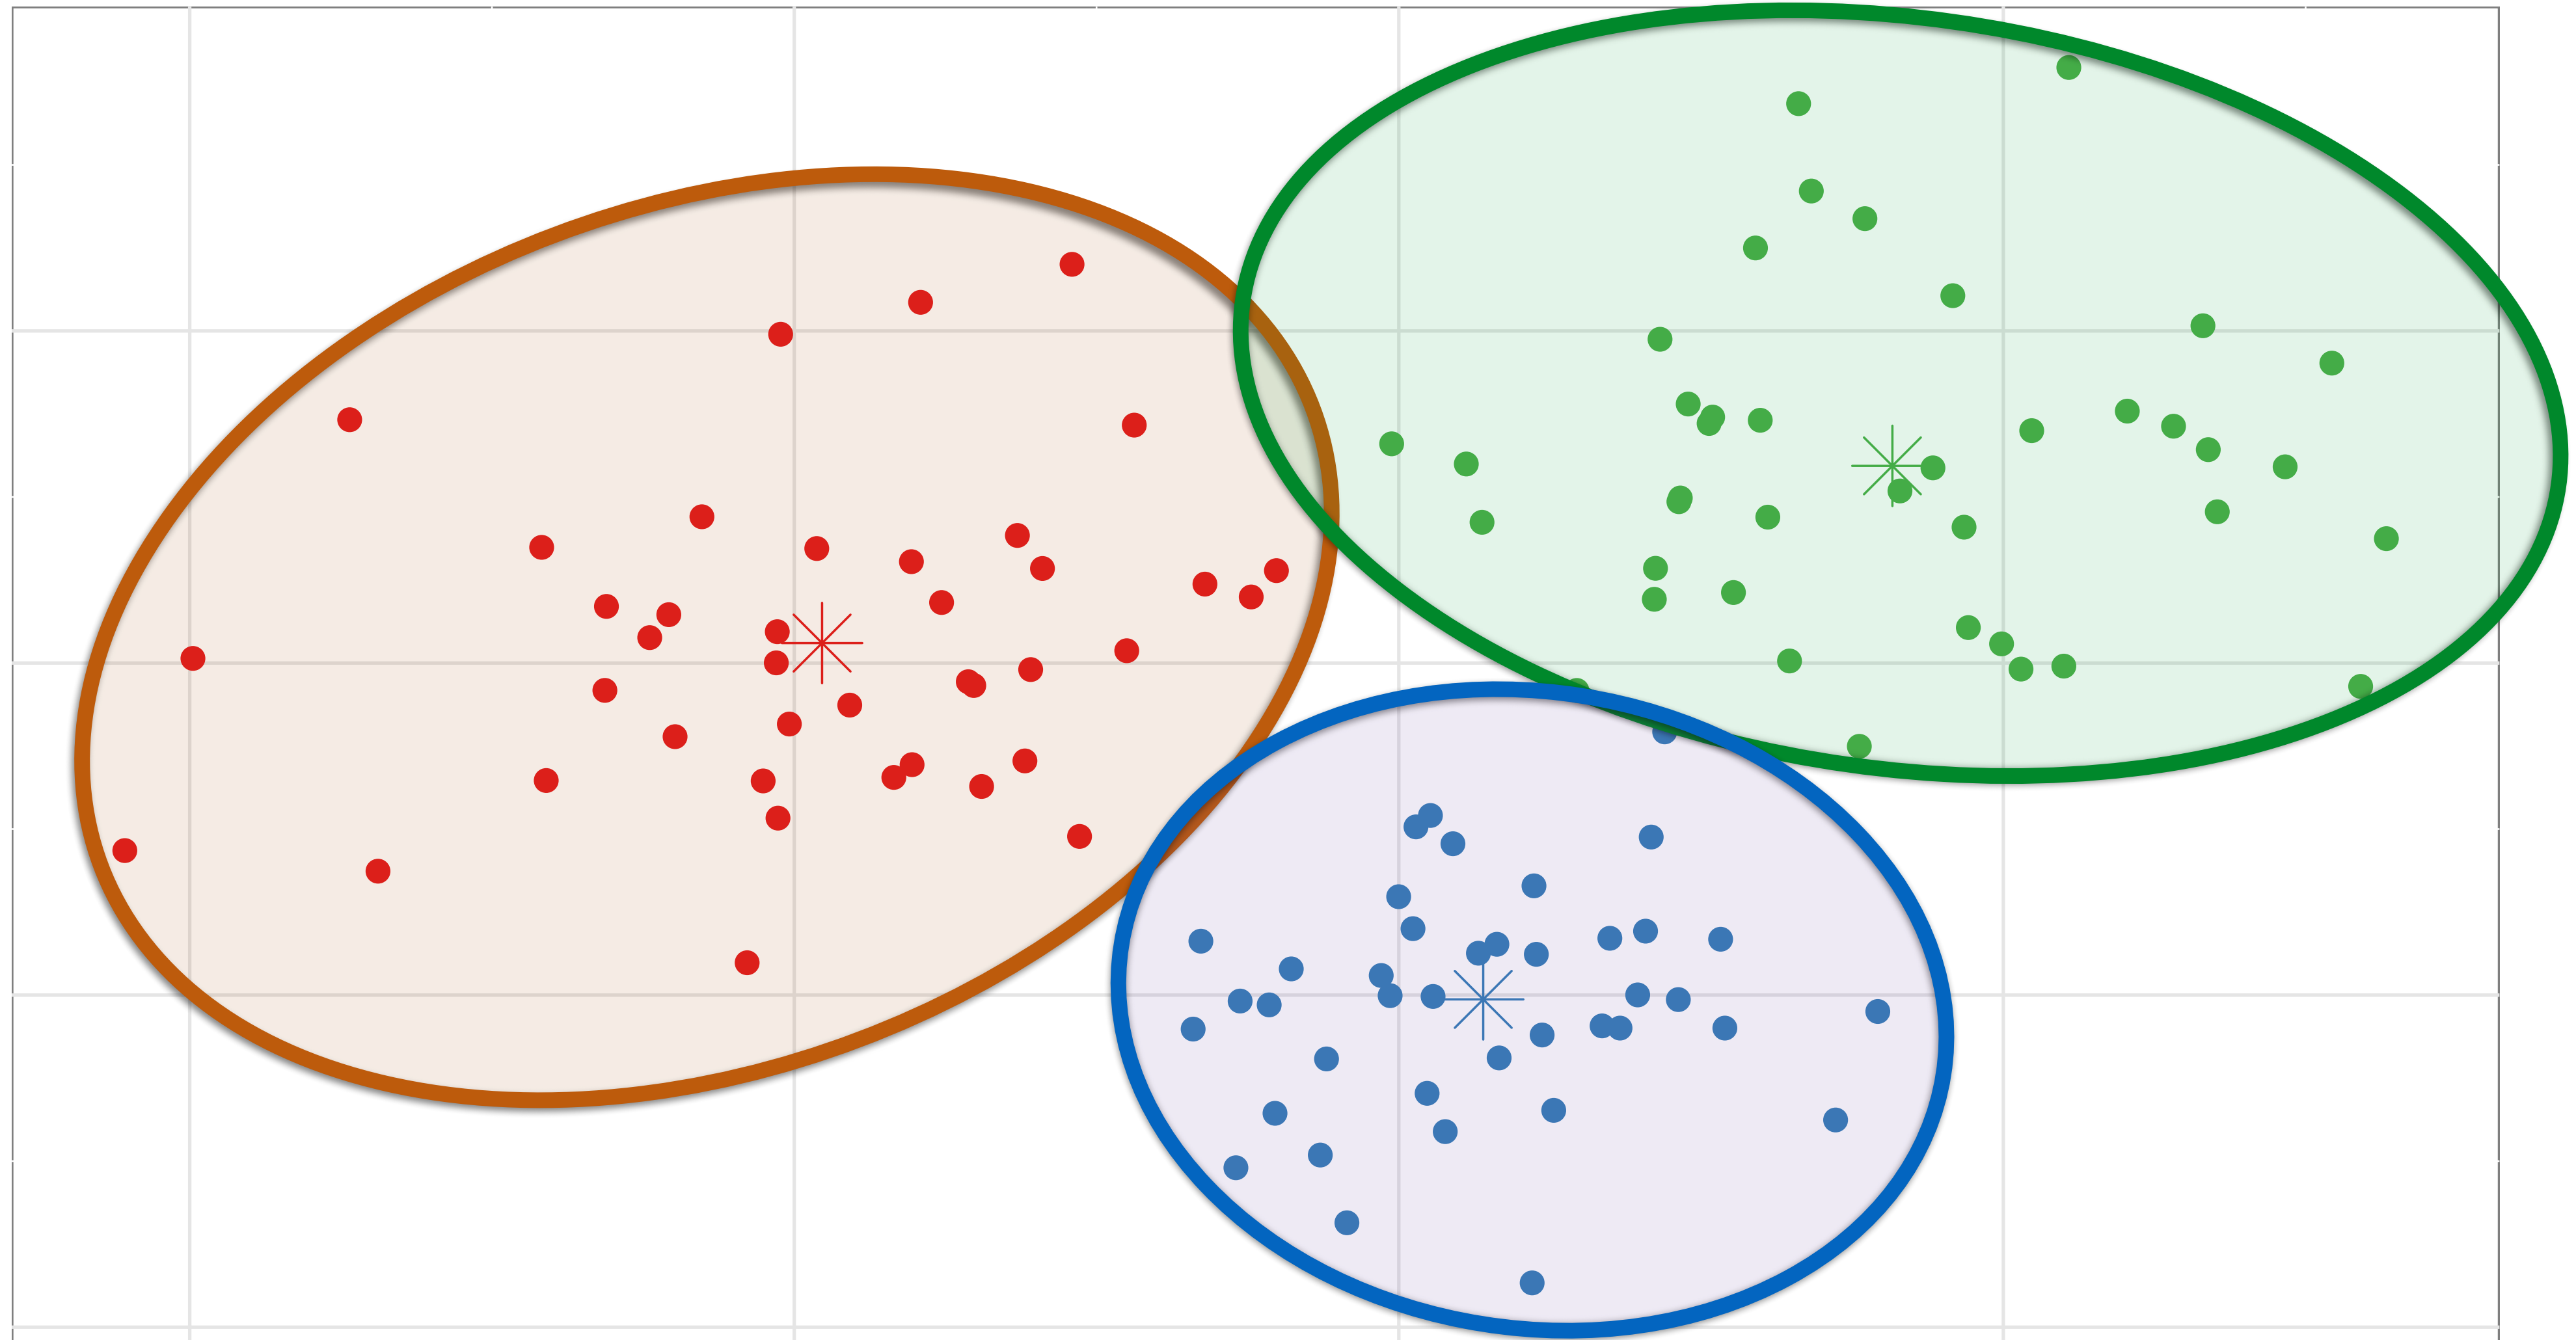
# Step 2: Assign each data point to the nearest centroid (5)

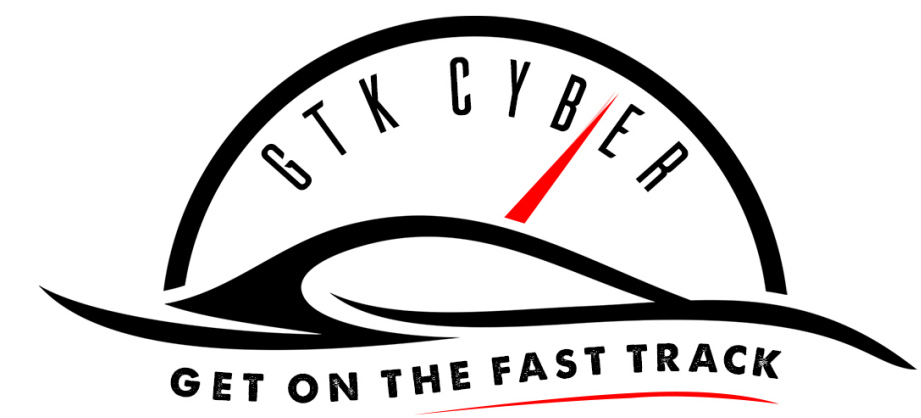


# Step 3: Move centroid to center of assigned points (5)



# Step 3: Move centroid to center of assigned points (5)

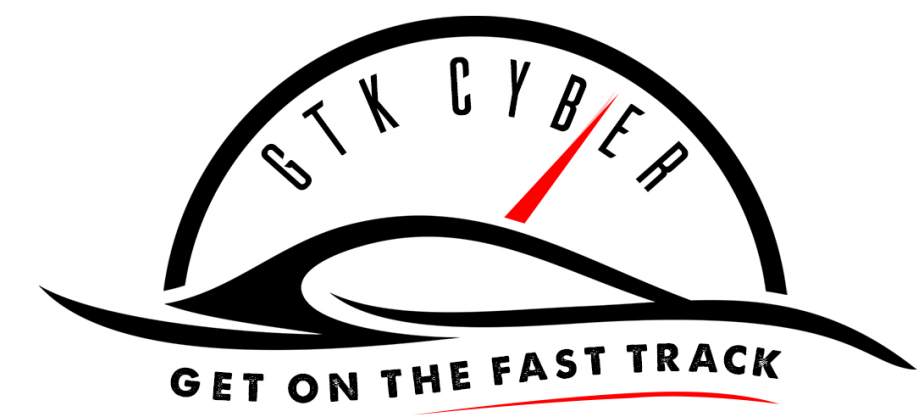




# K-Means: Got a problem with it?

Before starting, pick the number of clusters,  $K$

1. Pick  $K$  random centroids within data range
2. Assign each data point to the nearest centroid
3. Move centroid to center of assigned points
4. Repeat steps 2 and 3 until centroid stops shifting



# K-Means: Got a problem with it?

Before starting, pick the number of clusters, K *Subjective*


1. Pick K random centroids within data range *Not Repeatable*

2. Assign each data point to the nearest centroid *Sensitive to Scale*


3. Move centroid to center of assigned points


4. Repeat steps 2 and 3 until centroid stops shifting






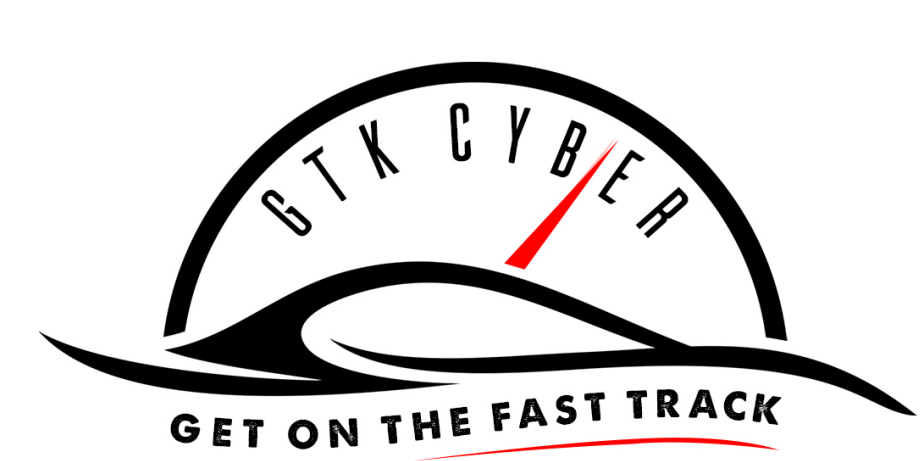
# How many clusters?







# Random Start...






# Random Start...







# Random Start...







# K-Means




“...it’s too easy to throw k-means on your data, and nevertheless get a result out (that is pretty much random, but you won't notice).”

— Anony-Mousse


<http://stats.stackexchange.com/questions/133656/how-to-understand-the-drawbacks-of-k-means>



# Which one is correct?




A





B


# Which one is correct?



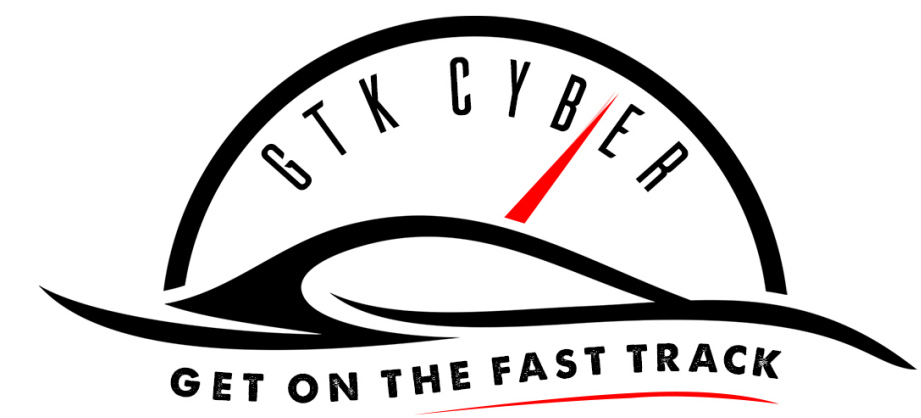
## BOTH



# Pain of optimization...Being stuck at sub-optimal local minimum...



Initial guess matters!  
Same outcome  
cannot be guaranteed




# K-Means in practice (Python version)

```
#Import from Scikit-learn  
from sklearn.cluster import KMeans
```

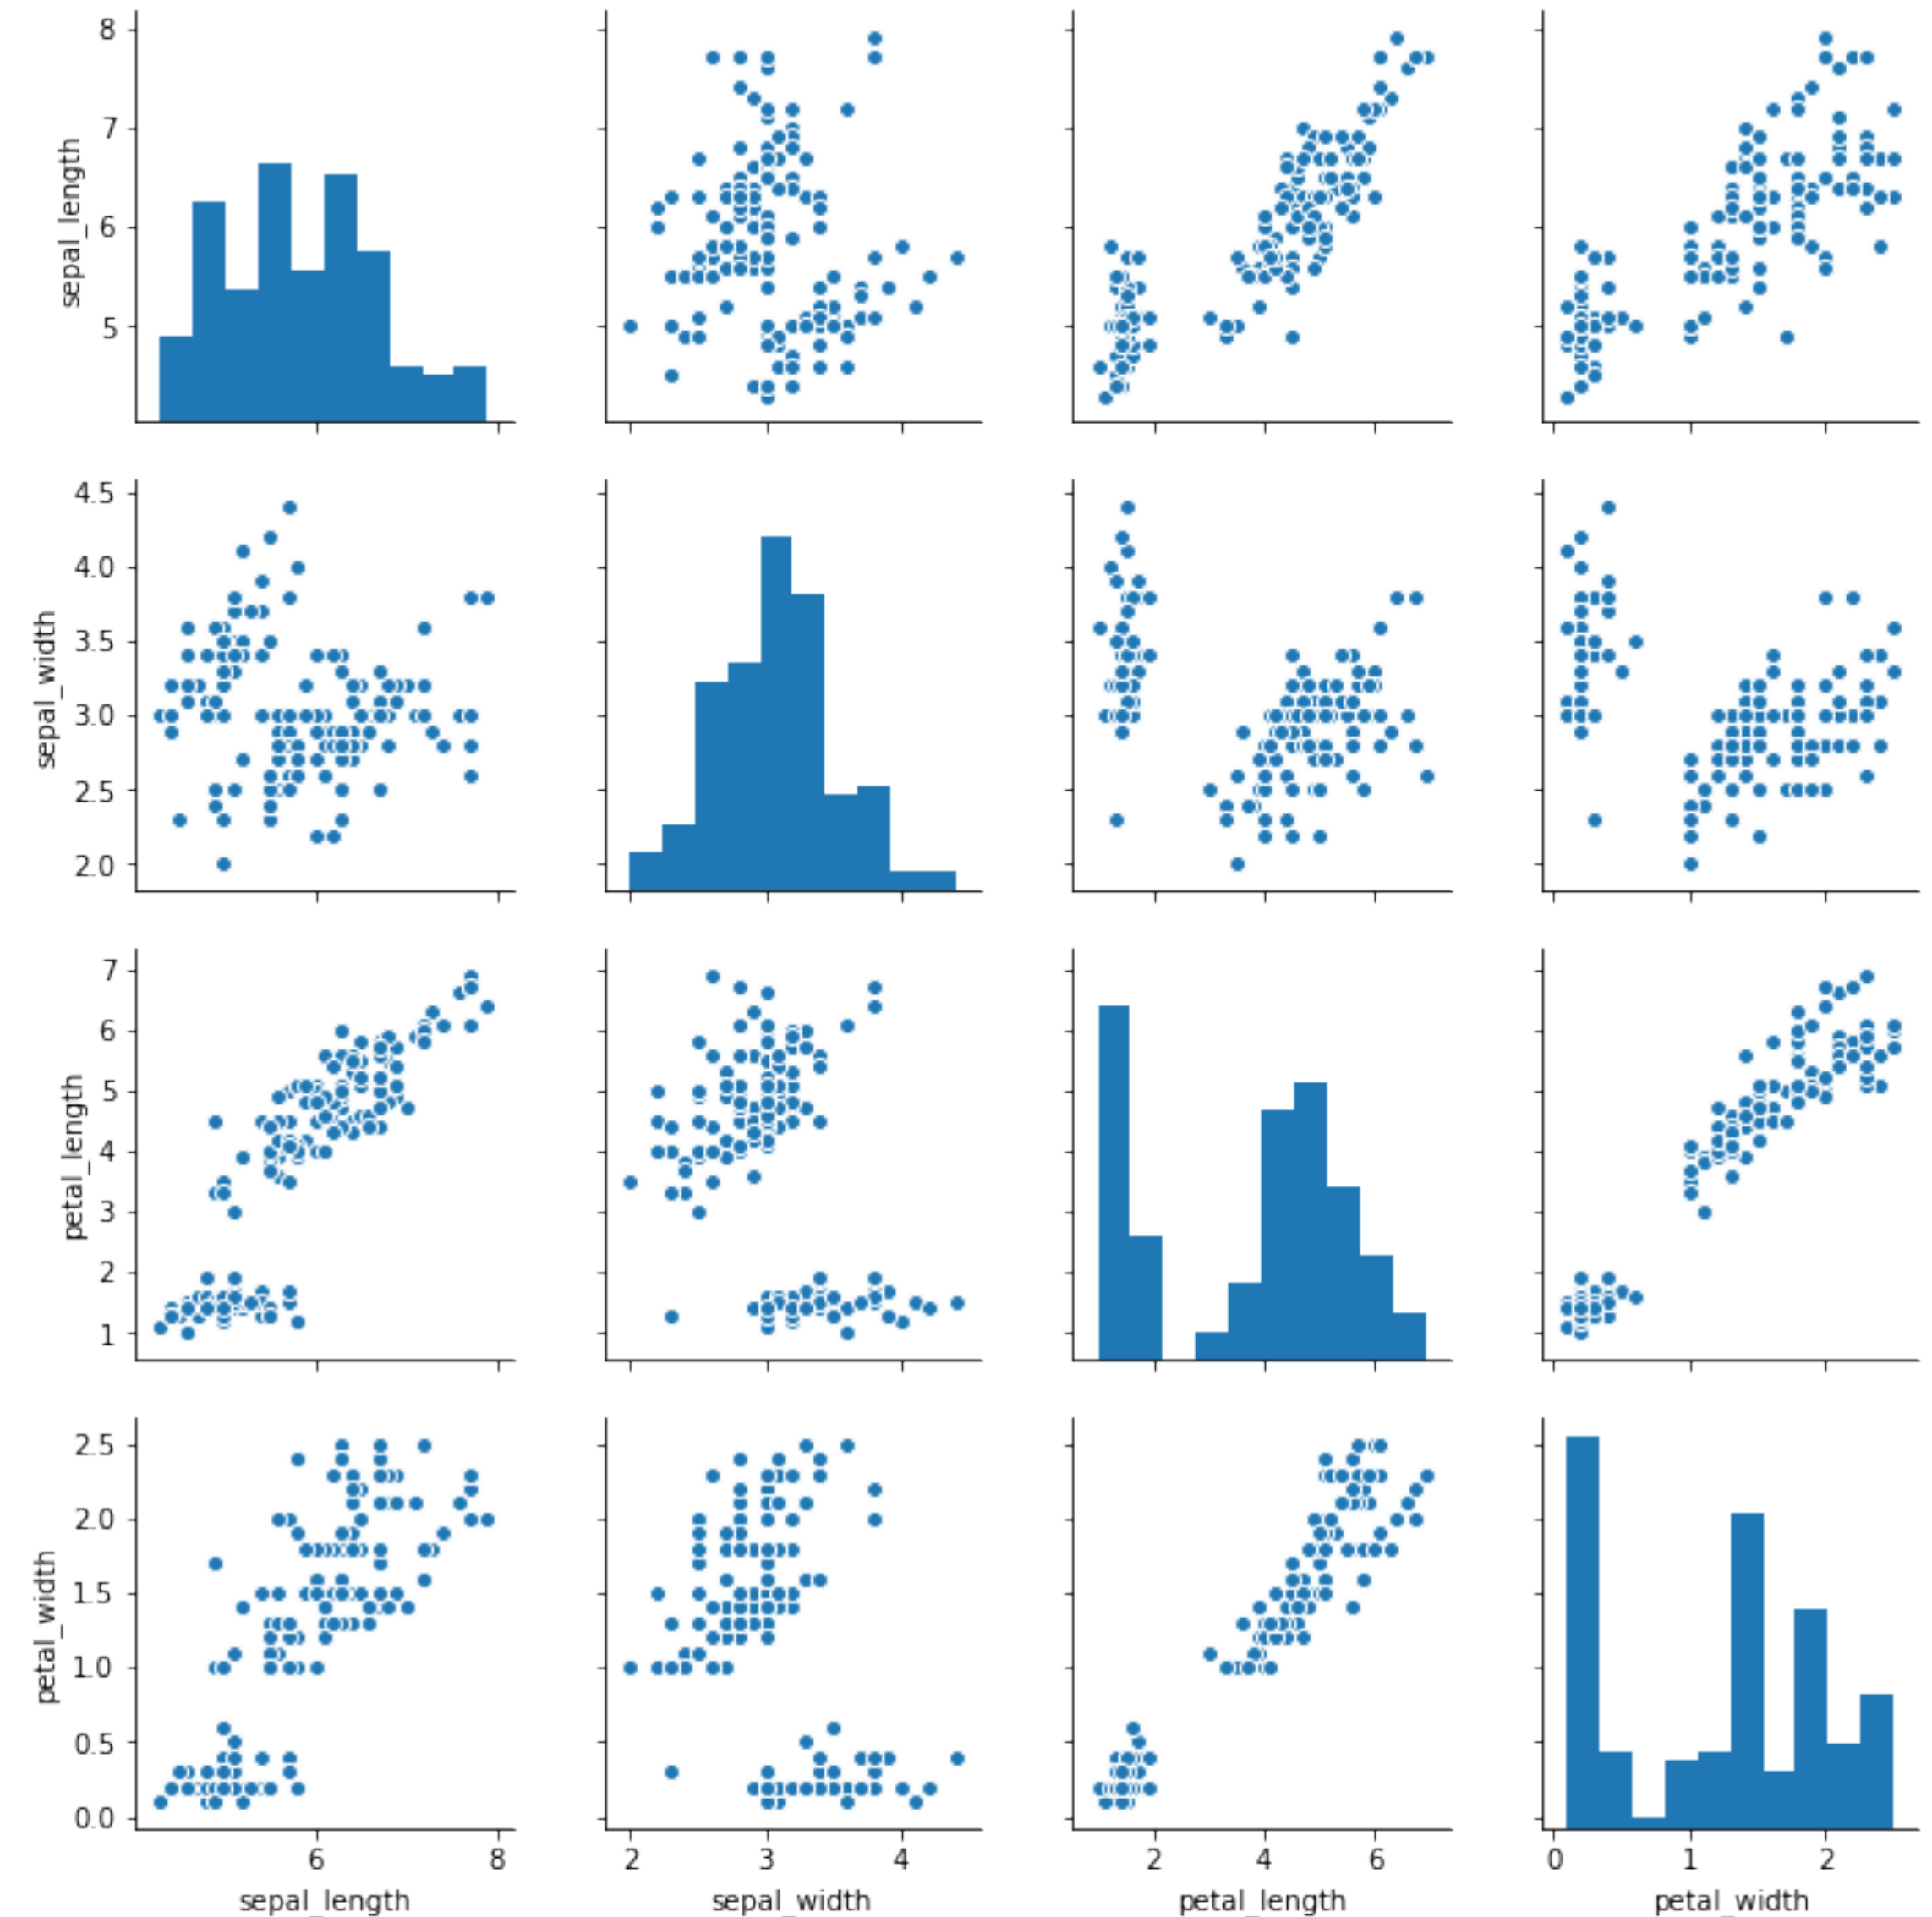
```
kmeans = KMeans(n_clusters=2)  
kmeans.fit(data)
```

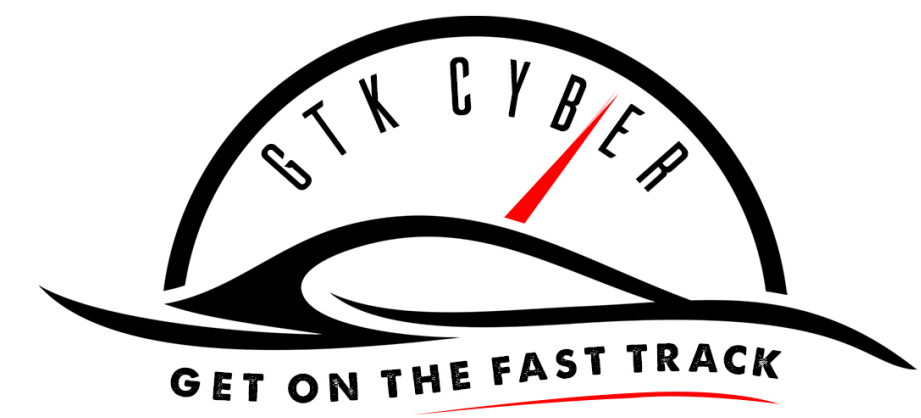
```
centroids = kmeans.cluster_centers_  
labels = kmeans.labels_f
```



# The Dataset

```
sns.pairplot(<data>)
```





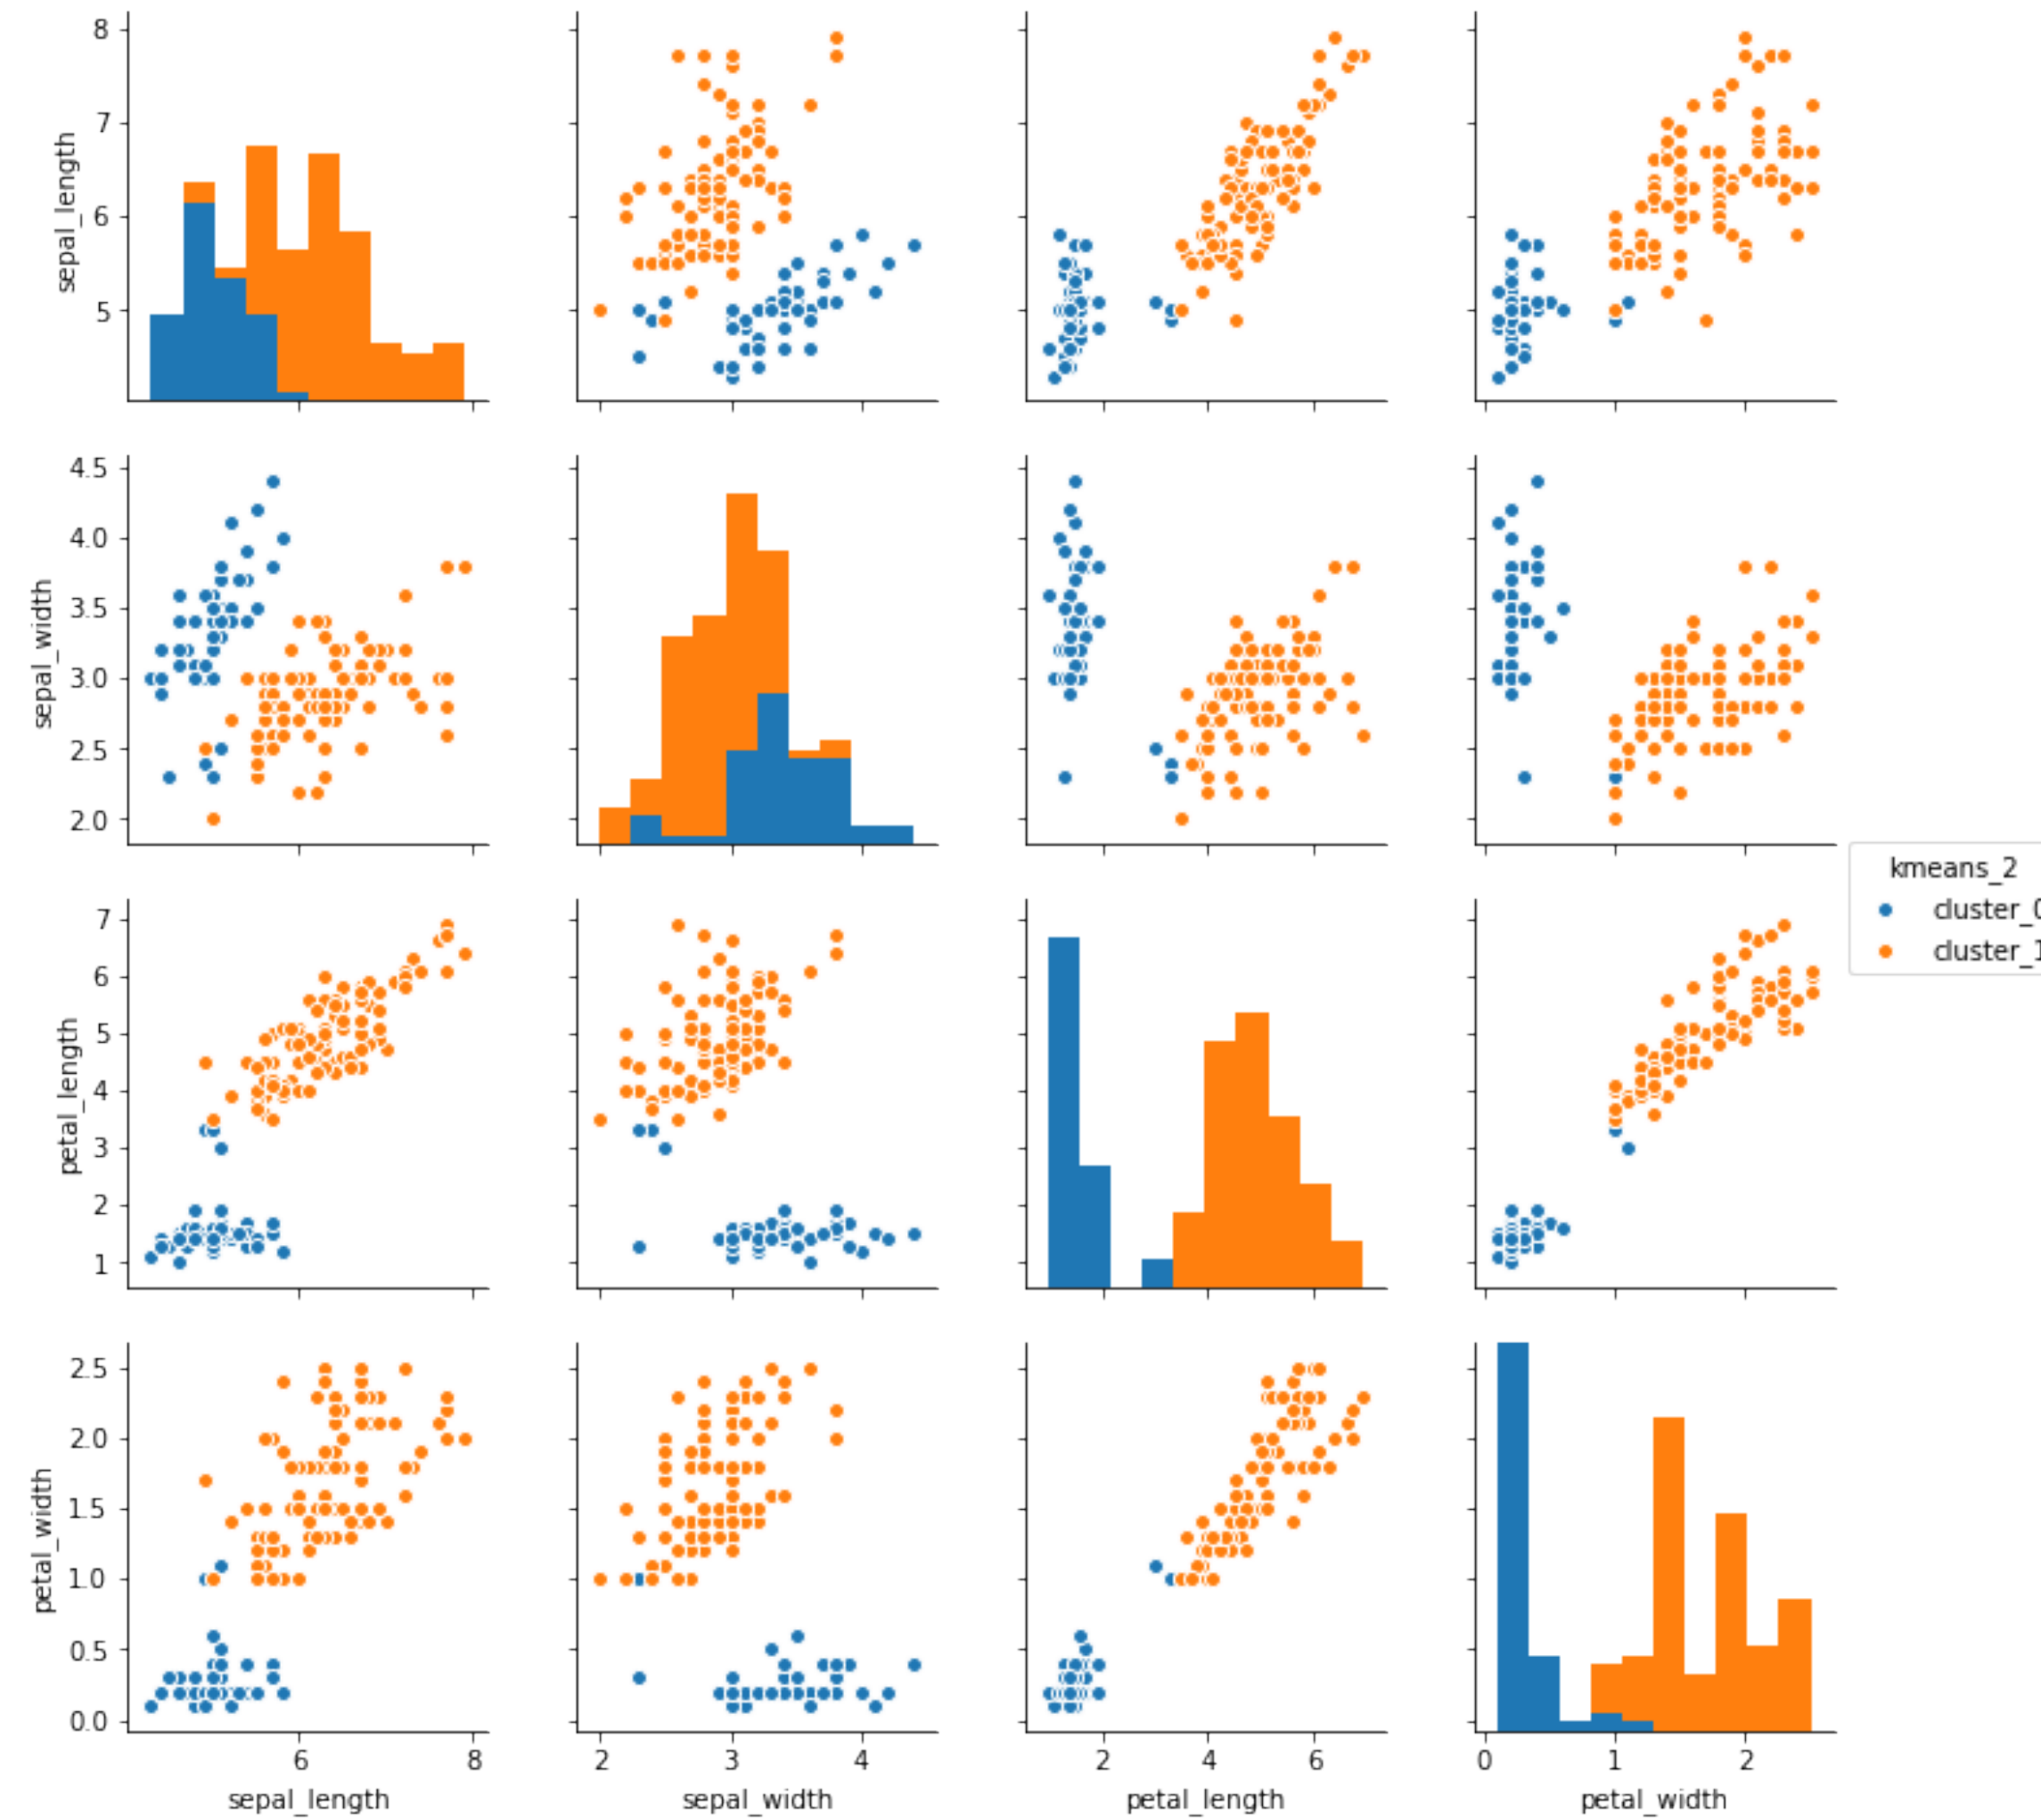
# K-Means Clustering

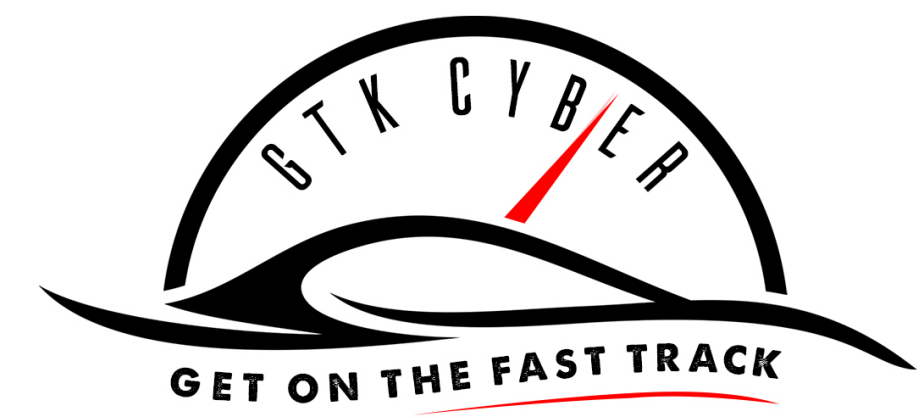
```
kmeans = KMeans( n_clusters=2 )  
kmeans.fit( <data> )
```



# K-Means Clustering


```
sns.pairplot(<data>, hue="kmeans_2")
```






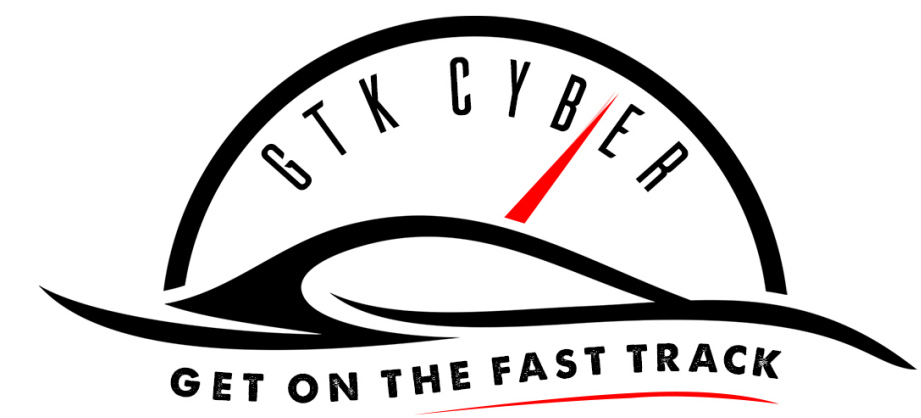
# K-Means Clustering

```
sns.pairplot(<data>,x_vars="col_1",y_vars="col_2",hue="kmeans_2",size=6)  
plt.scatter(<cluster_centers>,<col_2>, linewidths=3, marker='x', s=200,  
c='black' )
```





K-Means is affected by the scale of every feature.

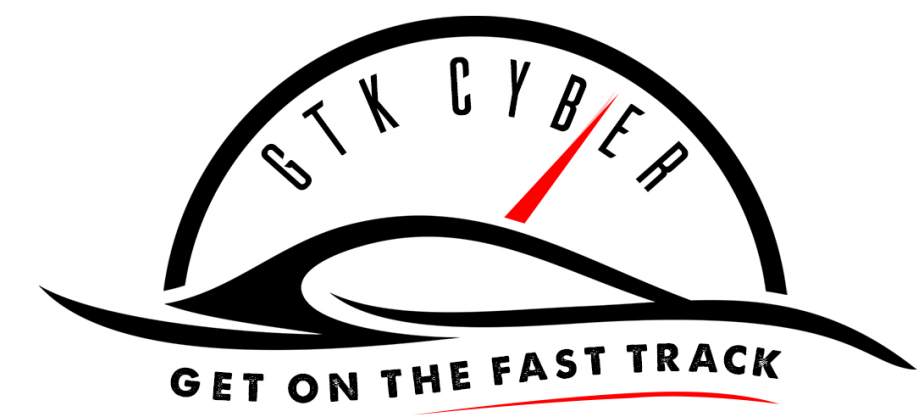


# Feature Scaling

For k-means clustering, features must be scaled to the same ranges of values to contribute "equally" to the euclidean distance calculation.

Each row is transformed per-column by:

- Subtracting from the element in each row the mean for each feature (column) and then taking this value and
- Dividing by that feature's (column's) standard deviation.

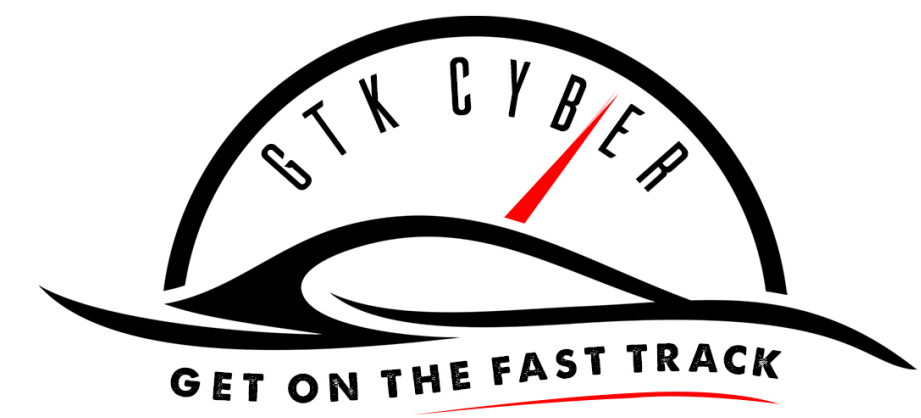


# Feature Scaling

```
# center and scale the data  
scaler = StandardScaler()
```

```
raw_data_scaled = scaler.fit_transform( <data> )
```

```
data_scaled = pd.DataFrame( raw_data_scaled, columns=features )
```



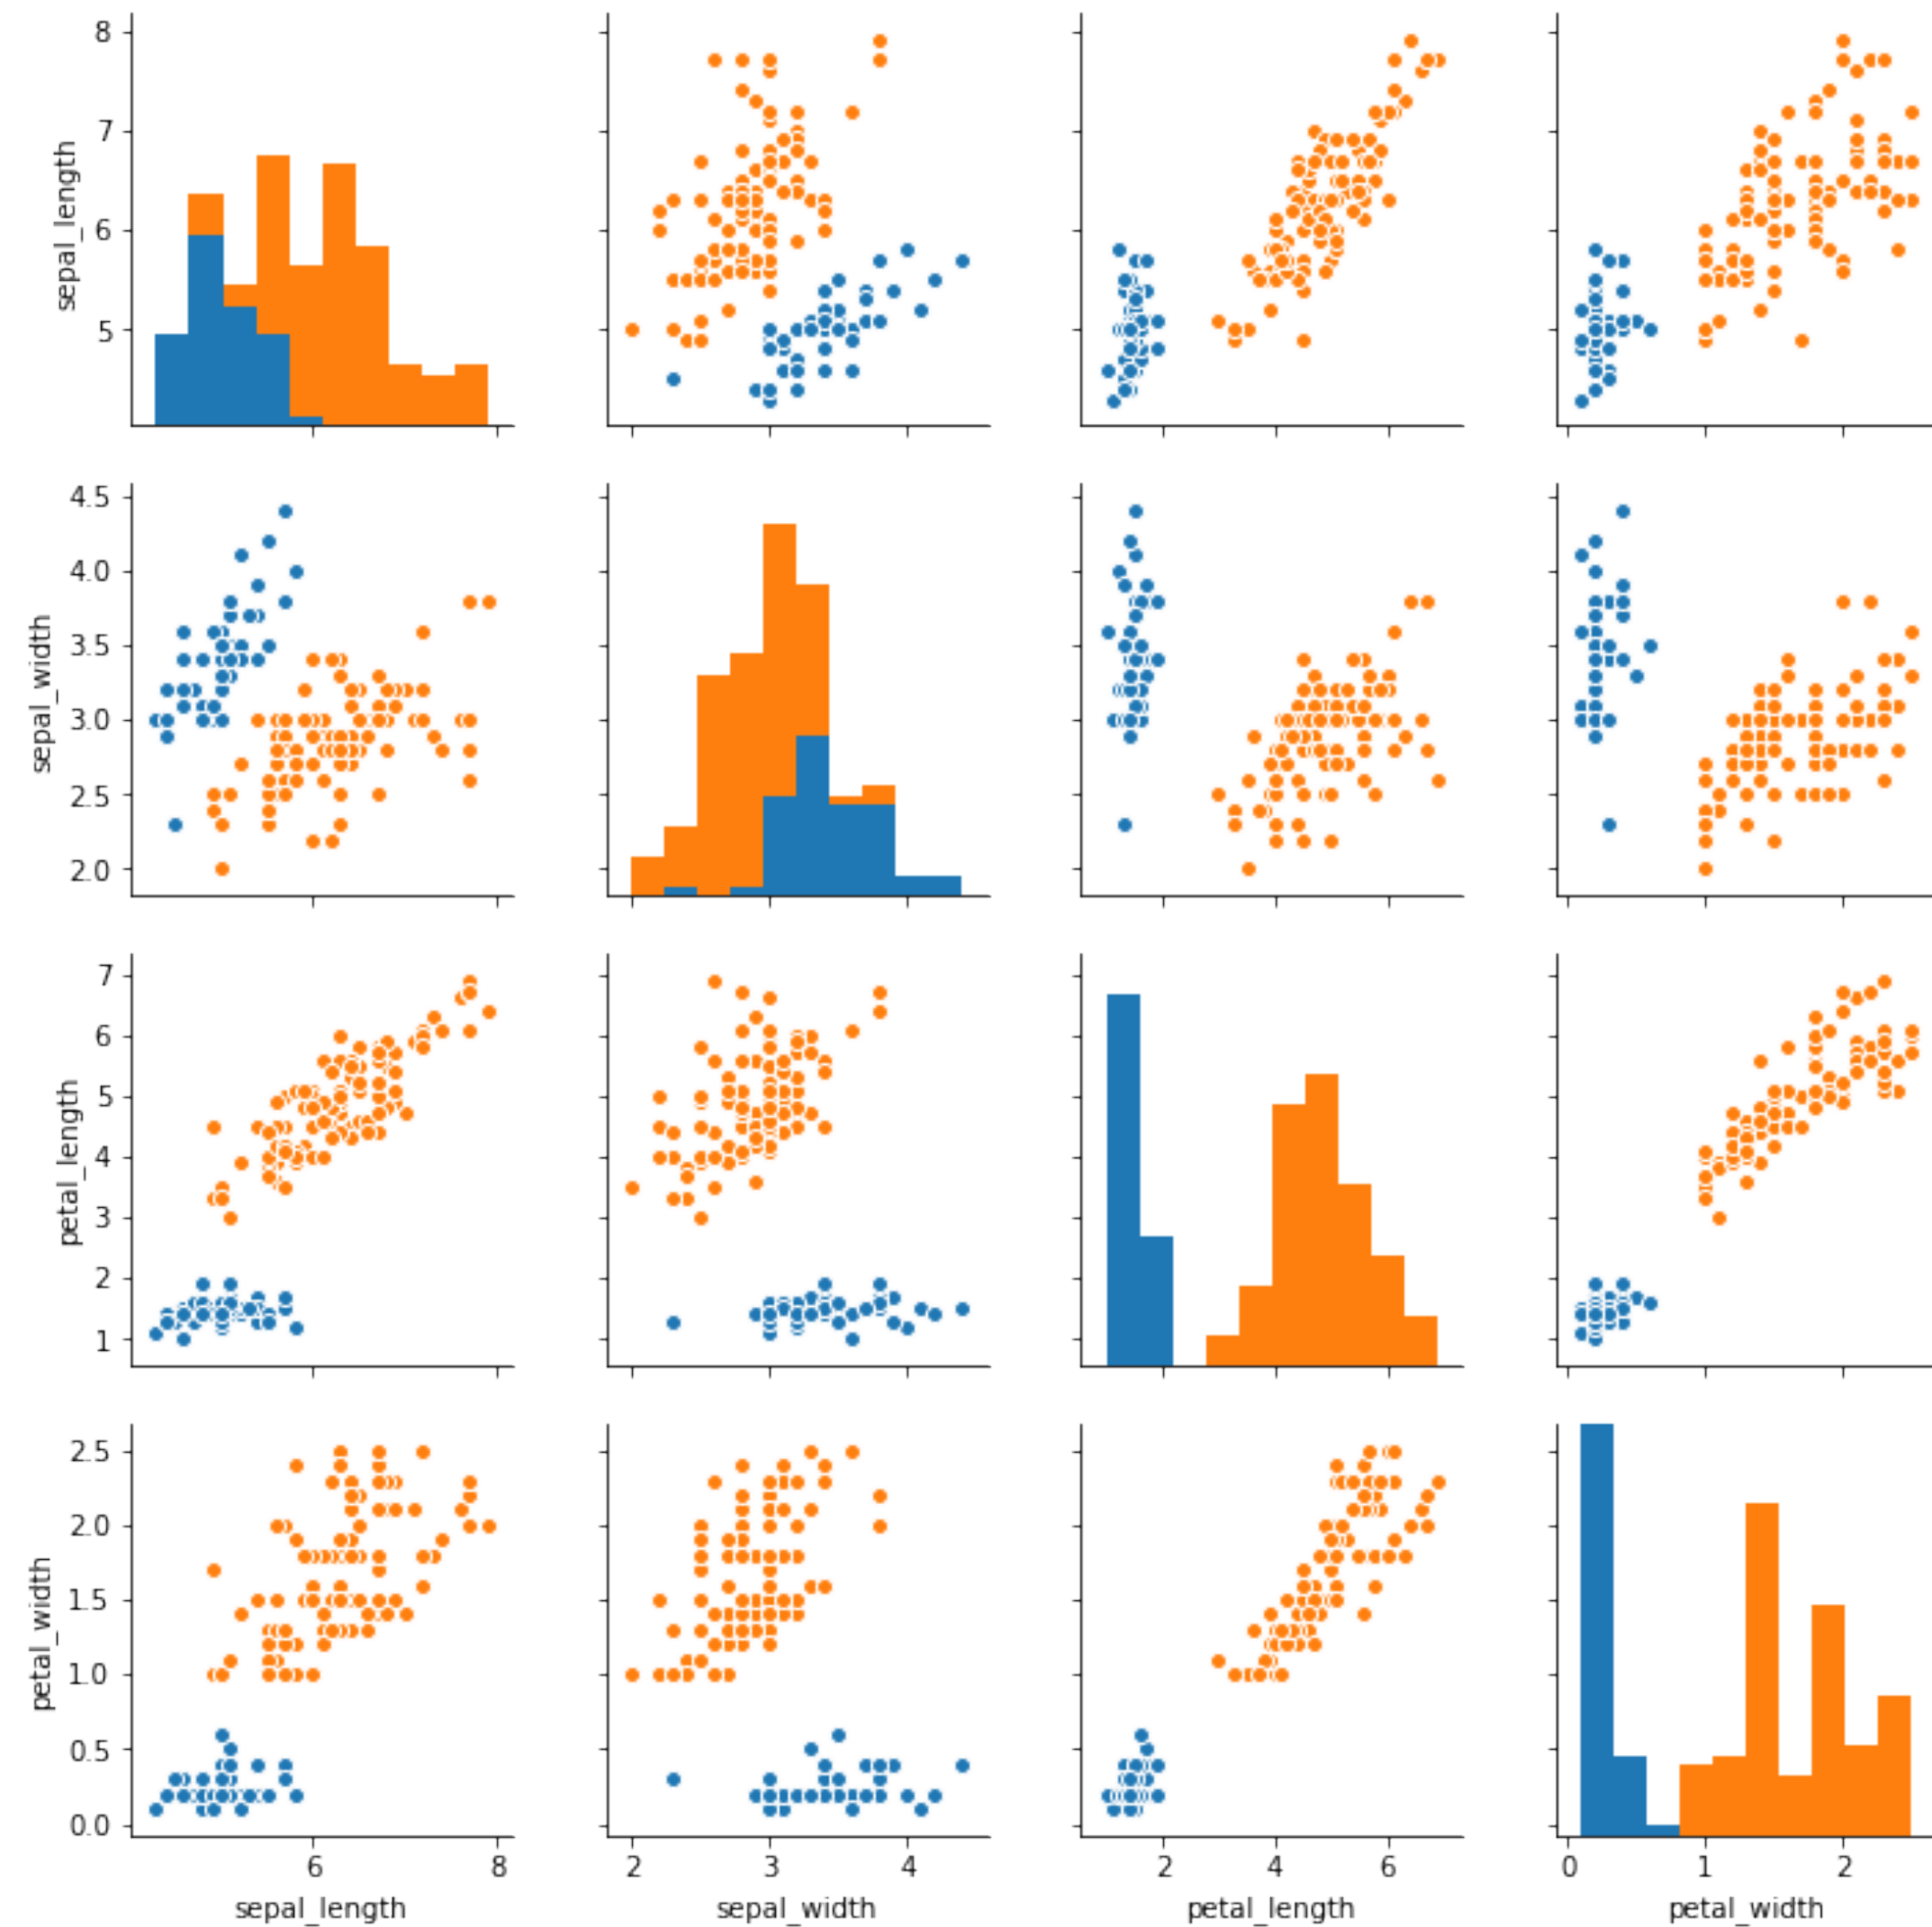
# Feature Scaling

```
# K-means on scaled data  
km = KMeans( n_clusters=2 )  
km.fit( <scaled_data> )
```

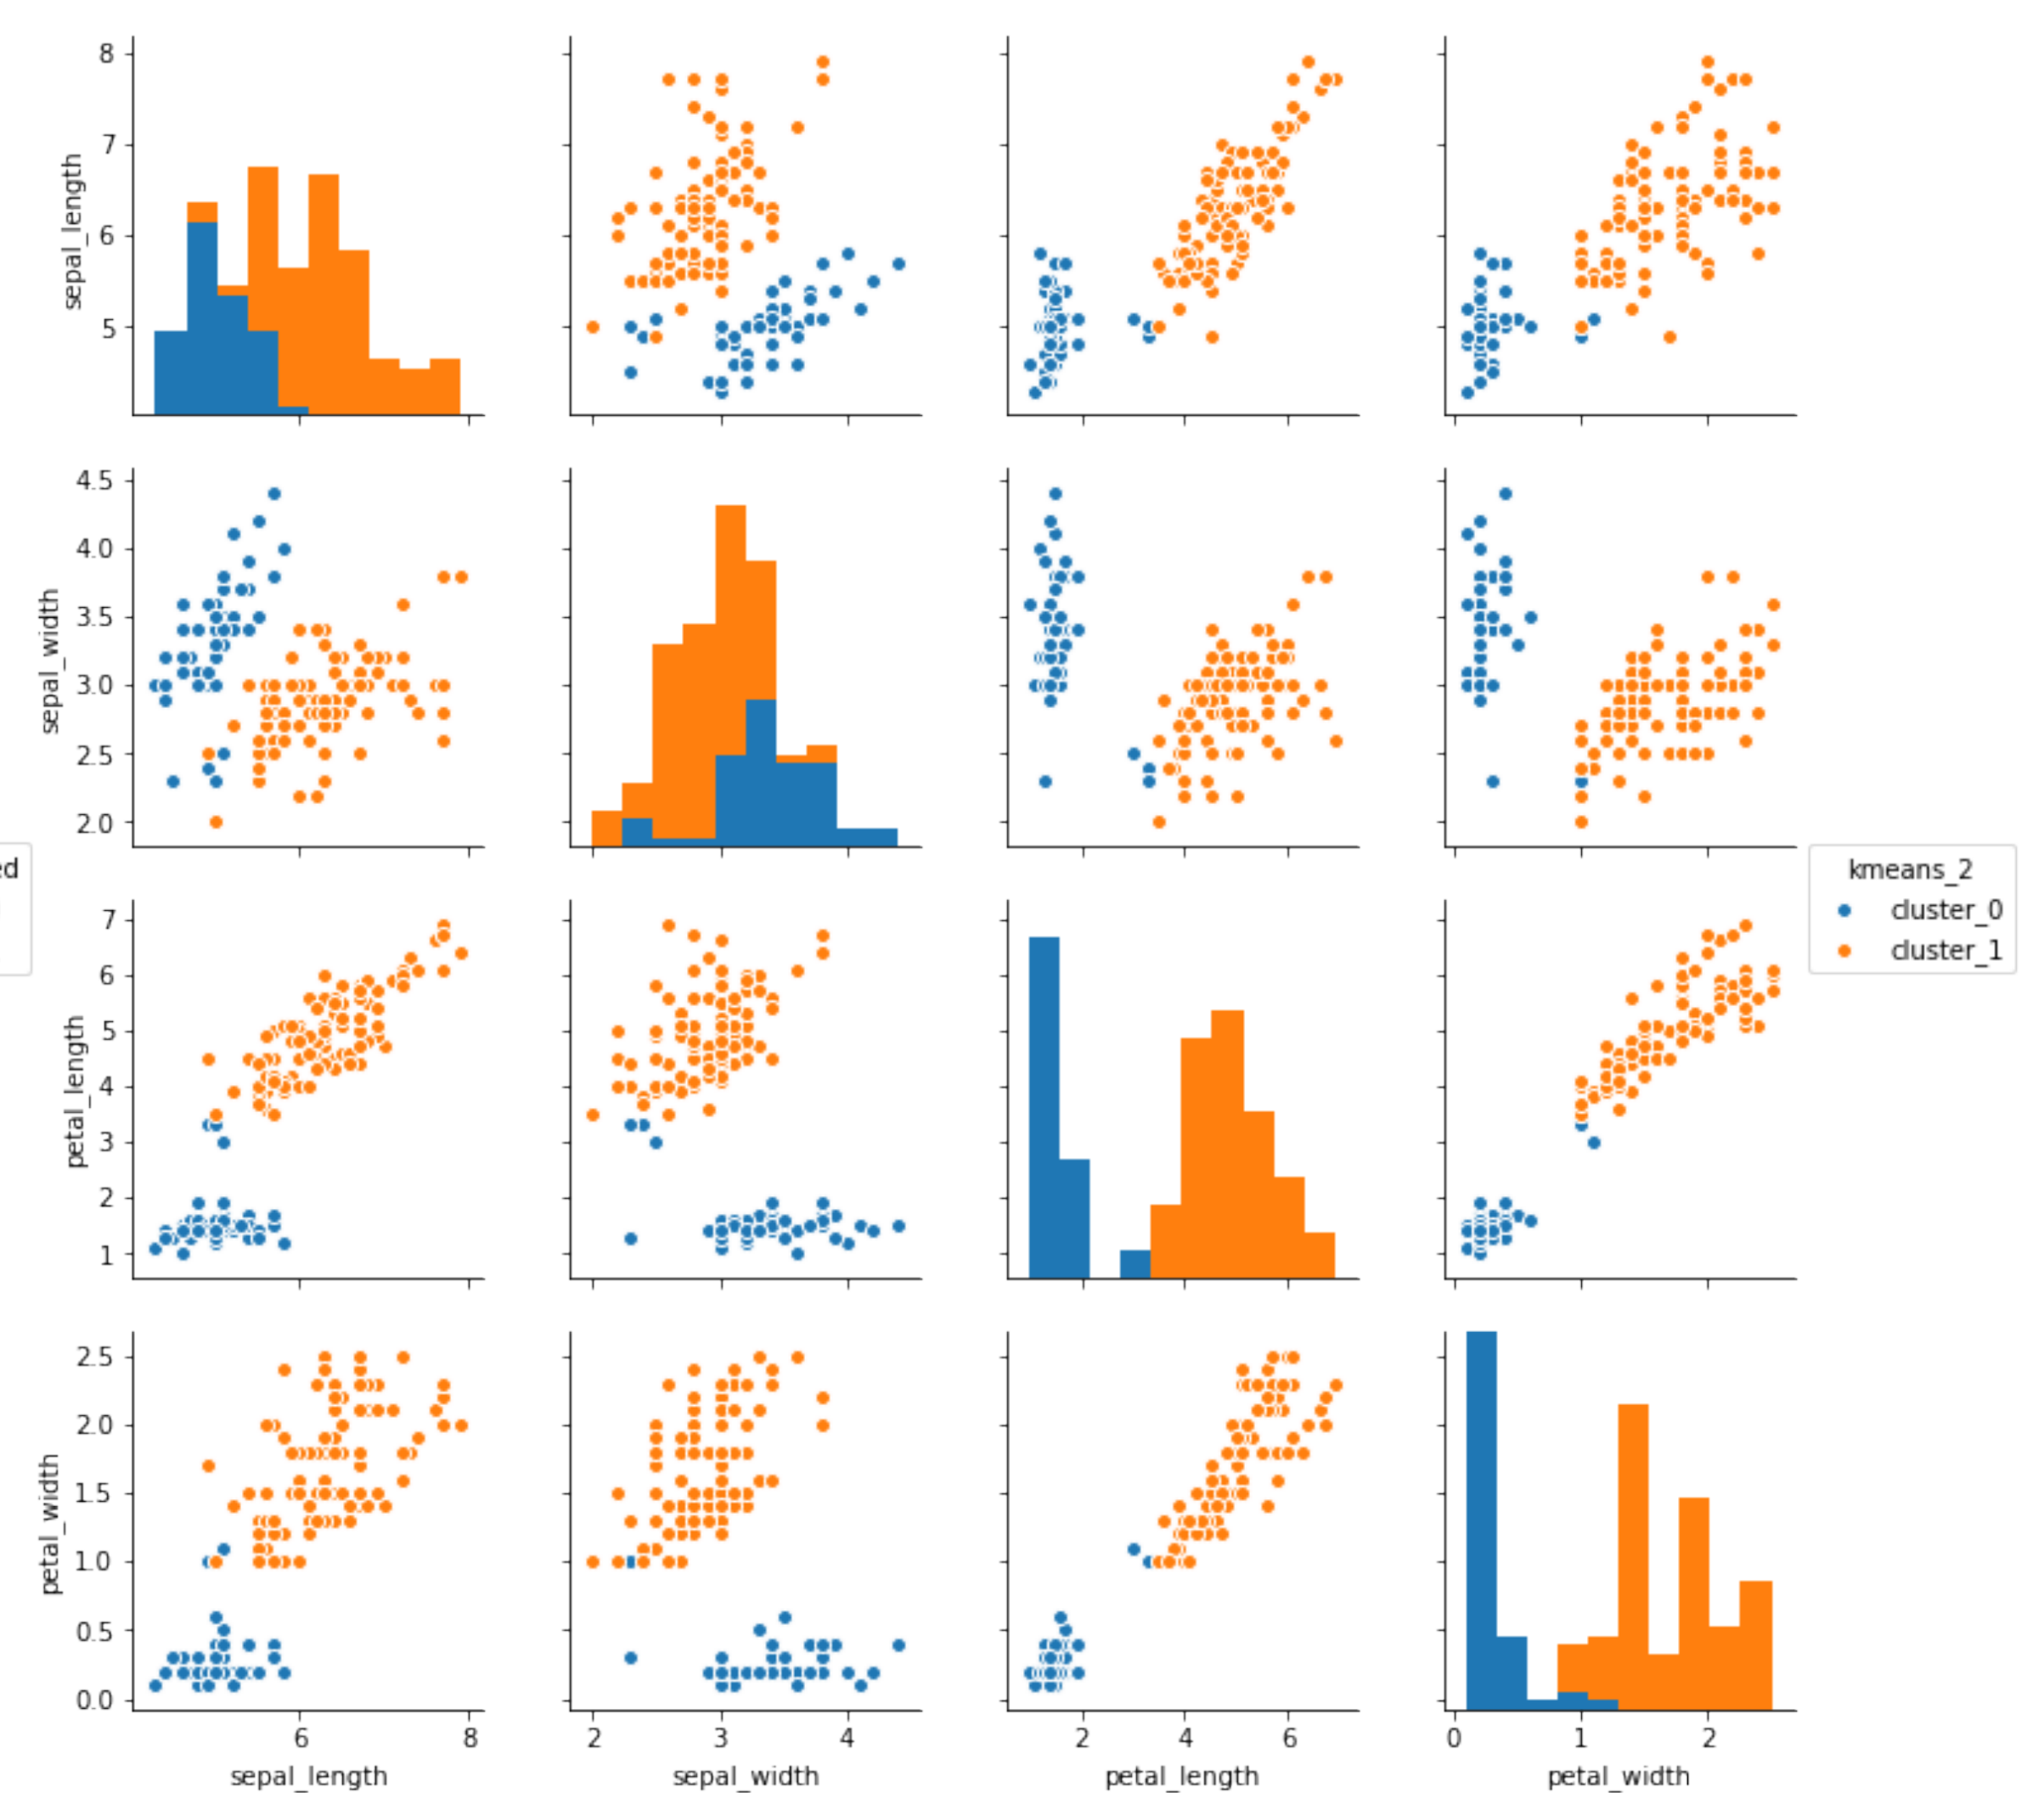


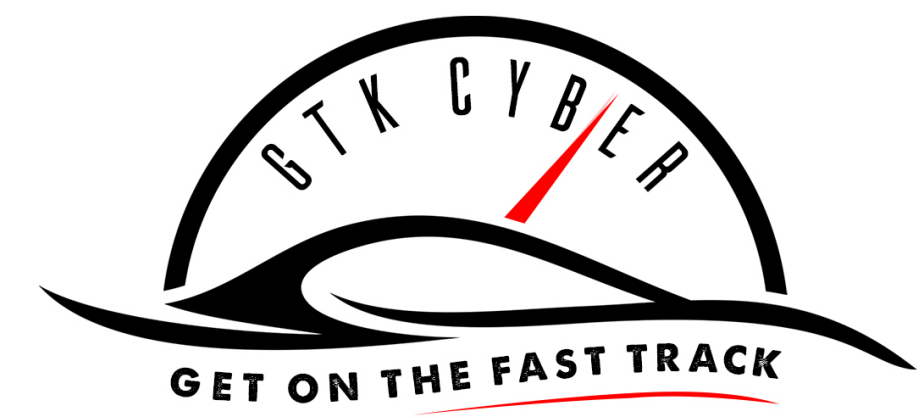
# Feature Scaling

Scaled Features




Unscaled Features





# More Clusters


```
km3 = KMeans(n_clusters=3)  
km3.fit(scaled_data)
```



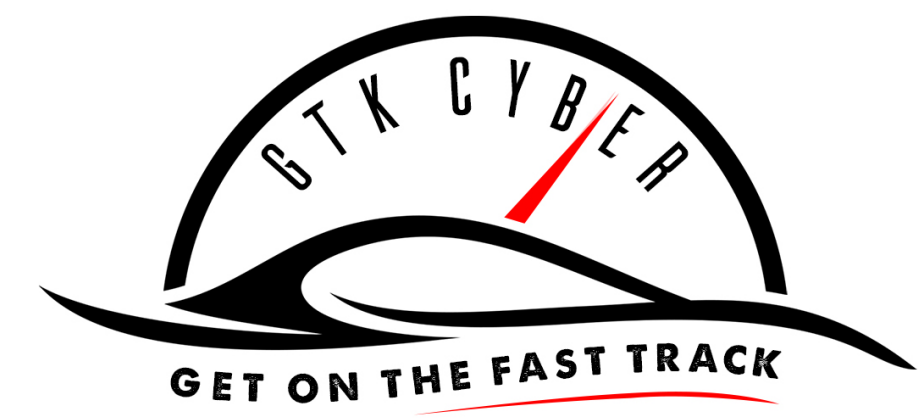
# Outlier Detection

```
clf = svm.OneClassSVM( tol=0.001, nu=0.1)
clf.fit(X)
target_pred_outliers=clf.predict(X)
```

Delete n% of “outlier data”,  
here ~10%







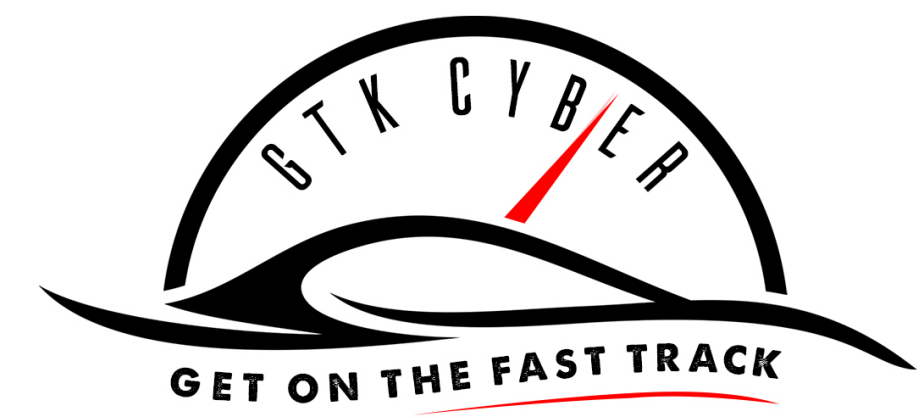
# Evaluating your Model

The Silhouette Coefficient is a common metric for evaluating clustering "performance" in situations when the "true" cluster assignments are not known.

b = mean distance to next nearest cluster

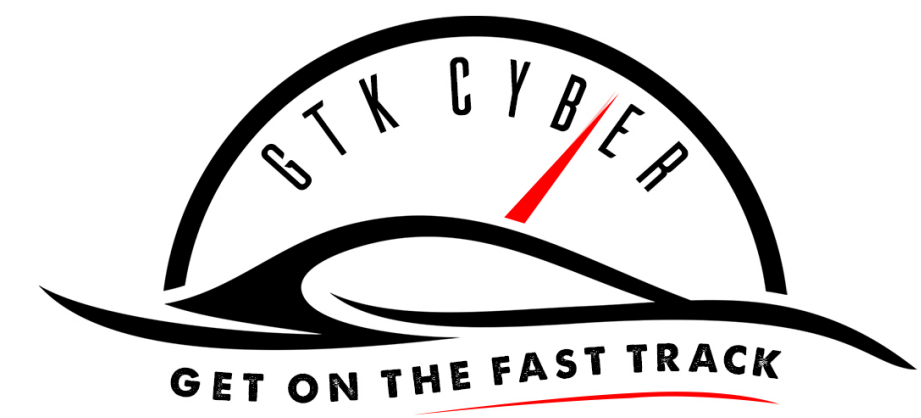
a = mean distance to other points in cluster

$$\text{silhouette\_coeff} = (b - a) / \max(a, b)$$




# Evaluating your Model

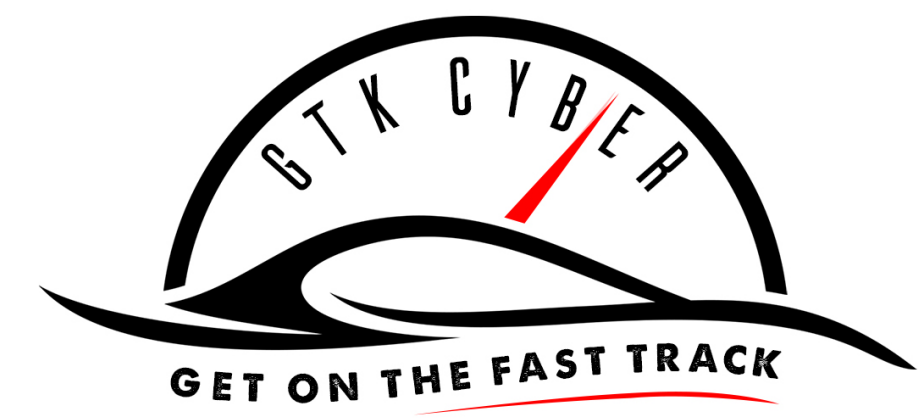
```
k_range = range(2,16)
scores = []
for k in k_range:
    km_ss = KMeans(n_clusters=k, random_state=1)
    km_ss.fit(iris_data_scaled)
    scores.append(silhouette_score(<data>,
km_ss.labels_))
```



# Evaluating your Model






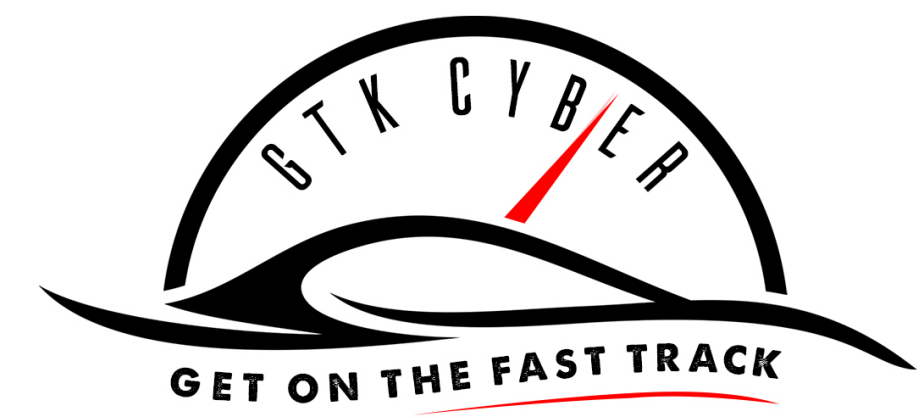


# Evaluating your Model

```
from yellowbrick.cluster import KElbowVisualizer  
visualizer = KElbowVisualizer(KMeans(), k=(4, 12))
```

```
visualizer.fit(X)  
visualizer.poof()
```

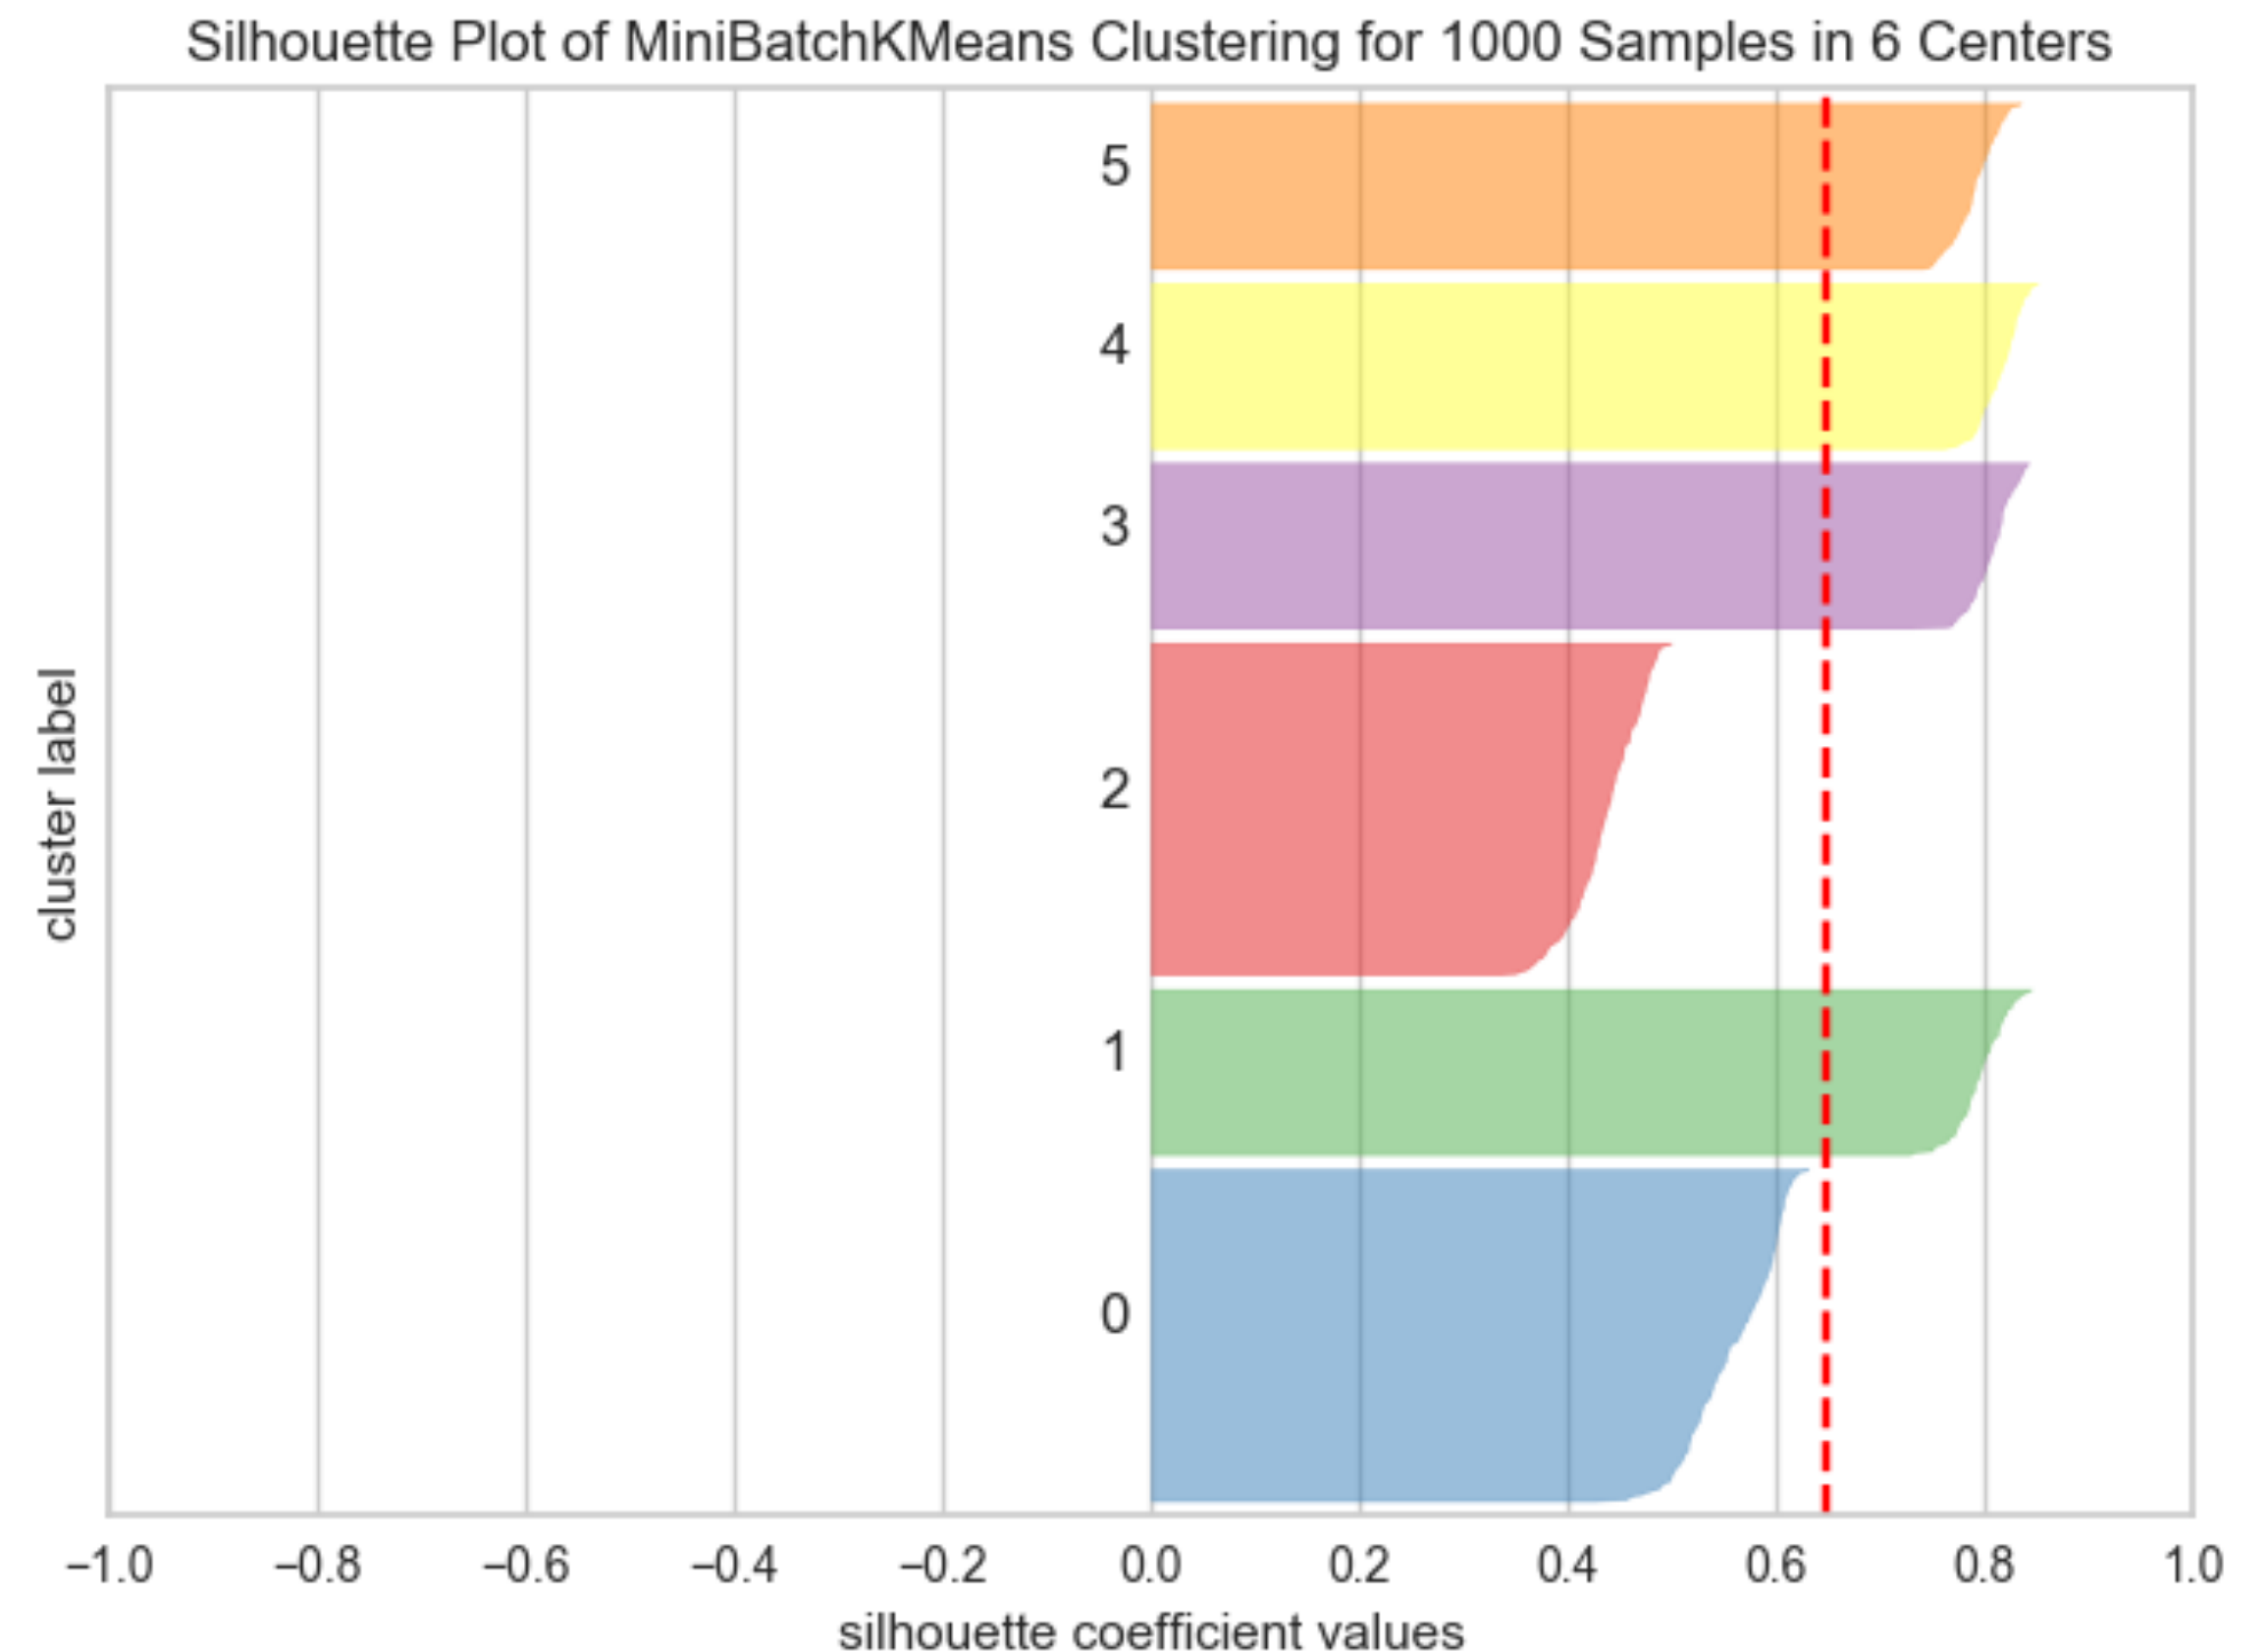


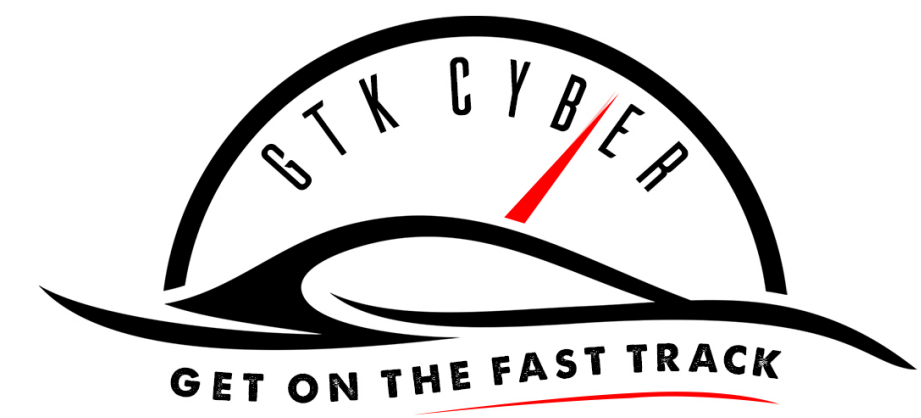


# Evaluating your Model

```
from yellowbrick.cluster import SilhouetteVisualizer  
model = MiniBatchKMeans(6)  
visualizer = SilhouetteVisualizer(model)
```

```
visualizer.fit(X)  
visualizer.poof()
```

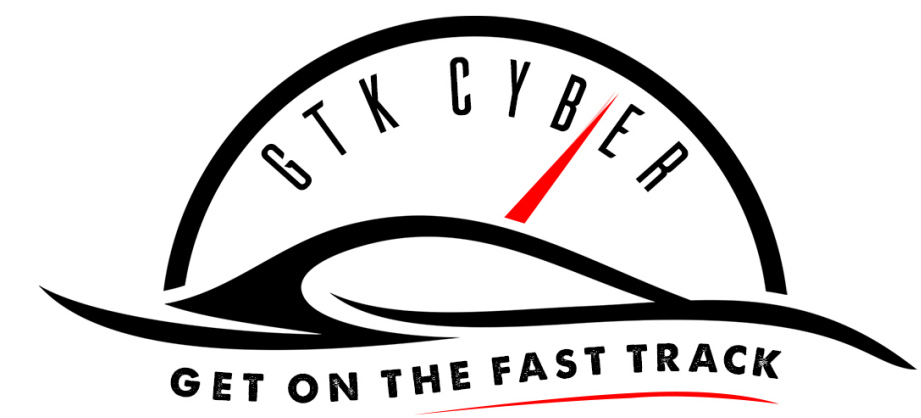




# DBSCAN

DBSCAN stands for **D**ensity-**B**ased **S**patial **C**lustering of **A**pplications with **N**oise.

Whereas K-means does not care about the density of data, DBSCAN does, under the assumption that regions of high density in your data should be treated as clusters.

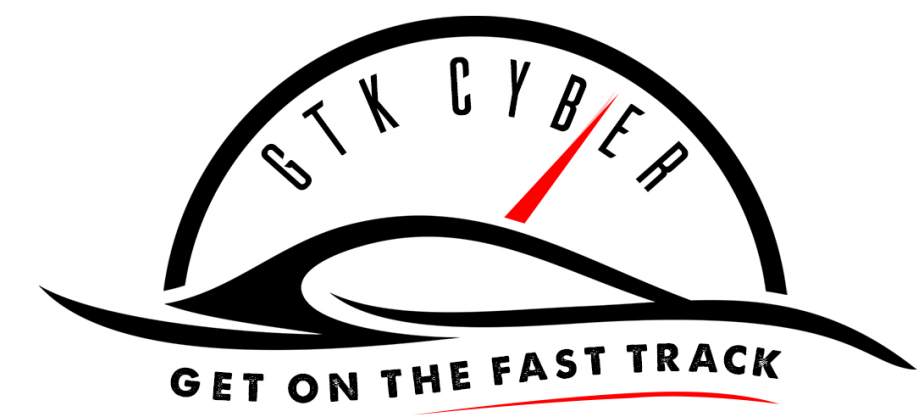


# DBSCAN

DBSCAN does not allow you to specify how many clusters you want. Instead, you specify 2 parameters:

- **e (epsilon)**: This is the maximum distance between two points to allow them to be neighbors
- **min\_samples**: The number of neighbors a given point is allowed to have to be able to be part of a cluster

Any points that don't satisfy the criteria of being close enough to other points are labeled outliers and all fall into a single "cluster" (their cluster label by default is -1).




# DBSCAN

DBSCAN works as follows:

1. Choose an arbitrary starting point in your dataset that has not been seen.
2. Retrieve this point's  $\epsilon$ -neighborhood (all points that are within a distance  $\epsilon$  from it), and if it contains at least **\*min\_samples**, a cluster is started.
3. Otherwise, the point is labeled as an outlier (-1). Note: This point might later be found in a sufficiently sized  $\epsilon$ -environment of a different point and hence be made part of a cluster.
4. If a point is found to be a dense part of a cluster, its  $\epsilon$ -neighborhood is also part of that cluster. All points that are found within the  $\epsilon$ -neighborhood are added, as is their own  $\epsilon$ -neighborhood when they are also dense.
5. Continue until the density-connected cluster is completely found.
6. Find a new unvisited point to process and repeat.



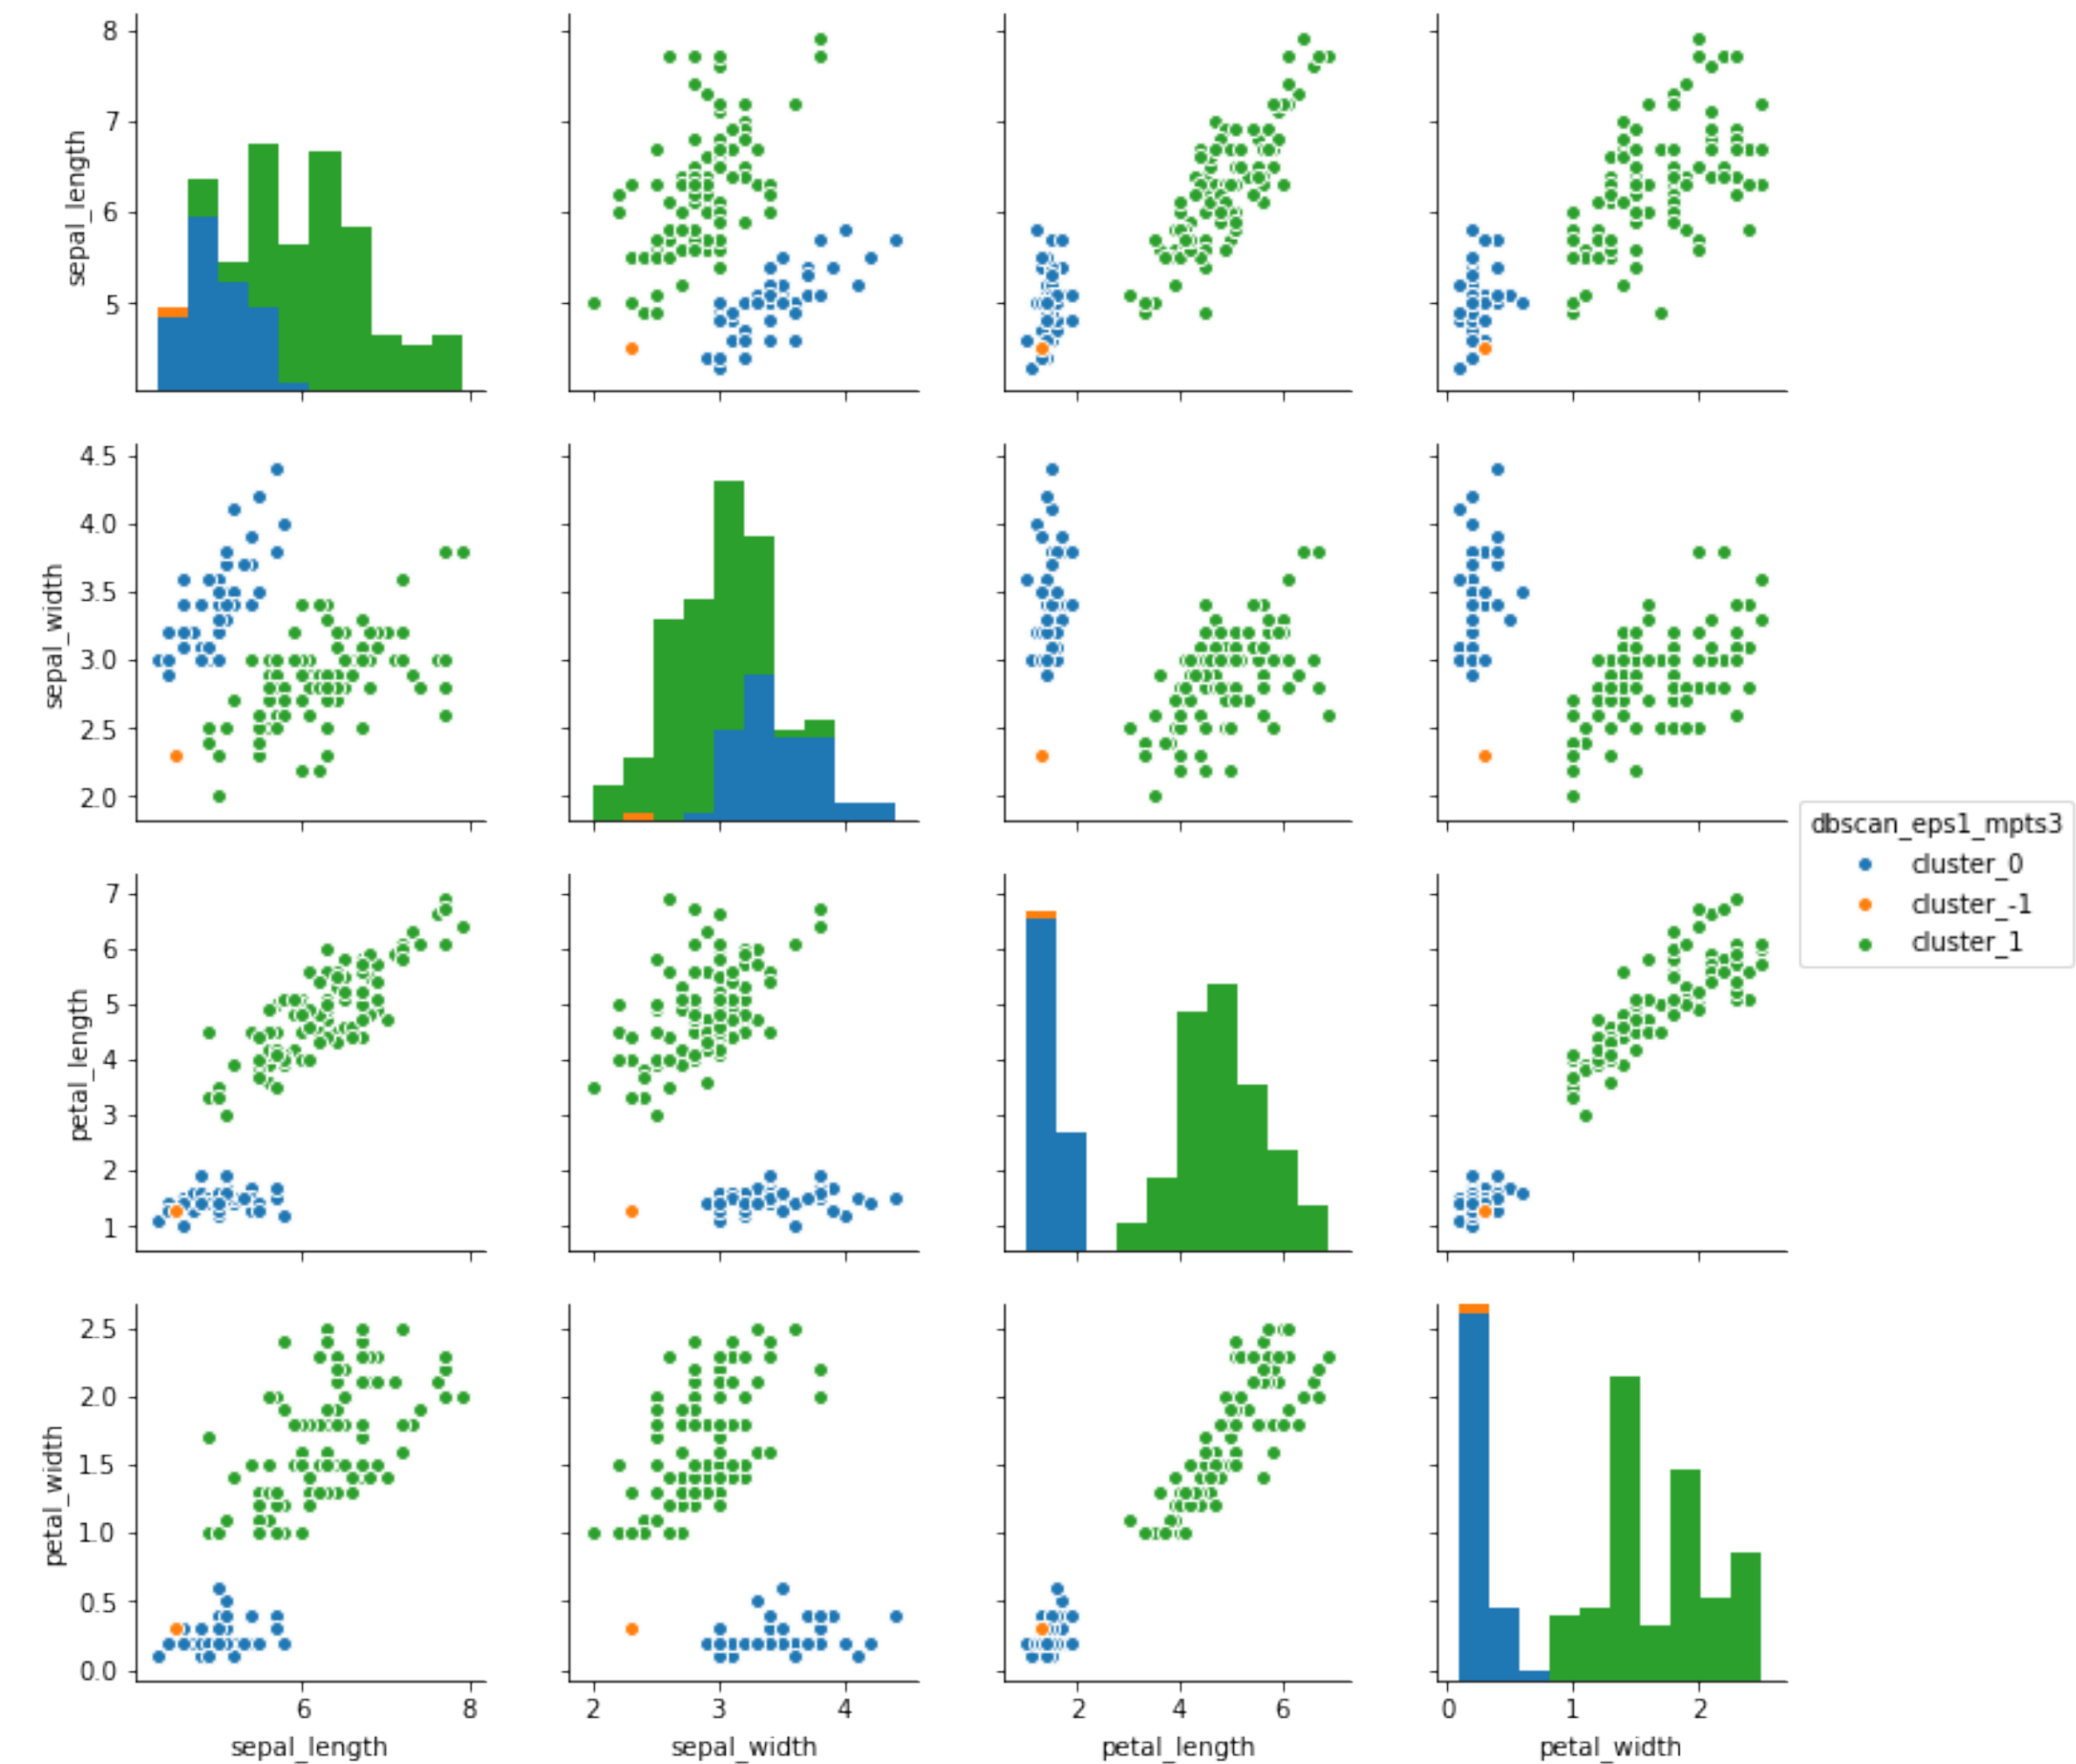


# DBSCAN

```
db = DBSCAN(eps=1, min_samples=3)
db.fit(<scaled_data>)
```

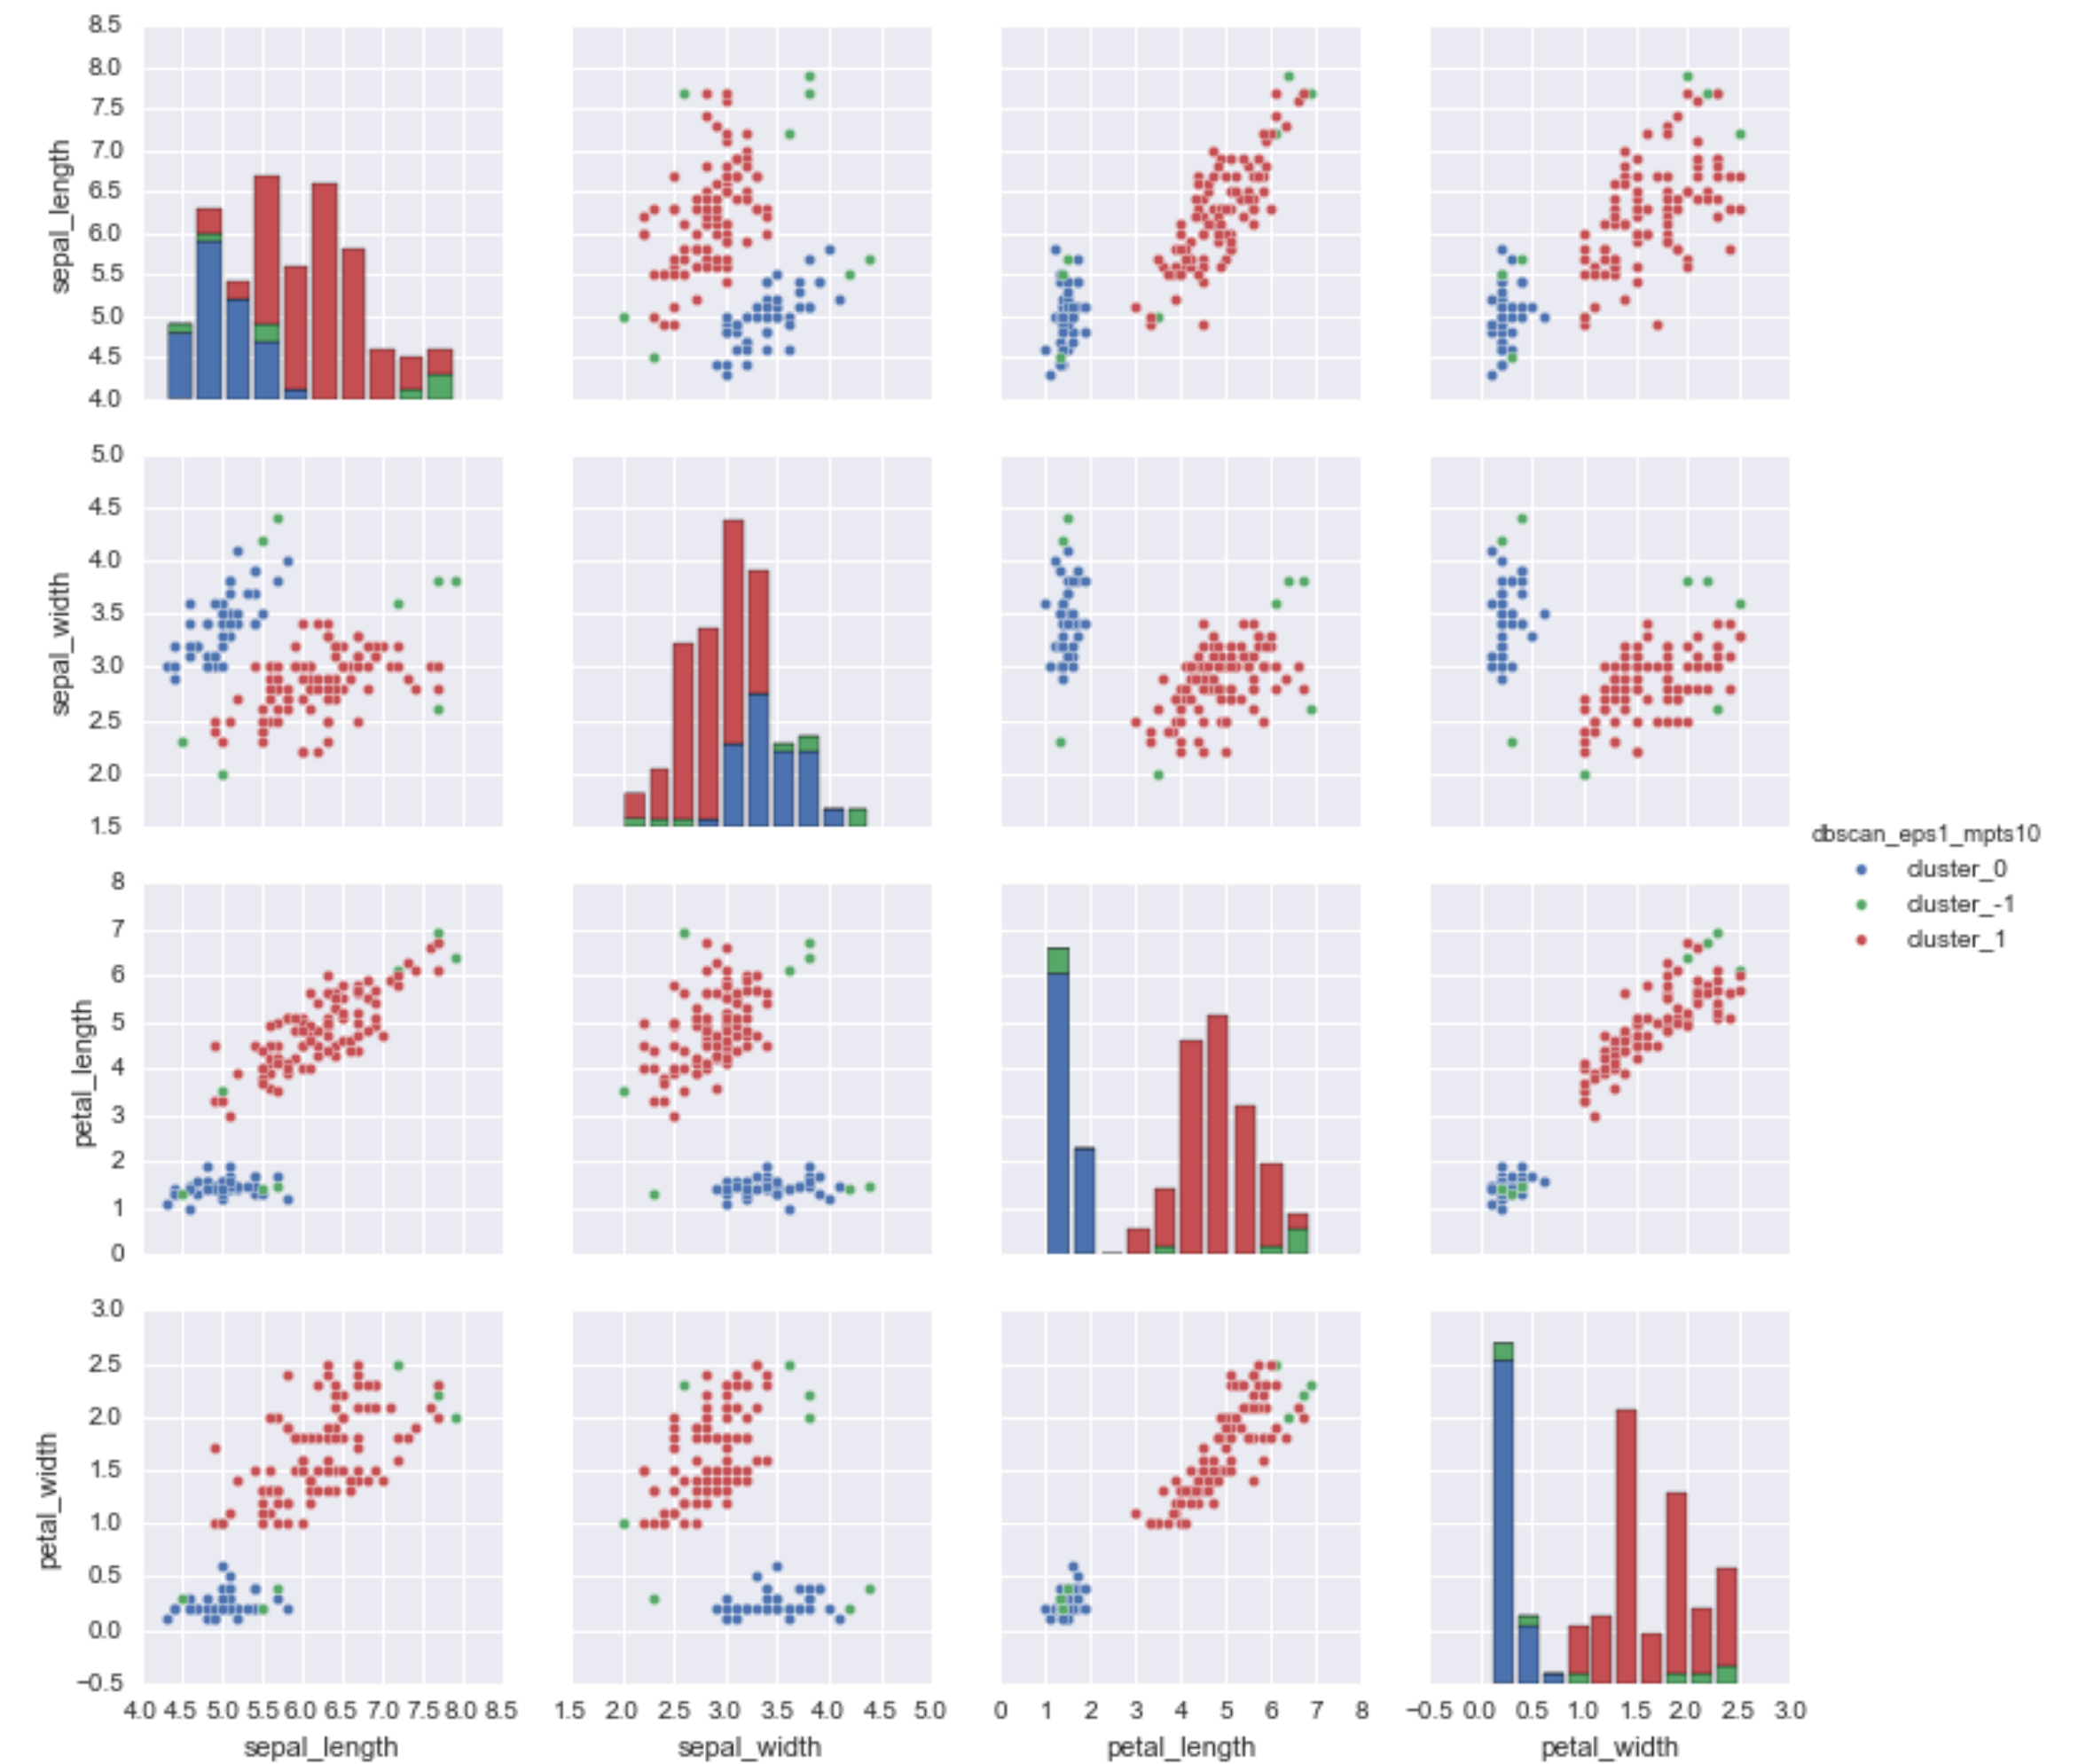
# DBSCAN

```
data_no_names['dbscan_eps1_mpts3'] = [ "cluster_" + str(label) for label in db.labels_ ]
sns.pairplot(data_no_names,hue="dbscan_eps1_mpts3")
```



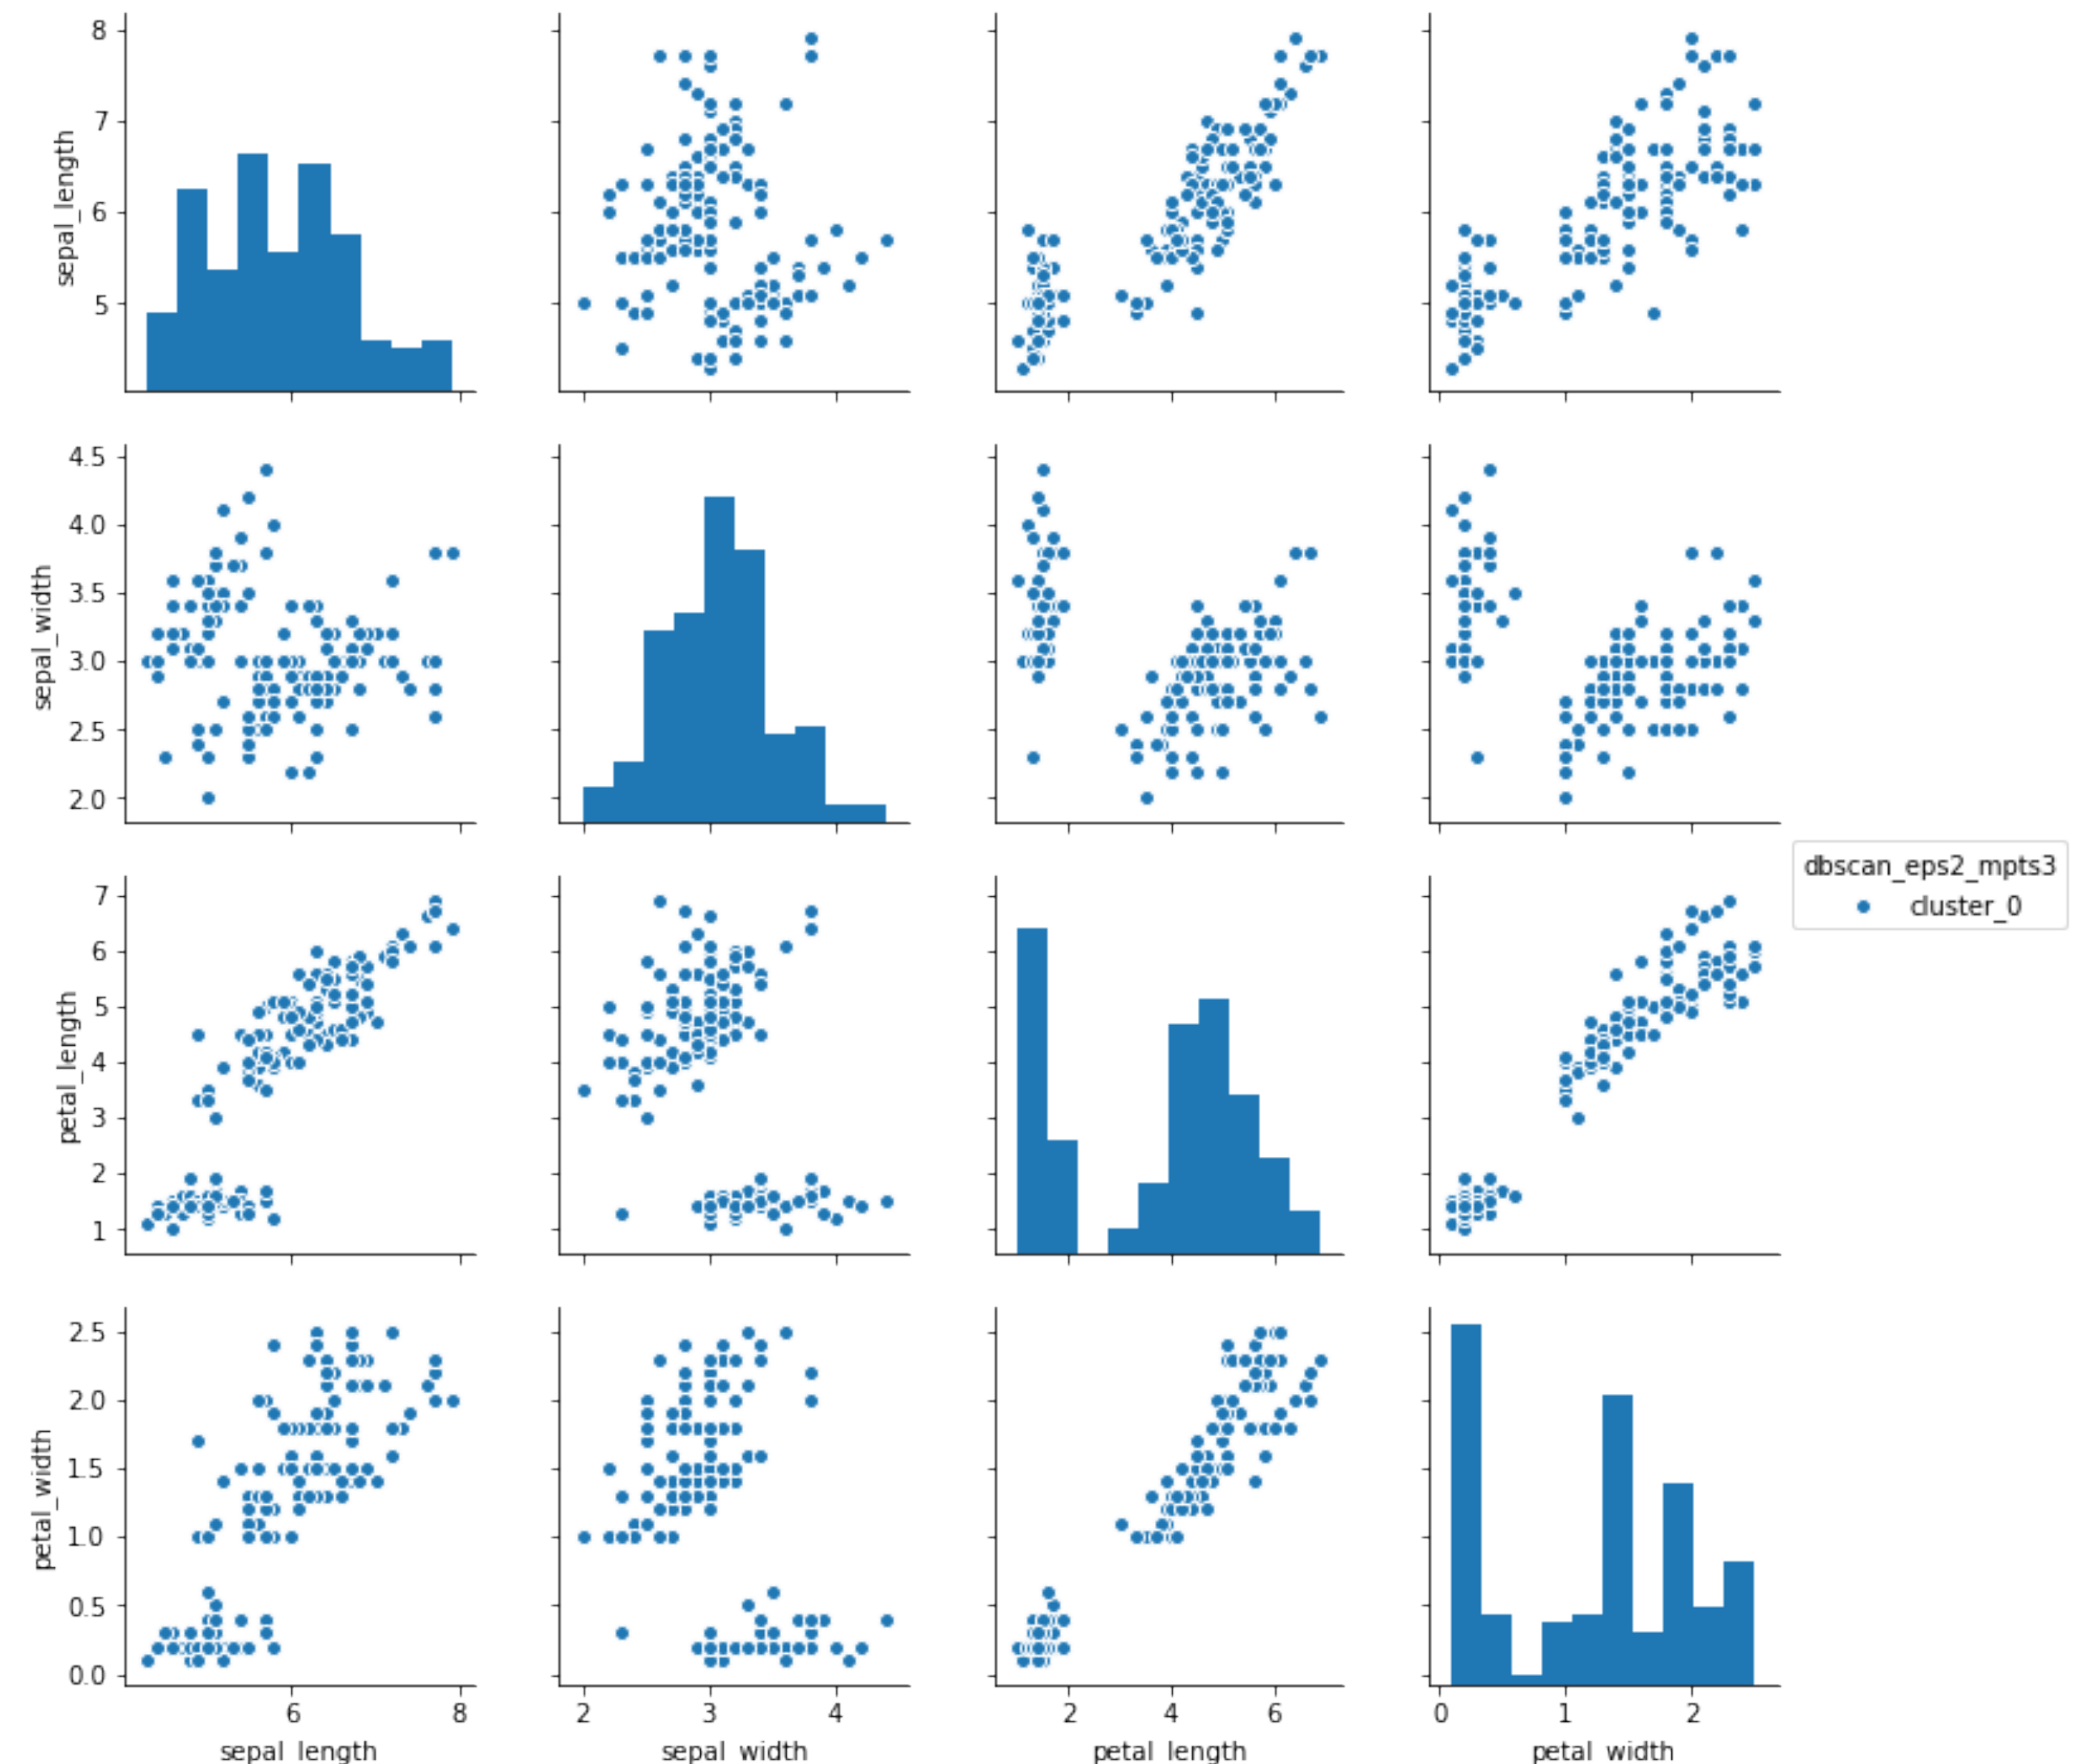
# DBSCAN

```
db2 = DBSCAN(eps=1, min_samples=10)
db2.fit(data_scaled)
```



# DBSCAN

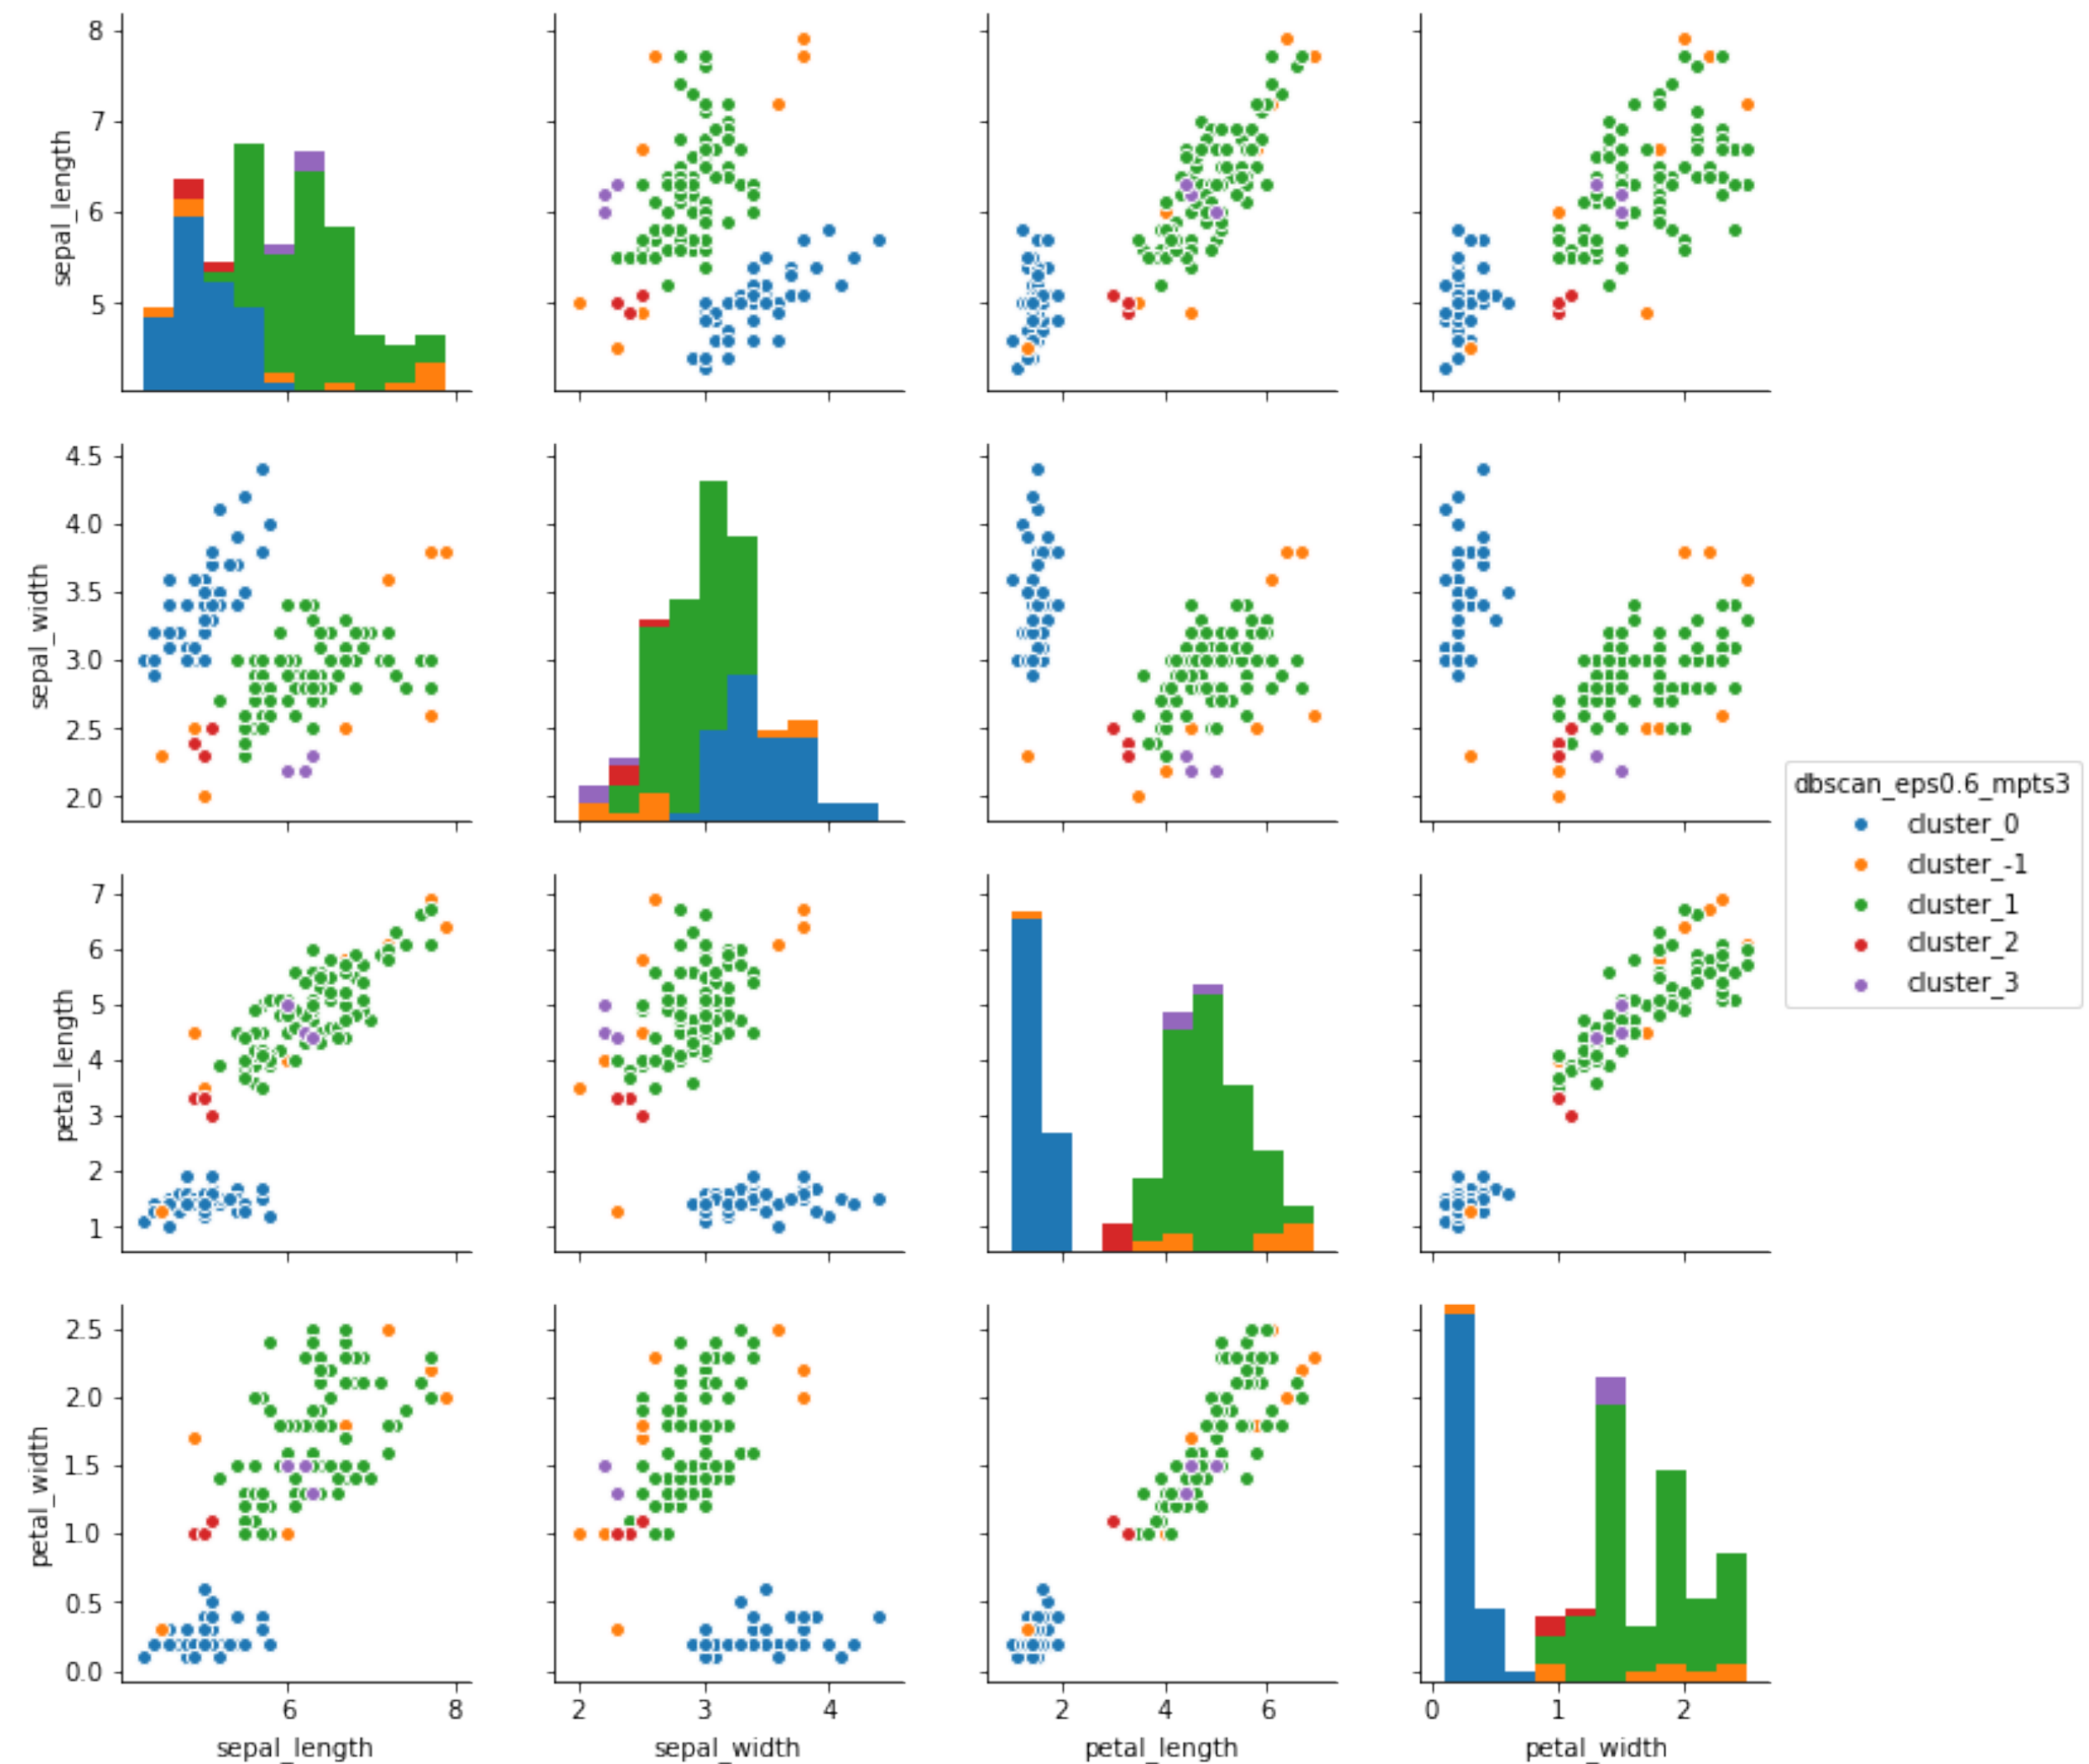
```
db2 = DBSCAN(eps=2, min_samples=3)  
db2.fit(iris_data_scaled)
```

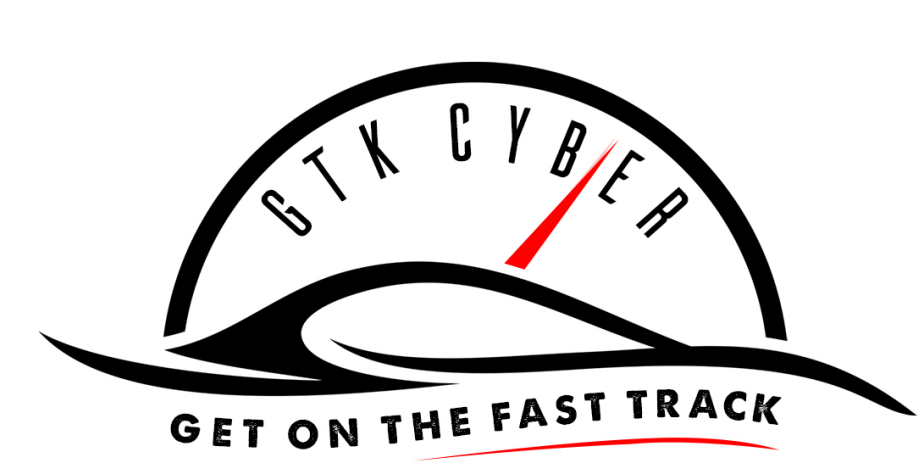




# DBSCAN

```
db2 = DBSCAN(eps=0.6, min_samples=3)  
db2.fit(iris_data_scaled)
```







# In Class Exercise

Please take 30 minutes and complete  
**Day 3: Clustering Worksheet**

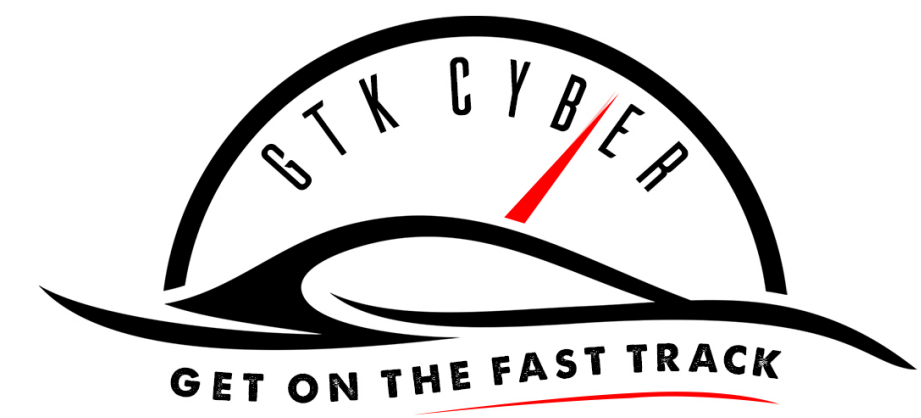




# Questions?



# Tuning Hyperparameters



# Grid Search

```
RandomForestClassifier(bootstrap=True,
```

```
class_weight=None,
```

```
criterion='gini',
```

```
max_depth=None,
```

```
max_features='auto',
```

```
max_leaf_nodes=None,
```

```
min_impurity_decrease=0.0,
```

```
min_impurity_split=None,
```

```
min_samples_leaf=1,
```

```
min_samples_split=2,
```

```
min_weight_fraction_leaf=0.0,
```


```
n_estimators=10,
```

```
n_jobs=1,
```

```
oob_score=False
```

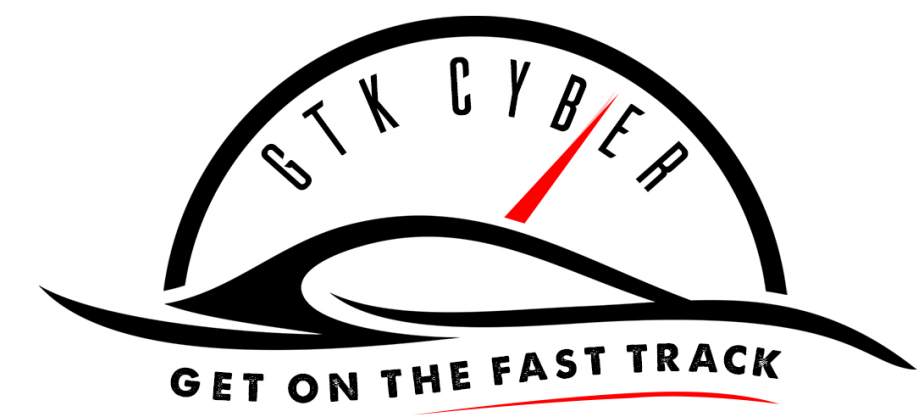
```
)
```

[gtkcyber.com](http://gtkcyber.com)



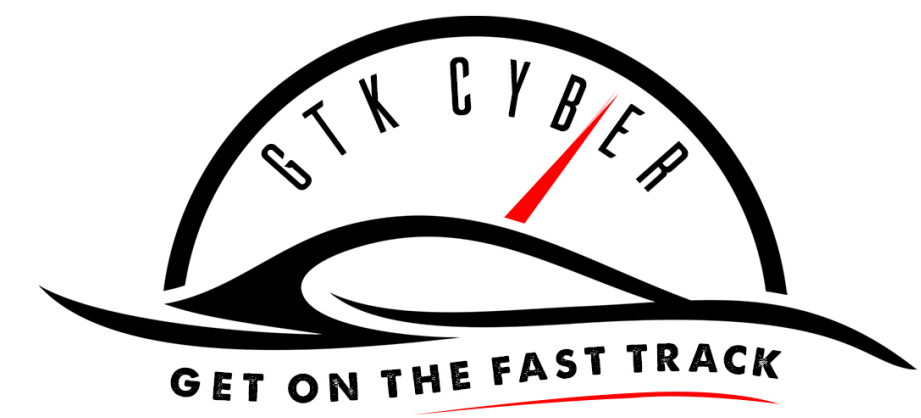
# Tuning these parameters

- GridSearchCV: You provide a list of possible parameters
- RandomizedSearchCV: Random combinations are searched



# Grid Search

```
param_grid = [  
    {'C': [1, 10, 100, 1000], 'kernel': ['linear']},  
    {'C': [1, 10, 100, 1000], 'gamma': [0.001, 0.0001],  
    'kernel': ['rbf']},  
]
```




# Grid Search

```
# use a full grid over all parameters
param_grid = {"max_depth": [3, None],
              "max_features": [1, 3, 10],
              "min_samples_split": [2, 3, 10],
              "min_samples_leaf": [1, 3, 10],
              "bootstrap": [True, False],
              "criterion": ["gini", "entropy"]}

# run grid search
grid_search = GridSearchCV(clf, param_grid=param_grid)
start = time()
grid_search.fit(X, y)

print("GridSearchCV took %.2f seconds for %d candidate parameter settings."
      % (time() - start, len(grid_search.cv_results_['params'])))
report(grid_search.cv_results_)
```



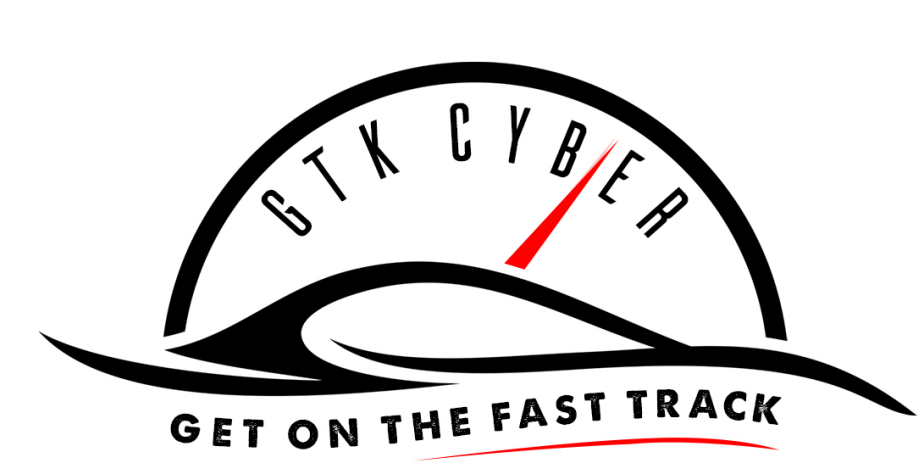


# Random Search

```
# specify parameters and distributions to sample from
param_dist = {"max_depth": [3, None],
              "max_features": sp_randint(1, 11),
              "min_samples_split": sp_randint(2, 11),
              "min_samples_leaf": sp_randint(1, 11),
              "bootstrap": [True, False],
              "criterion": ["gini", "entropy"]}

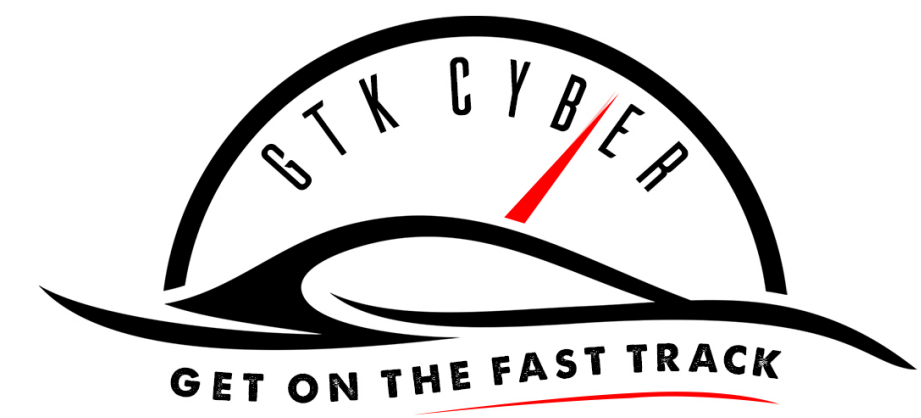
# run randomized search
n_iter_search = 20
random_search = RandomizedSearchCV(clf, param_distributions=param_dist,
                                   n_iter=n_iter_search)

start = time()
random_search.fit(X, y)
print("RandomizedSearchCV took %.2f seconds for %d candidates"
      " parameter settings." % ((time() - start), n_iter_search))
report(random_search.cv_results_)
```



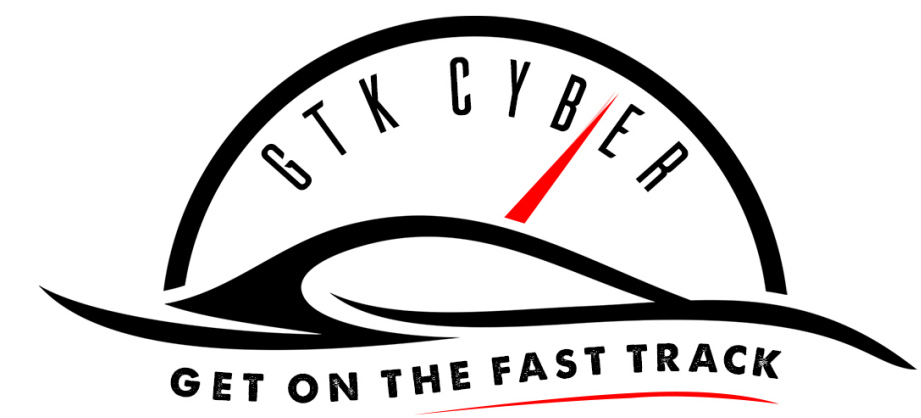
# What is a pipeline?





# Why use a pipeline?

- It makes code more readable
- You don't have to worry about keeping track data during intermediate steps, for example between transforming and estimating.
- It makes it trivial to move ordering of the pipeline pieces, or to swap pieces in and out.
- It allows you to do GridSearchCV on your workflow



# Without Pipeline

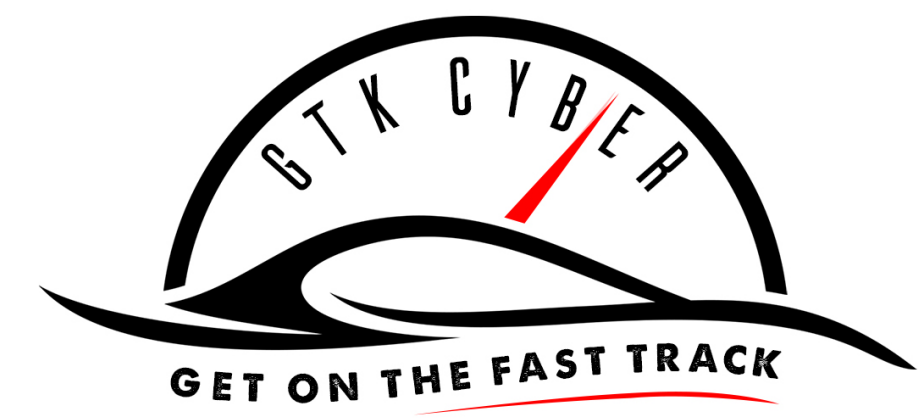
```
#get categorical features
#drop off last column because its unnecessary
X_categorical =
pd.get_dummies(df[categorical_columns]).astype(int).iloc[:, :-1]

#get and transform numeric features
X_numeric = df[numeric_columns]
X_numeric[numeric_columns] =
StandardScaler().fit_transform(X_numeric)

#get outcome variable
y = df[target]

#combine transformed categorical and numeric features
X_final = pd.concat((X_numeric, X_categorical), axis=1)
```

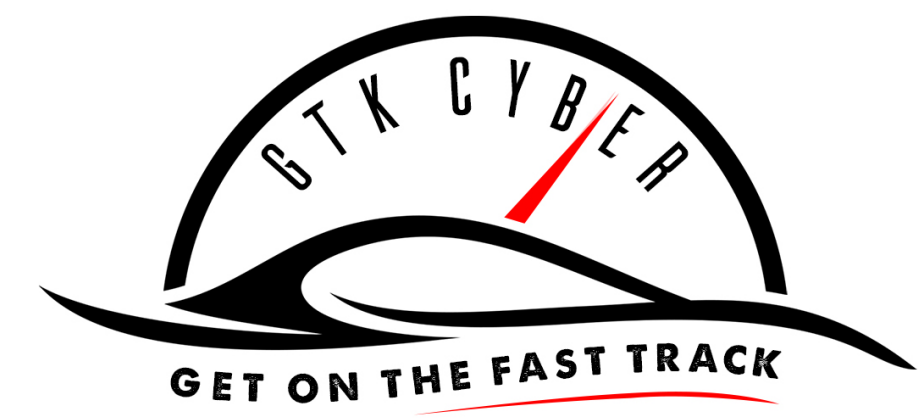




# Without Pipeline

```
#create rf regressor and check 10-fold RMSE
rf = RandomForestRegressor()
cross_val_scores =
np.abs(cross_val_score(rf,X_final,y,scoring =
"neg_mean_squared_error", cv=10))
rmse_cross_val_scores = np.sqrt(cross_val_scores)
```



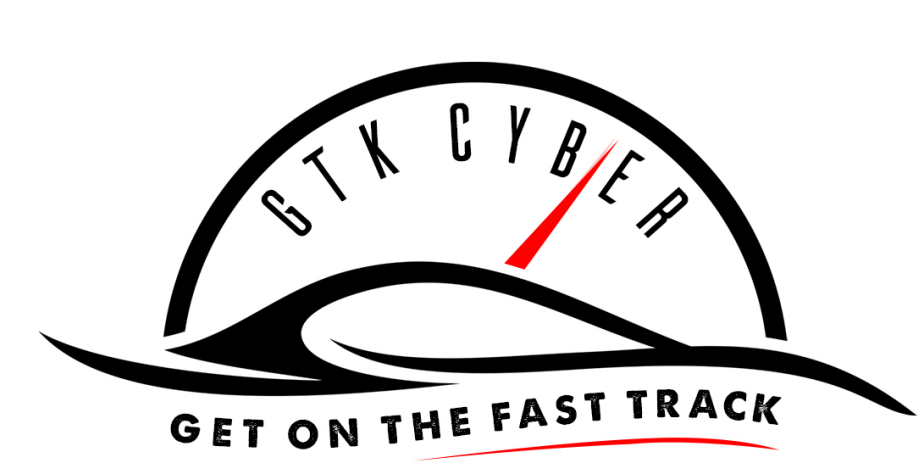


# With Pipeline

```
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier

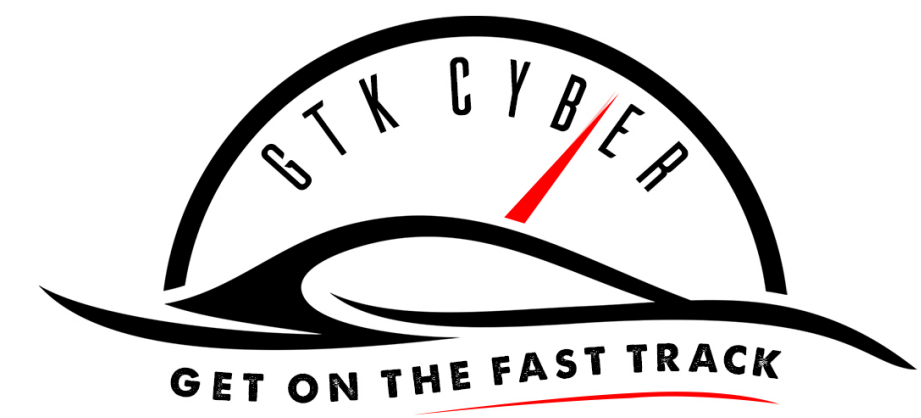
# it takes a list of tuples as parameter
pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('clf', KNeighborsClassifier())
])

pipeline.fit(X_train, y_train)
```



# With Pipeline

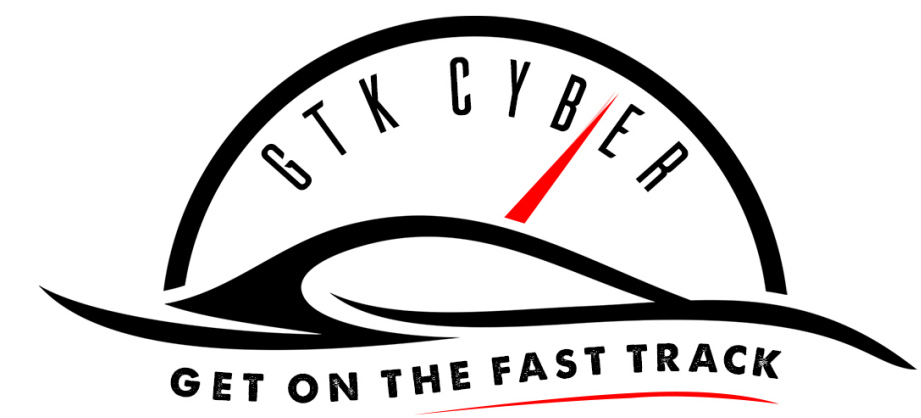
```
cross_val_scores =  
np.abs(cross_val_score(full_pipeline,X,y,cv=10,scoring="n  
eg_mean_squared_error"))  
rmse_cross_val_scores = np.sqrt(cross_val_scores)
```



# With Pipeline

## **full\_pipeline.steps**

```
[('all_features', FeatureUnion(n_jobs=1,
    transformer_list=[('categoricals', Pipeline(memory=None,
        steps=[('selector', ItemSelector(key=['rbc', 'pc', 'pcc', 'ba', 'htn',
            'dm', 'cad', 'appet', 'pe', 'ane'])), ('imputer', Imputer(axis=0, copy=True,
missing_values=0, strategy='most_frequent',
        verbose=0)), ('encoder', OneHotEncoder(cat ... tegy='median', verbose=0)),
('scaler', StandardScaler(copy=True, with_mean=True, with_std=True))])),
    transformer_weights=None)),
('rf_classifier',
    RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
        max_depth=None, max_features='auto', max_leaf_nodes=None,
        min_impurity_decrease=0.0, min_impurity_split=None,
        min_samples_leaf=1, min_samples_split=2,
        min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,
        oob_score=False, random_state=None, verbose=0,
        warm_start=False))]
```



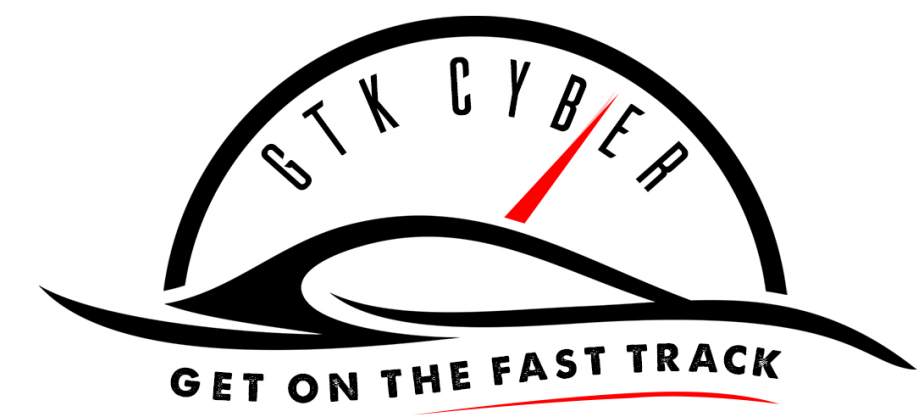
# Pipelines

```
from sklearn.feature_selection import SelectKBest
from sklearn.ensemble import RandomForestClassifier
from sklearn.pipeline import Pipeline
```

```
select = SelectKBest(k=100)
clf = RandomForestClassifier()
```

```
steps = [('feature_selection', select),
         ('random_forest', clf)]
```

```
pipeline = Pipeline(steps)
```



# Pipelines


```
pipeline.fit( X_train, y_train )
```

```
y_prediction = pipeline.predict( X_test )
```

```
report = classification_report( y_test, y_prediction )
```

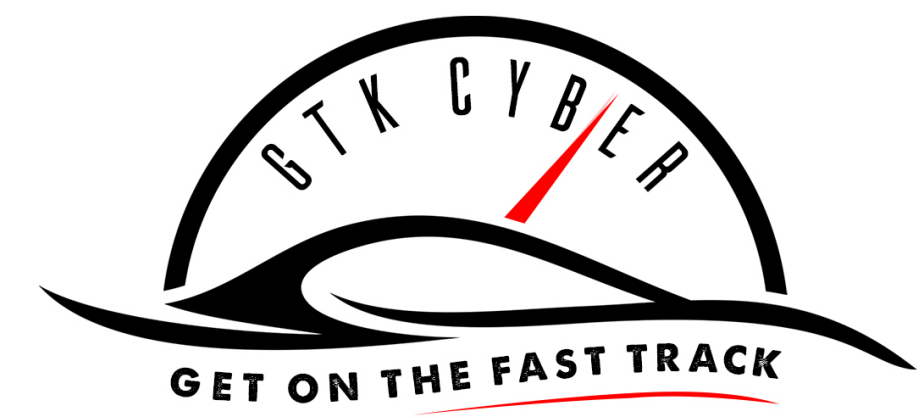
```
print(report)
```





# In Class Exercise

Please take 30 minutes and complete  
**Day 3: Optimizing your Model**

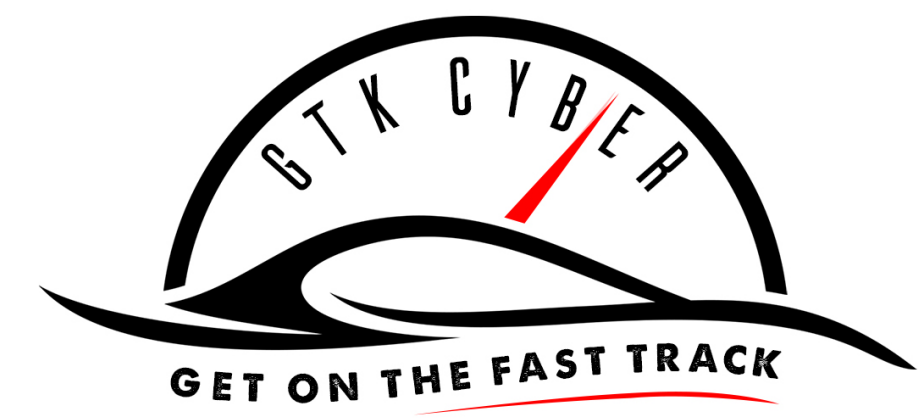


# Pickling Your Model

- Pickling your model allows you to preserve your model to be used later.





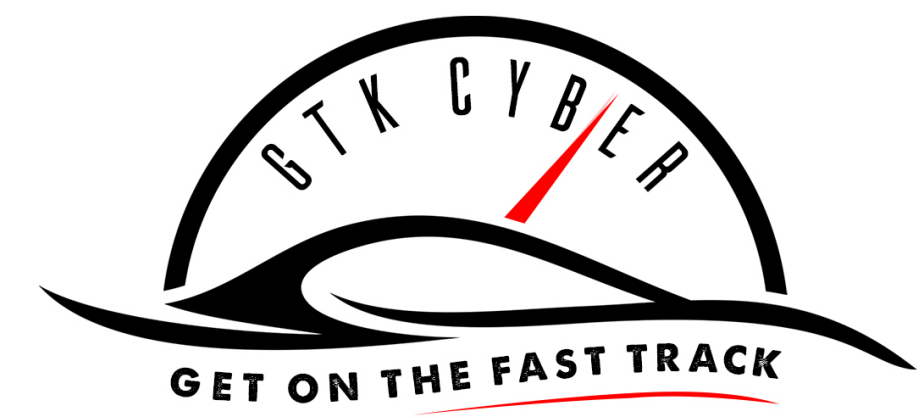


# Pickling Your Model

In order to rebuild a similar model with future versions of scikit-learn, additional metadata should be saved along the pickled model:

- The training data, e.g. a reference to a immutable snapshot
- The python source code used to generate the model
- The versions of scikit-learn and its dependencies
- The cross validation score obtained on the training data





# Pickling Your Model


#Saving your model

```
from sklearn.externals import joblib  
joblib.dump(clf, 'filename.pkl')
```

#Loading your model

```
clf = joblib.load('filename.pkl')
```





# Case Studies in Machine Learning & Cyber Security

GET ON THE FAST TRACK

A stylized speedometer graphic in the background. The speedometer has a semi-circular arc with letters 'C', 'Y', 'B', 'F', 'R' along the top. A red needle points towards the 'F'. Below the speedometer, the text 'GET ON THE FAST TRACK' is written in a light gray, sans-serif font, with 'FAST' underlined in red.

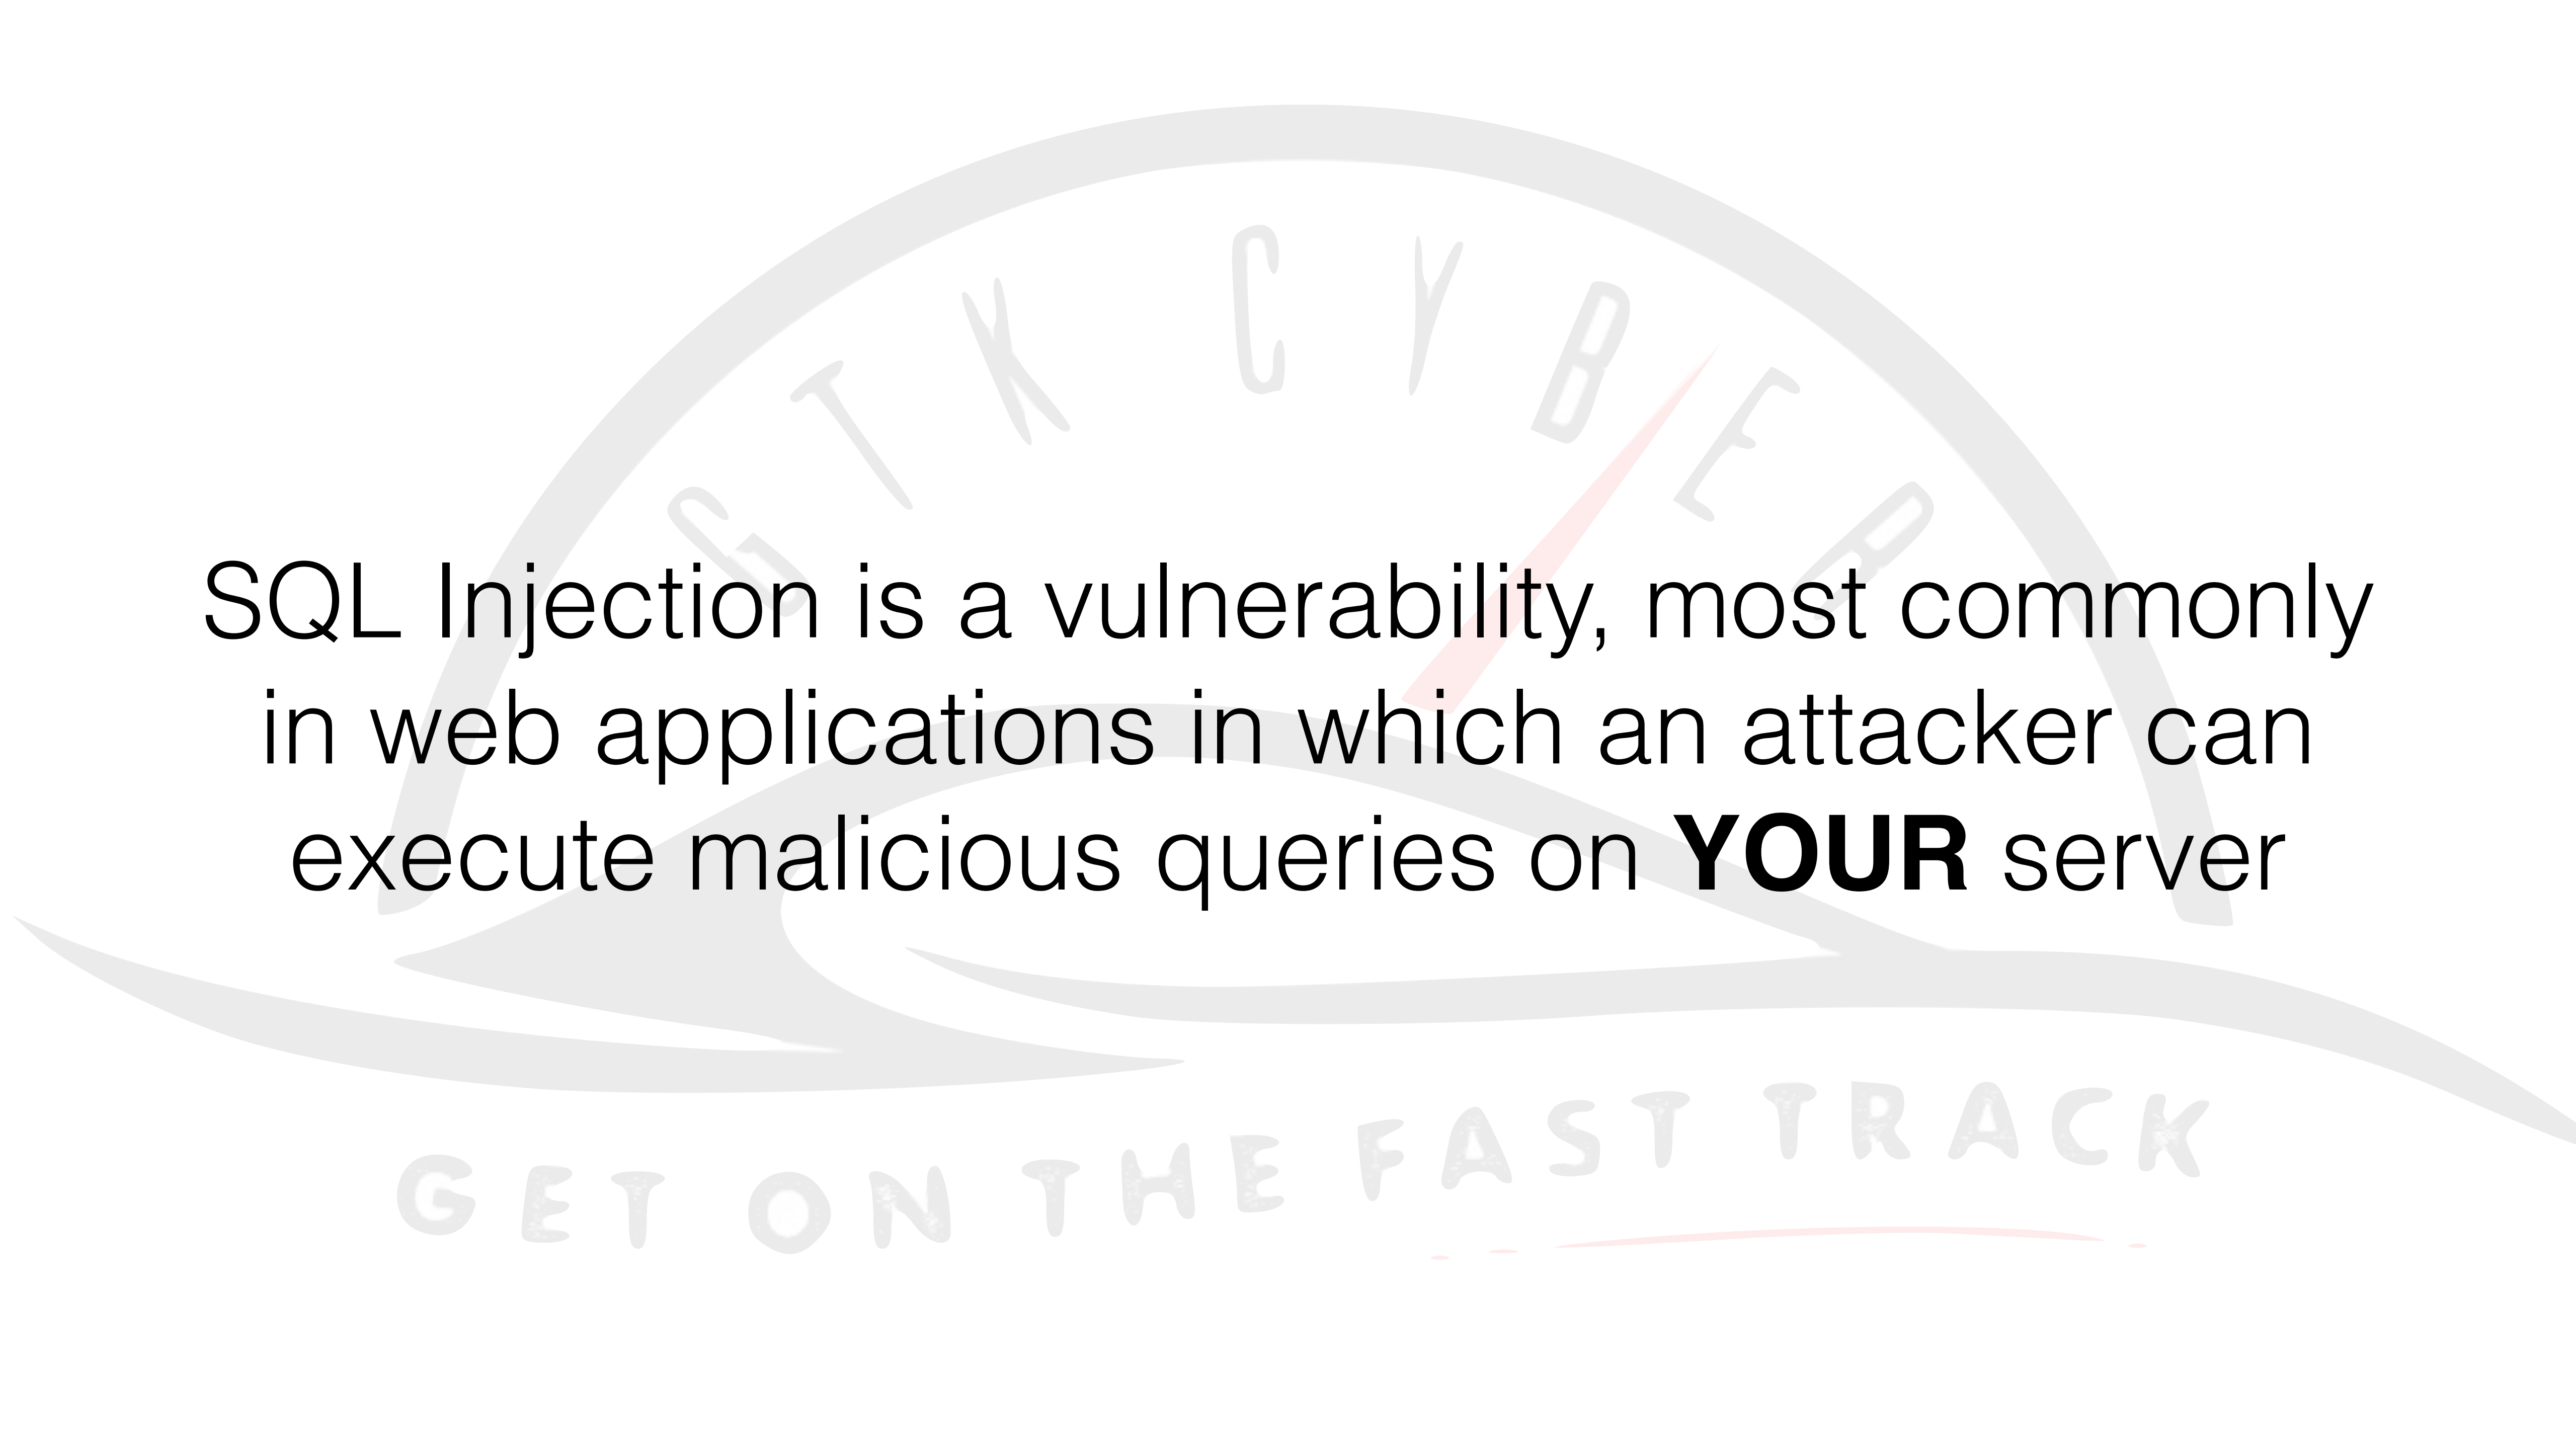
# Finding SQL Injection Attempts



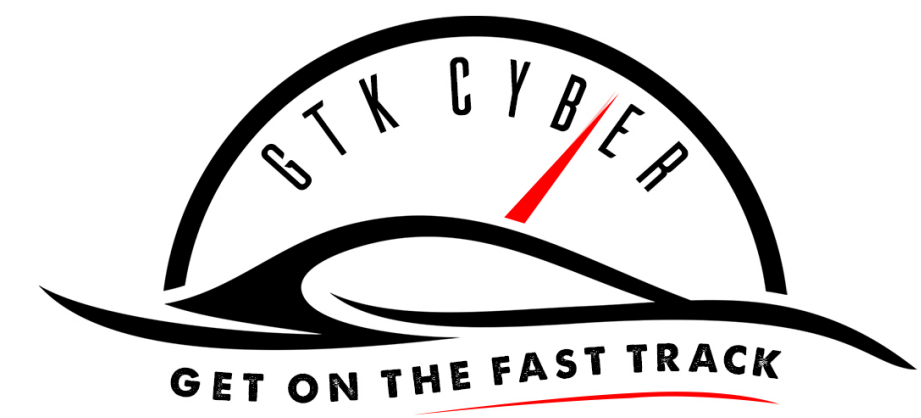
A stylized speedometer with a red needle pointing to 'F'. The speedometer has a semi-circular arc with letters 'C', 'Y', 'B', 'F', 'A' visible. Below the speedometer, the text 'GET ON THE FAST TRACK' is written in a bold, sans-serif font, with 'FAST' underlined.

What is SQL injection?

GET ON THE FAST TRACK

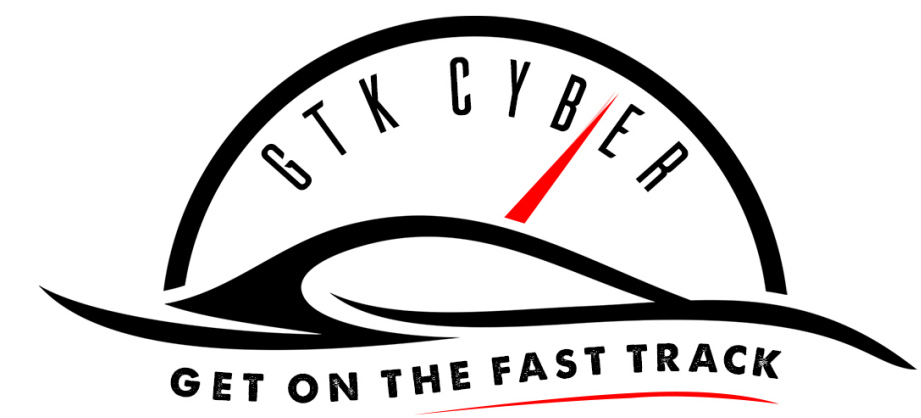


SQL Injection is a vulnerability, most commonly in web applications in which an attacker can execute malicious queries on **YOUR** server



# A normal SQL Query

```
SELECT *  
FROM users  
WHERE username=charles AND  
password=pass1234
```



# Pseudo Code for Web App Authentication

```
username = <from user>
```

```
password = <from user>
```

```
query = "SELECT * FROM users WHERE username =  
username AND password = password"
```

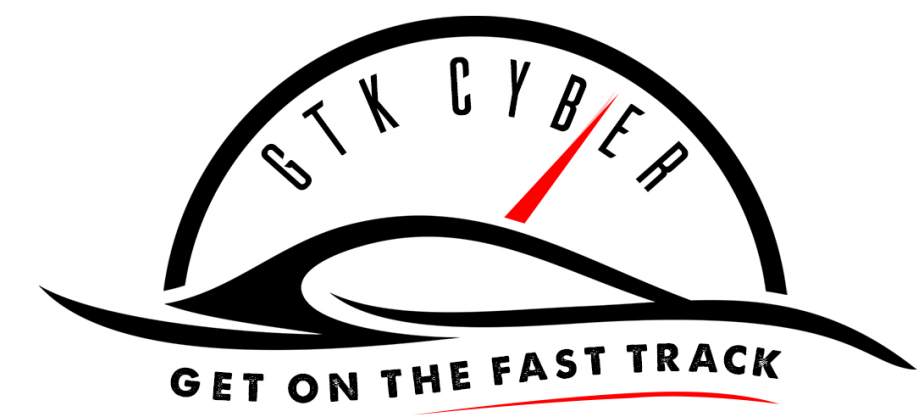
```
query_result = db.execute(query)
```

```
if len( query_result > 0 )
```

```
    //Authenticate user
```

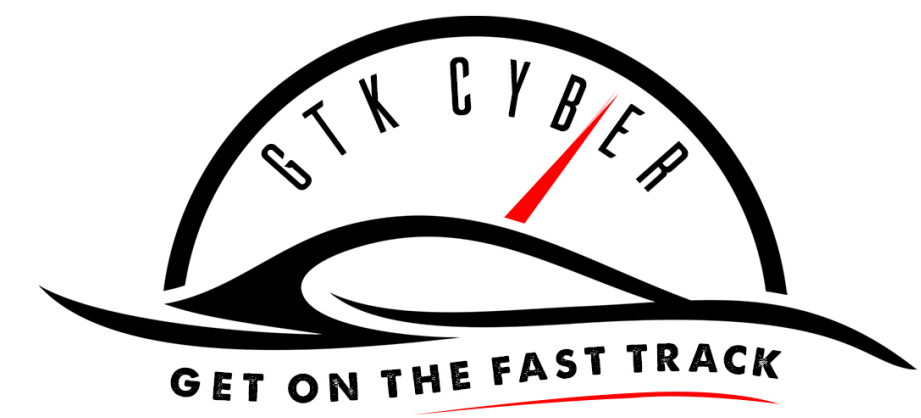
```
else
```

```
    //Boot them out
```



# Pseudo Code for Web App Authentication

```
username = "charles"  
password = "12345" //Combination an idiot would use on their luggage  
query = "SELECT *  
FROM users  
WHERE username = charles AND  
password = 12345"  
  
query_result = db.execute(query)  
if len( query_result > 0 )  
    //Authenticate user  
else  
    //Boot them out
```

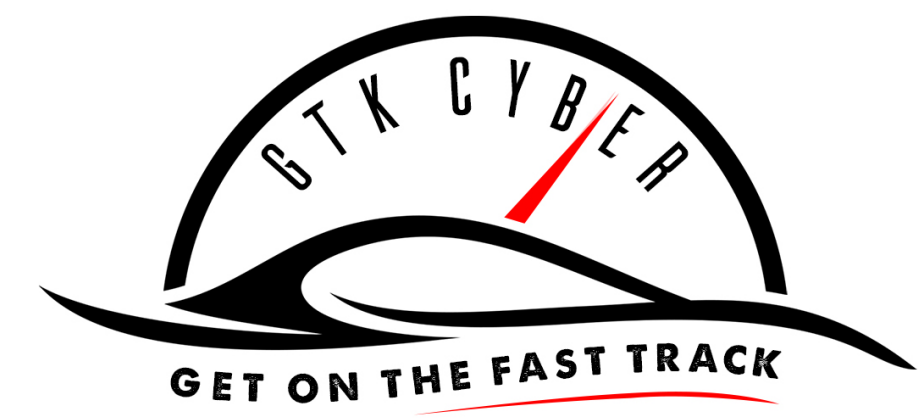


# Pseudo Code for Web App Authentication

```
username = "charles"
password = "12345 OR 1=1" //Combination an idiot would use on
their luggage
query = "SELECT *
FROM users
WHERE username = charles AND
password = 12345 OR 1=1"

query_result = db.execute(query)
if len( query_result > 0 )
    //Authenticate user
else
    //Boot them out
```



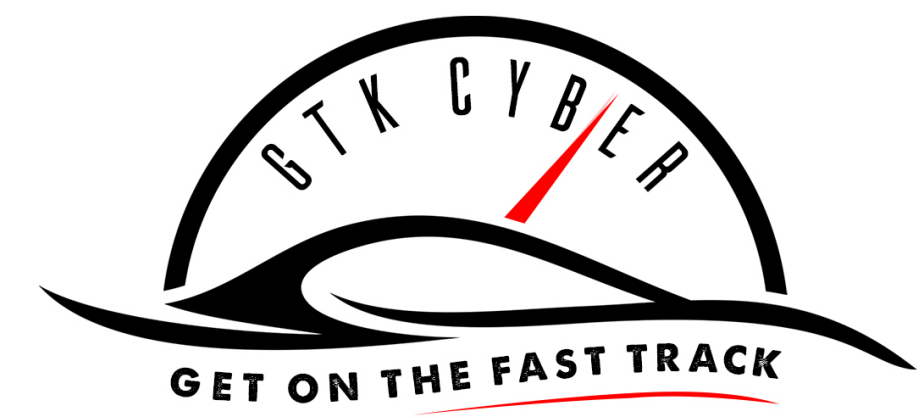


# Legit vs. Malicious

```
SELECT count(category) FROM Product WHERE price >  
' $20 '
```

```
SELECT ctid, xmin, * FROM lockdemo
```

```
SELECT current_date FROM dual
```




# Legit vs. Malicious

' or 'unusual' = 'unusual'

' or 'something' = 'some'+'thing'

' or 'text' = n'text'

' or 'something' like 'some%'




# Step 1: Tokenize SQL


```
import sqlparse
def parse_it(raw_sql):
    parsed = sqlparse.parse(unicode(raw_sql, 'utf-8'))
    return [token._get_repr_name() for parse in parsed for token in
parse.tokens if token._get_repr_name() != 'Whitespace']

dataframe['parsed_sql'] = dataframe['raw_sql'].map(lambda x: parse_it(x))
```

	raw_sql	type	parsed_sql
0	; exec master..xp_cmdshell 'ping	malicious	[Single, Identifier, Float, Float, Float,
1	create user name identified by	malicious	[DDL, Keyword, Identifier, Keyword,
2	create user name identified by pass123	malicious	[DDL, Keyword, Identifier, Keyword,
3	exec sp_addlogin 'name' , 'password'	malicious	[Keyword, Identifier, IdentifierList]
4	exec sp_addsrvrolemember 'name' ,	malicious	[Keyword, Identifier, IdentifierList]

# Step 2: Create N-Grams

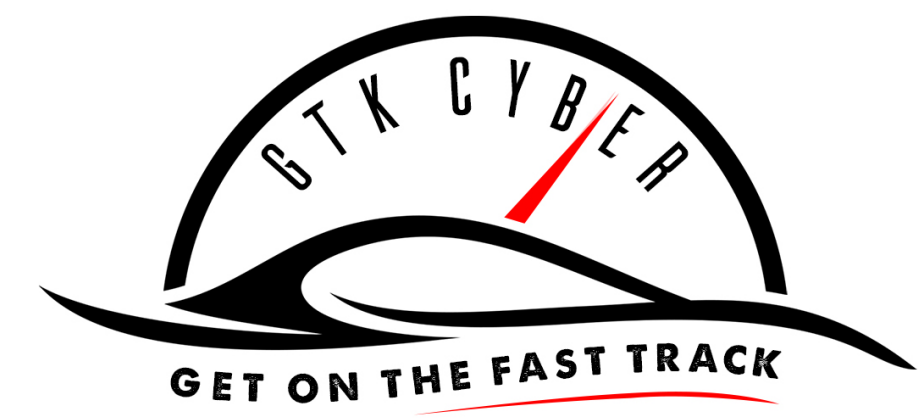




# Step 2: Create N-Grams

	raw_sql	type	parsed_sql	sequences
0	; exec master..xp_cmdshell 'ping 10.10.1.2'--	malicious	[Single, Identifier, Float, Float, Float, Erro...	[('Single',), ('Identifier',), ('Float',), ('F...
44	anything' or 'x'='x	malicious	[Identifier, Single, Identifier, Single, Ident...	[('Identifier',), ('Single',), ('Identifier',)...]
49	; exec master..xp_cmdshell 'ping aaa.bbb.ccc....	malicious	[Single, Identifier, Error, Single]	[('Single',), ('Identifier',), ('Error',), ('S...
54	; if not(select system_user) <> 'sa' waitfor ...	malicious	[Single, Identifier, Single, Integer, Placehol...	[('Single',), ('Identifier',), ('Single',), ('...
55	; if is_srvrolemember('sysadmin') > 0 waitfor...	malicious	[Single, Identifier, Single, Integer, Placehol...	[('Single',), ('Identifier',), ('Single',), ('...




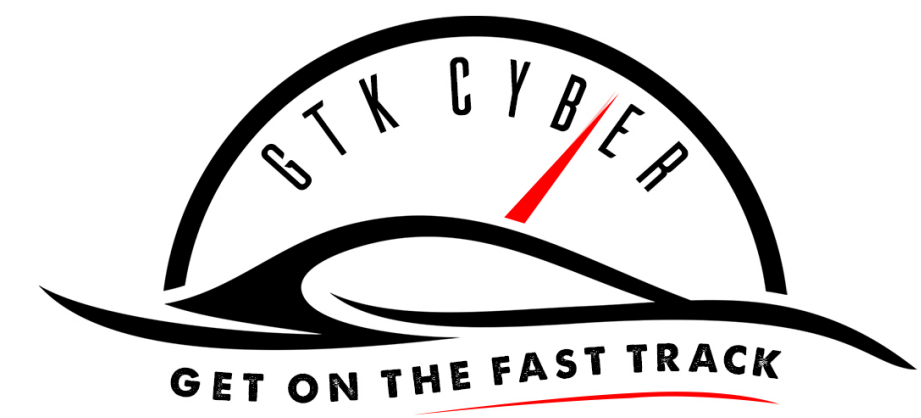


# Step 3: Build Feature Vector


- Entropy
- Length
- G-Score Test

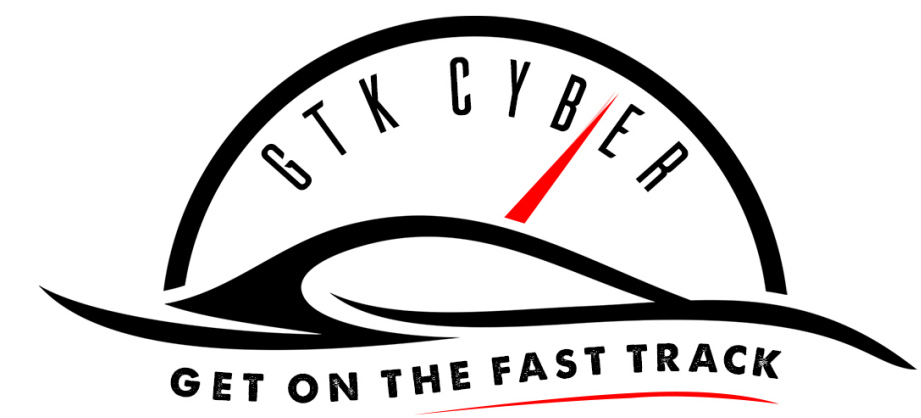
# Step 3: Build Feature Vector






# Step 3: Build Feature Vector

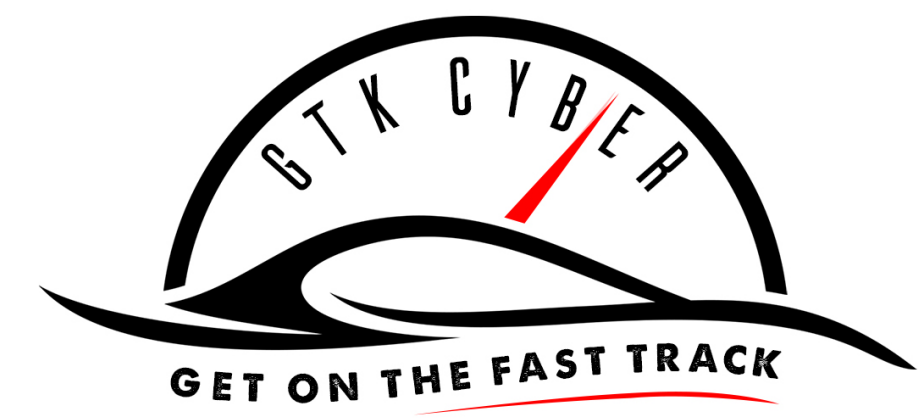





# Step 3: Build Feature Vector

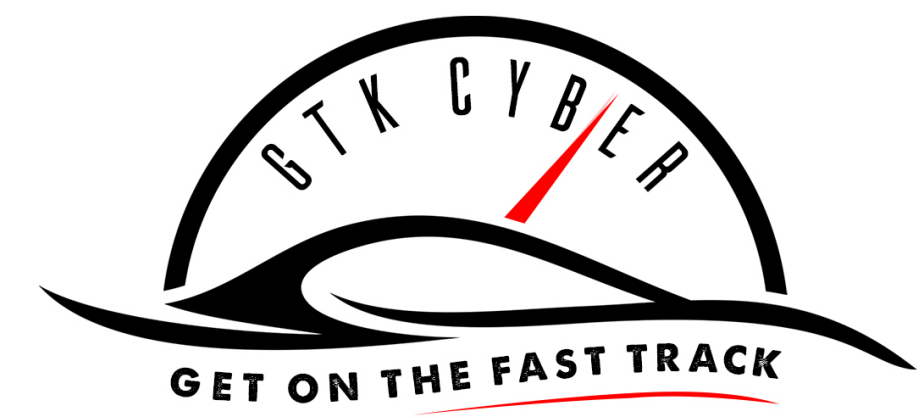






# Step 3: Build Feature Vector



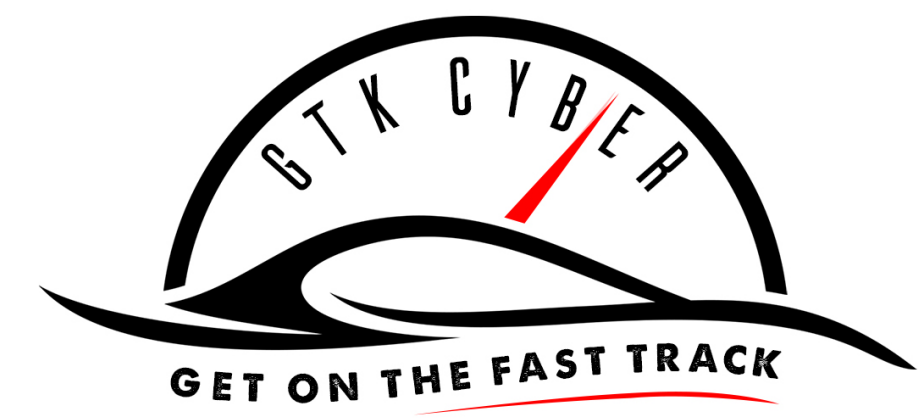


# Step 4: Train Classifier

```
import sklearn.ensemble
clf =
sklearn.ensemble.RandomForestClassifier(n_estimators=20)

scores = sklearn.cross_validation.cross_val_score(clf, X, y,
cv=10, n_jobs=4)
print scores
```







# Step 5: Evaluate Model


```
scores = sklearn.cross_validation.cross_val_score(clf, X, y,  
cv=10, n_jobs=4)
```

```
[ 0.99784173  0.99784173  1.          0.99784173  
0.99856115  0.99784017  
0.99640029  0.99856012  0.99784017  0.99784017]
```



# Step 5: Evaluate Model





Thank you!